STAT 330-001

40142034 47

Doney, Samuel Jay

895072291 doney1 425-530-3128

Test: 640586 Form: A Wadsworth, Spencer Gordon

Entry: 03:49pm 10/21/2026 Checked out:

Employee: Megan S.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly.
Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

047

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

Lasso and Ridge regression are sometimes referred to as "shrinkage" methods because they "shrink" or reduce the magnitude of the coefficients in a linear model. This is done by satisfying a separate equation than standard SSE to introduce bias and therefore reducing variance in our model.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\hat{\beta}$?

In matrix notation

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

(B) What is the conceptual difference between $\hat{\beta}$ and β?

$\hat{\beta}$ represents estimated coefficients calculated from our sample to minimize the SSE. β represents the actual population parameter, representing the true coefficient that associates our predictors with the response.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least squares can be used as a measure to test differing values of $\hat{\beta}$. In particular, we aim to minimize the value of $\sum_{i=1}^{n} e_i^2$ by selecting $\hat{\beta}$s that do so. In a sense, we fit the model, find the sum of the residuals for all $\hat{\beta}$s, and select the $\hat{\beta}$s with the smallest sum of squared errors. (setting derivative = 0 and solving).

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

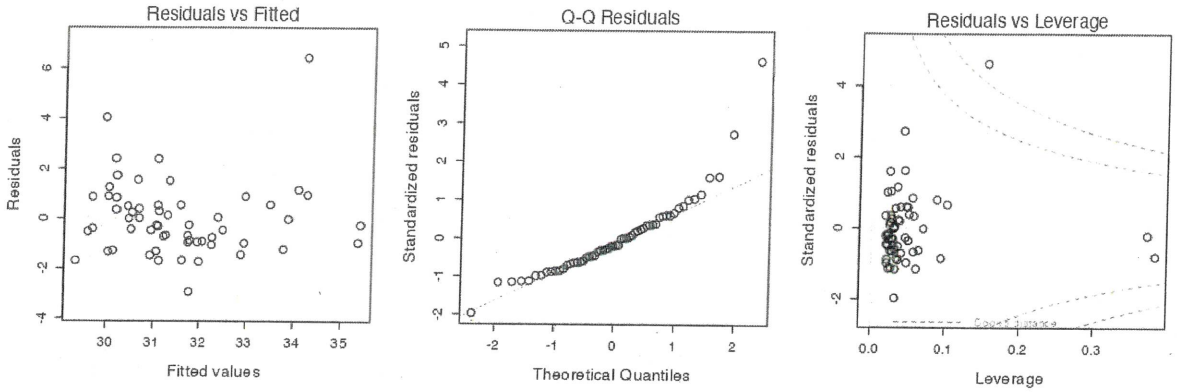

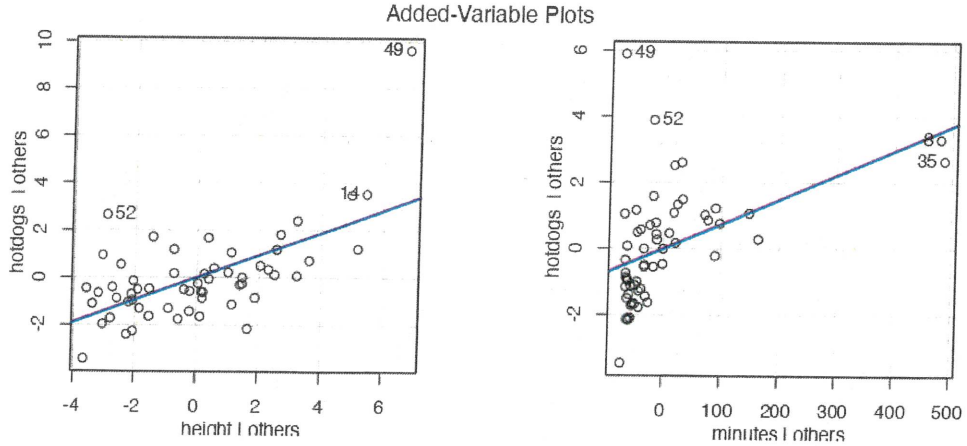

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 \approx 0.473$

Holding minutes constant, a 1-inch increase in height is associated with an estimated 0.473 increase in hotdogs consumed in the 10-minute window, on average.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

- In general, the residual vs. fitted values appear to be random with no clear trend up or down
- The height AV plots seems like a cloud following the line, but the minutes has a tight cluster of points. Nothing terrible, but proceed with caution.

(B) There is independence between each observation

The data was not randomly collected, as all contestants chose to join the competition. They also may learn from, practice, and observe one another. I would say independence is violated here.

(C) The errors are normally distributed

The Q-Q plot suggests very strong tails and skew, and the shapiro-wilk test had a p-value well below 0.05 despite only having 58 samples. Both of these point to a violation of normally distributed errors condition.

(D) The errors have constant variance (homoscedasticity)

Using the residuals vs. fitted plot, there is no clear trend suggesting the variance increases or decreases with the fitted values. The Breusch pagan test has a p-value of ≈ 0.07, which is larger than 0.05. There doesn't seem to be a violation of the constant variance assumption here.

(E) There are no unduly influential points

There seems to be about 3 points with significant residuals, and 3 with high leverage. Although all are worth looking into, there is 1 point with both high leverage and large residual that is worth investigating. This assumption is violated → need to investigate.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

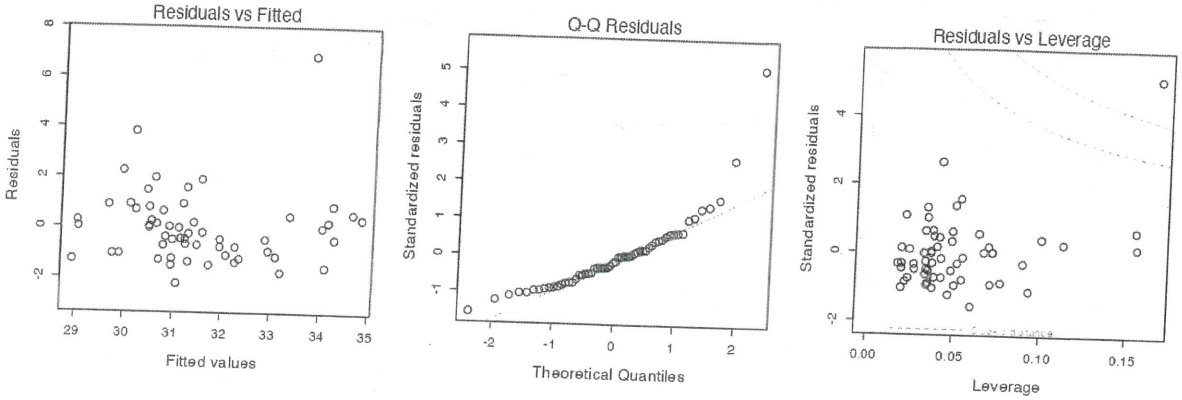

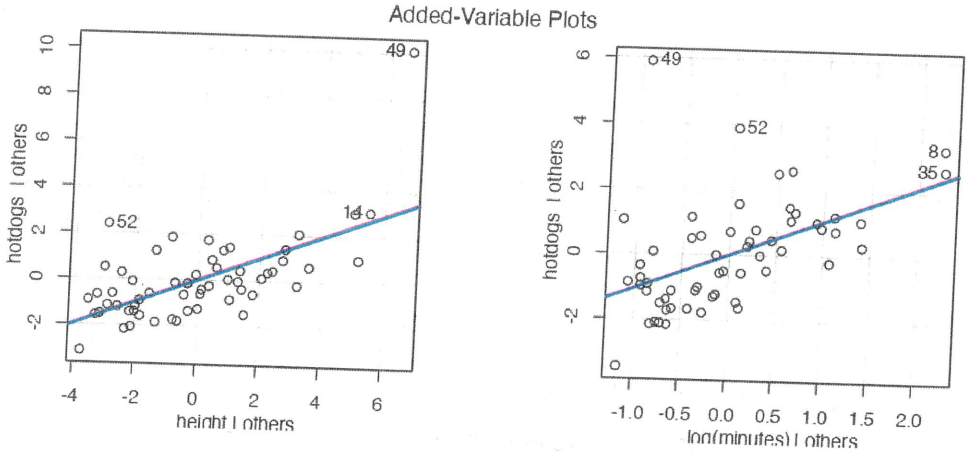

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

This is an exaggerated statement. A p-value below 0.05 does not justify choosing another model over this one. There are other LINE assumptions we need to check. This difference is small to begin with; we don't have enough information to make a claim like this yet.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages:
- Larger Shapiro-wilk p-value
- Smaller CV RMSE
- Better AV and fitted vs resid pots (better linearity)

Disadvantages
- smaller Breusch pagan p-value (equal variance at risk)
- More complex model because log is introduced.
- A point with a very large Cook's distance (very influential)

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$

Accounting for height, a 10% increase in weight is associated with an estimated increase of $1.06702 \log(1.10) \sim 0.1017$ in the mean of hotdogs eaten.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$\hat{y}_{friend} = -6.05589 + 0.48834(68) + 4.369$

≈ 31.52023

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

$n - p - 1$

$58 - 2 - 1 = 55$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

estimate $\pm$ t^* · SE

$31.52023 \pm 2.004 \cdot 1.485$

(28.54429, 34.49617)

lower upper

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

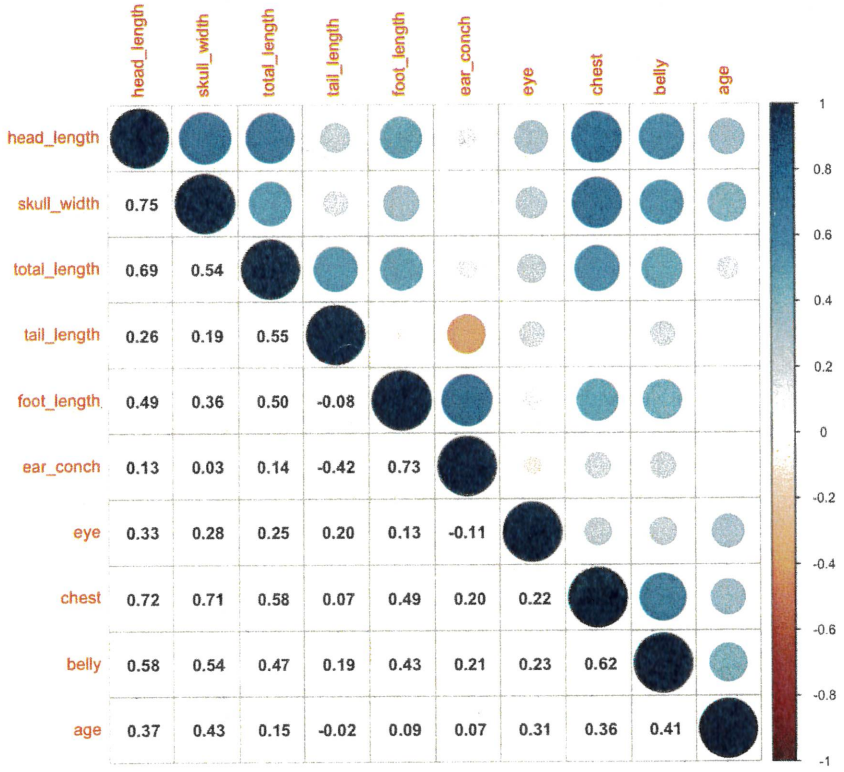

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 2.01 \cdot 0.10599$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

Holding all other predictors constant, we are 95% confident that there is no significant association between belly width and age, given the interval contains 0.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

F-tests and t-tests are not testing the exact same thing. An F-test tests whether ALL coefficients are 0, while a t-tests performs this one at a time. An F-test only needs a single coefficient to be not 0 in order to justify a small p-value. Thus, it makes sense the F-test is significant.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets evaluates every possible linear model (2^p) to determine which one has the most desirable result based on the metric of choice (AIC, in this case). Thus, it calculates the AIC for all possible models and recommends the models with the best criterion.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 46

Ventura, Coltran D

457489496 cventur4 385-210-9309

Test: 640426 Form: A Wadsworth, Spencer Gordon

Entry: 03:17pm 10/21/2026 Checked out:

Employee: Ty B.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly.
Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

046

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

LASSO and Ridge are refferred to as shrinkage methods because they are shrinking the coefficient estimates towards zero. This is accomplished by computing Ridge $= SSE + \lambda \Sigma \beta_i^2$ and LASSO $= SSE + \lambda \Sigma |\beta_i|$ where λ is the penalty factor that is chosen by how much to shrink the estimates by

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}_i$ is the estimated coefficient value associated with variable X_i

β is the true coefficient value associated with the variable X_i

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

least squares is used to estimate regression coefficients in matrix form $SSE = (y-\hat{y})^T(y-\hat{y})$ $\hat{y} = X\hat{\beta}$. Using the SSE formula we can derive the normal equations $(X^TX)\hat{\beta} = X^TY$ if (X^TX) is invertible, therefore, $\hat{\beta} = (X^TX)^{-1}X^TY$ and we can use the least squares to estimate the values of each $\hat{\beta}_j$. To be clear the least squares minimizes the sum of the residuals (differences between y values and fitted values)

who considers this? I have never met someone who thought this ...

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

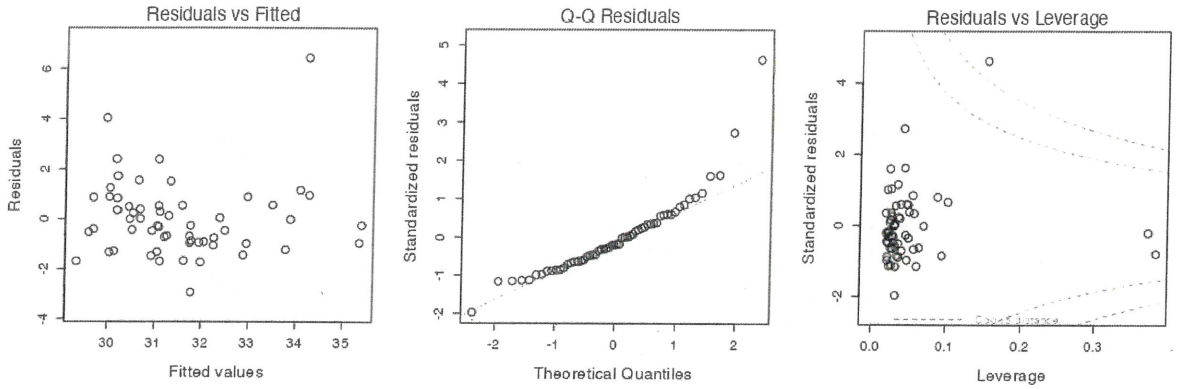

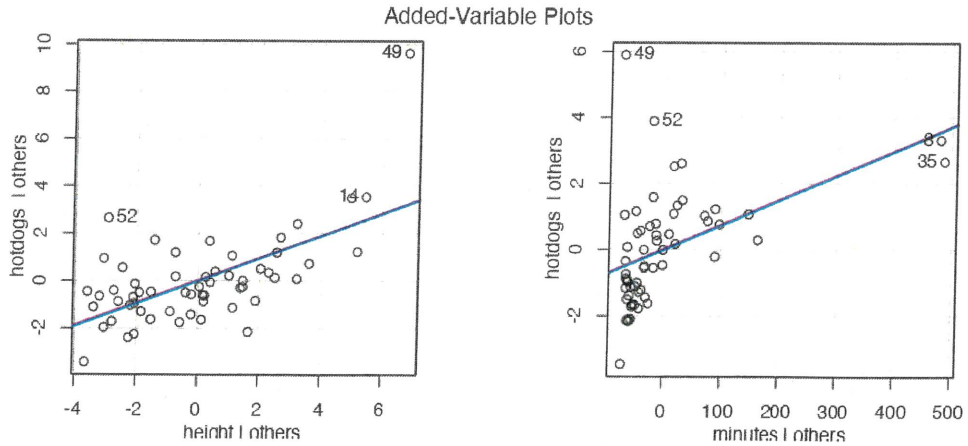

Coefficient	Estimate	Std. Error	t value	Pr($>\|t\|$)
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

Holding minutes constant, a one-unit increase in height in inches results in a 0.472892 increase in amount of hotdogs that are eaten.

The estimated value of β_1 is 0.472892

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0 I'm assuming alpha = .05

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The residuals vs fitted plot has its data points centered around zero for the most part and looks mostly randomly scattered. I could see an argument to be made for a slight funnel-like taper to the right though. I will say the model is mostly linear.

(B) There is independence between each observation

Yes there is independence between the observations. the amount of hot dogs eaten by one individual is not dependent upon the amount eaten by another.

(C) The errors are normally distributed

The QQ-plot shows most of the data on the line with slight upward curves on both end. Maybe it's normally distributed but lets look at Shapiro-Wilkes. The Shapiro-Wilkes test has a small p-value well below .05 or any other reasonable alpha-level. This gives reason to reject H_0: that the errors are normally distributed.

(D) The errors have constant variance (homoscedasticity)

Again, the residuals vs fitted plot has the data centered at zero and mostly random. To better check for homoscedasticity I will look at the results of the Breusch-Pagan test. The H_0: of the BP test is equal variance using an alpha-level of 0.05 The p-value of .071 suggests we fail to reject H_0: and conclude that this model does have equal variance

(E) There are no unduly influential points

There are no unduly influential points. On the residuals vs leverage plot there are points with large residuals and points with large leverage, but no points that stick out with both a large residual and large leverage. I'm determining this as a lack of influence using the 5 rooks's distance calculation.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

level-log

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

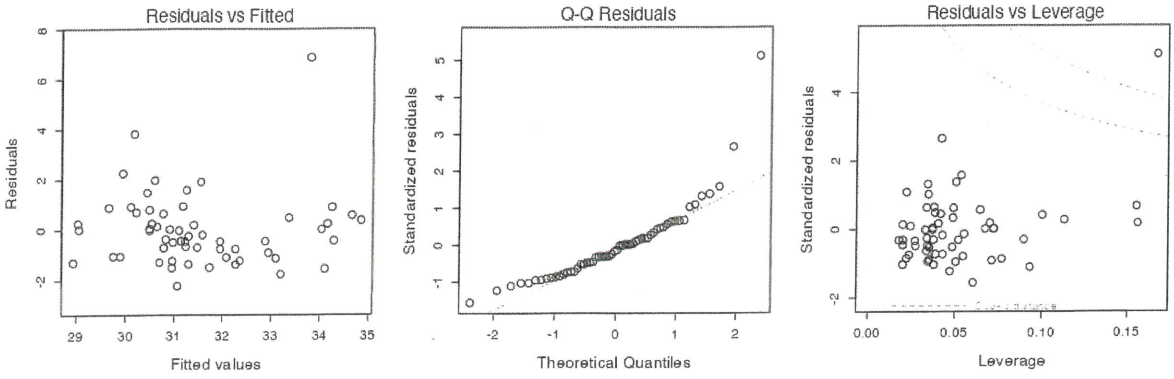

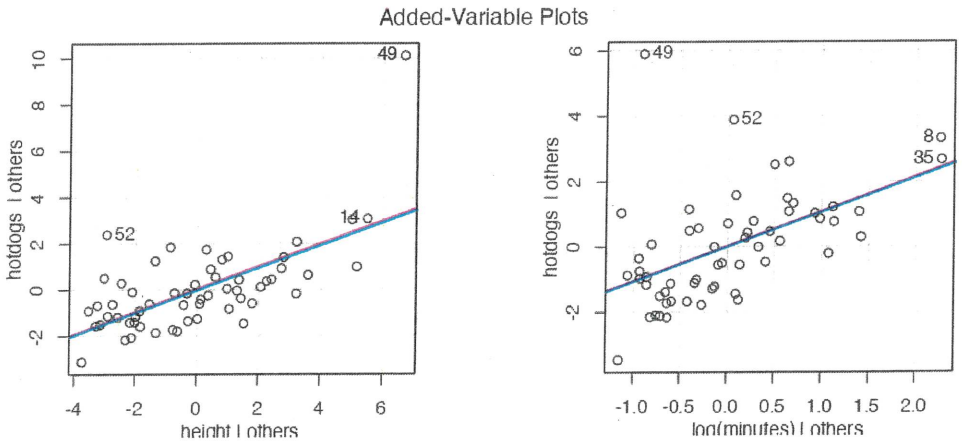

Coefficient	Estimate	Std. Error	t value	Pr($>\|t\|$)
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond? "Clearly Superior" may be too strong of wording/conclusion. The new BP p-value of 0.028 suggest that we reject H_0 and say that there could be some unequal variance in the new level-log model. I would tell the research the new level-log model could actually be inferior to the previous model.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages of Model 2 over Model 1

No real advantages I can see. I doesn't seem that the transformation improved the model.

Disadvantages of Model 2 compared to model 1

- Unequal variance
- Potentially influential points (using Cook's distance)

Question 16: State the estimated value of β_2 and provide an interpretation in context

Holding height constant, p% increase in minutes yields a $.06702 \ln\left(1+\frac{p}{100}\right)$ unit change in number of hot dogs. For a small increase the change is equivalently a .01β increase in the number of hot dogs eaten for every 1% increase in minutes.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes.

$\widehat{hotdogs}_{minutes} = \hat{\beta}_0 + \hat{\beta}_1 height + \hat{\beta}_2 log(minutes)$

$.01(1.06702) = 0.0106702$

$= -6.05589 + 0.148841(68) + 0.106702(4.369)$

$= 27.198$ hot dogs

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

$CI = \pm t \cdot SE(\hat{\beta})$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^TX)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^TX)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$PI = \hat{y} \pm ME$

$ME = t^* SE$

$SE = \frac{s}{\sqrt{n}}$

$n = 58$

$x(x^Tx)^{-1}x = H = leverage$

$27.2 \pm \frac{1.485}{.204} 2.004$

$= 27.2 \pm (1.28) 2.004$

27.2 ± 14.6

$= (12.6, 41.8)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

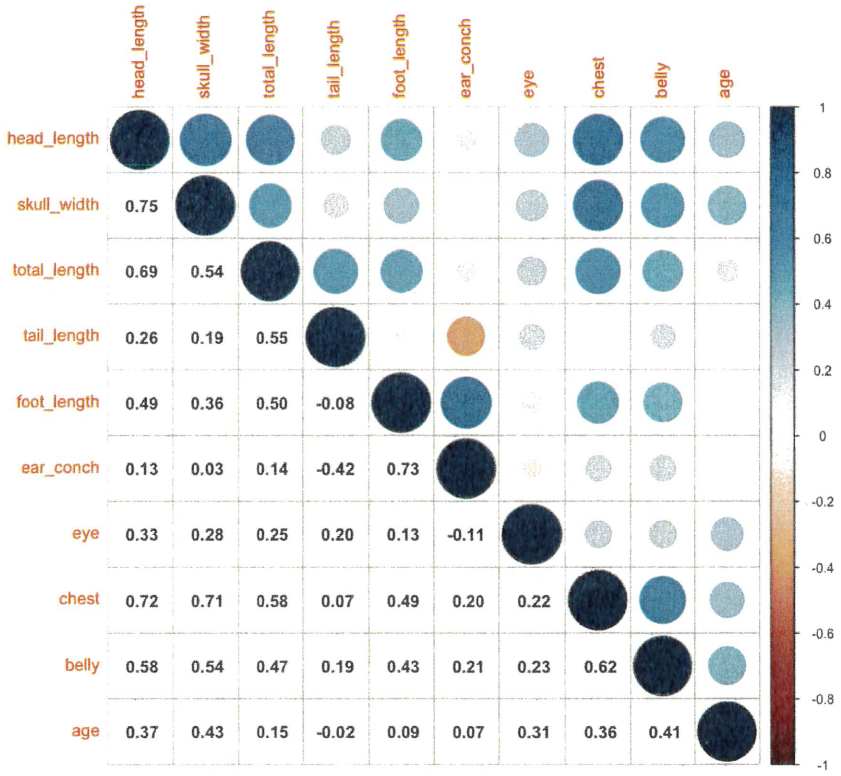

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

skull_width estimate = .16559 ± 2.01 x (.16559)

(-0.167 , 0.498)

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

I'm confident that if repeated many times the true coefficient on belly width would be between -0.023 and .443 about 95% of the time since 0 is in the interval.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

This is an indicator that there is probably multicollinearity in this model. The model is over-fit, meaning, that the data fits well with our sample but the model as is wouldn't do a good job at predicting age because of the redundant predictors creating noise in the model.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The best subsets algorithm uses $AIC = n \log\left(\frac{SSE}{n}\right) + 2k$ with k as the penalty value to looks at every possible combination of predictors in the model which give RMSE values. Using AIC which model gives the smallest RMSE as the metric for which model to use. the smaller the AIC the better

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

B Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

B Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

Scratch Paper

Scratch Paper

STAT 330-001

40142034 45

Gardiner, Monte

822808742 mbgardin 801-836-6328

Test: 640419 Form: A
Entry: 03:16pm 10/21/2026
Employee: Ty B.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

5:08 → 5:10

Monte Gardner

40142034

045

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: **(Circle all that apply)** Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: **(Circle one answer)** What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: **(Circle all that apply)** In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

LASSO and ridge regression are referred to as "shrinkage" methods because the utilize a penalty λ to "shrink" the coefficients towards 0 in the case of ridge, and even to 0 in the case of LASSO. $\lambda \Sigma \beta_i^2$ Ridge $\lambda \Sigma |\beta_i|$ Lasso

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$D_i = \frac{2.4}{3+1}\left(\frac{.25}{1-.25}\right)$ $.6\ (.33) = .2$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is the fitted coefficient in our fitted model, whereas β is the predicted coefficient in our predicted model.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

OLS minimizes the function $\sum e_i^2$ helping us estimate our regression coefficients similarly to best subsets.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

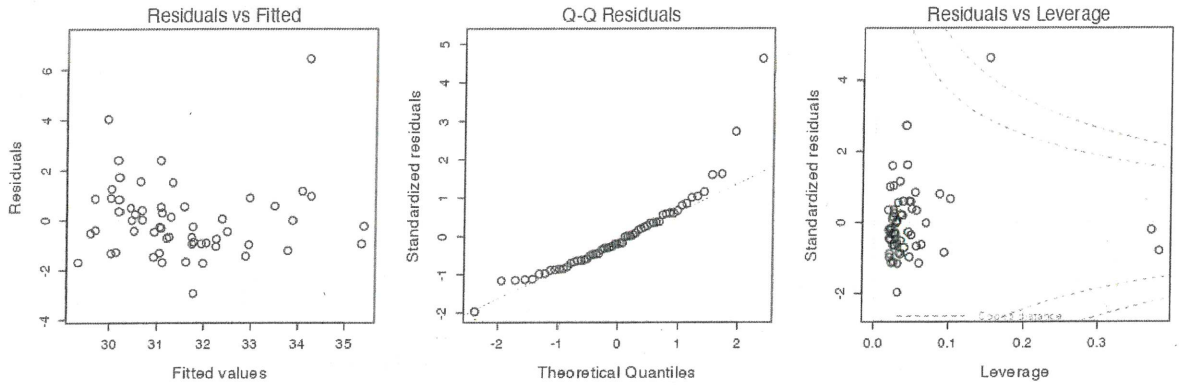

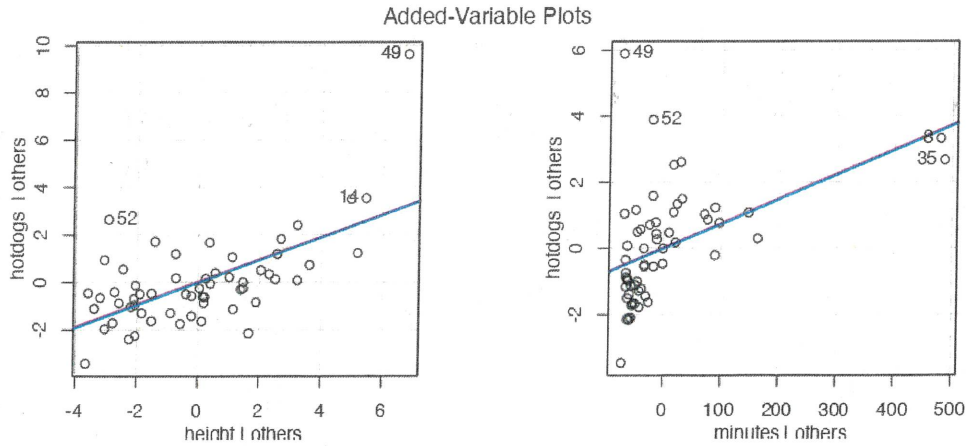

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

Estimated $\beta_1 \simeq .472892$, for each additional 1 inch (height), holding all other predictors constant, about $\hat{\beta}_1 \simeq .472892$ more hot dogs are consumed.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

$t = 5.885$

$SE = .080352$

$t \cdot SE = 5.885(.080352) = .47287152$

$.472892 - \hat{\beta} = .47287152 = c_1$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

Looking at the Residuals vs Fitted and feeling the vibes, I see no warning signs of curvature, so it satisfies linearity.

(B) There is independence between each observation

Looking at the design, many of the 58 are from the 4 largest counties, something to keep in mind but still satisfies independence.

(C) The errors are normally distributed

Looking at the Shapiro-Wilk p-value, I am deeply concerned by $2.43e-8$, which is much smaller than .05, doesn't satisfy very well. Investigation required.

(D) The errors have constant variance (homoscedasticity)

Looking at the Breusch-Pagan p-value, I see $p = .071$, which is greater than .05, so it is satisfied well.

(E) There are no unduly influential points

There are definately outliers, visually apparent in the datas, but none that appear to be unduly influential other than the single value seen outside of the Cooks distance in the Residuals vs Leverage chart.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

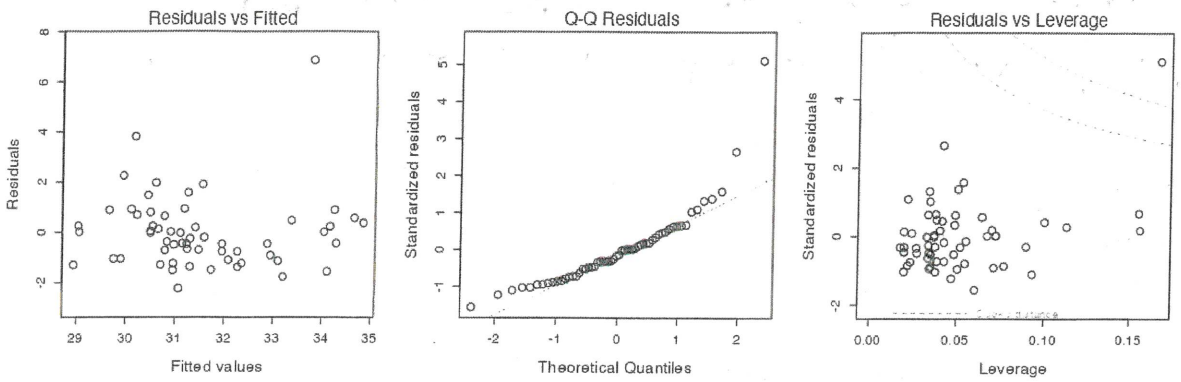

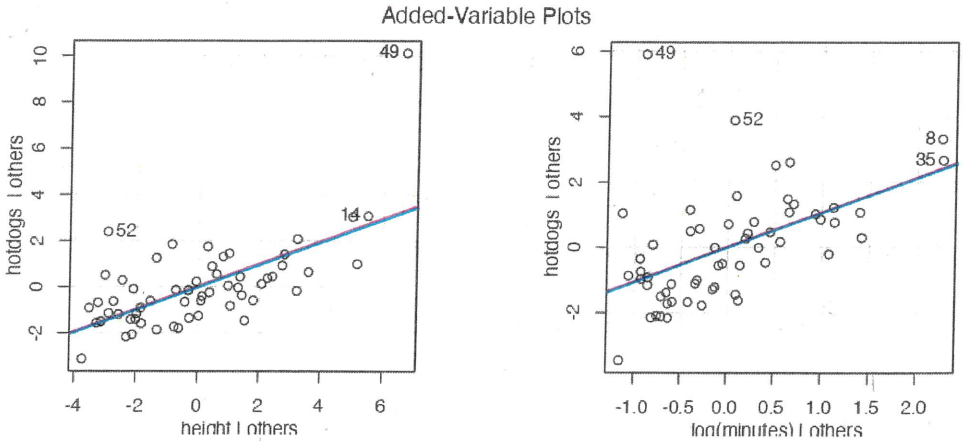

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would respond by kindly reminding them that the Breusch-Pagan test is not the only variable in play for determining model effectiveness. A p value less than .05 does not merit throwing out a model, but does merit farthur investigation.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

The model(2) has more linearity than model(1) in it's added variable plots, and a marginally smaller CV RMSE value, suggesting higher performance. However, the Breusch Pagan P value is too high and some of the other values are harder to interpret, conditional median, not necessarily mean.

Question 16: State the estimated value of β_2 and provide an interpretation in context

Estimated β_2 is 1.06702, meaning that for each additional log(minute), holding all other factors constant, 1.06702 more hot dogs are consumed.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

My friend can consume a whopping 31.813 68, 4.369 or ≈ 32 hot dogs in 10 minutes. Impressive.

$hotdogs_i = \beta_0 + \beta_1 height_i + \beta_2 \log(minutes) + \epsilon_i$

$hotdogs = -6.05589 + (.48834)(68) + (1.06702)(4.369) = 31.81304038$

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

$df = n - p - 1$

$58 - 2 - 1$

55

Question 19: You're given that $s\sqrt{x^*(X^TX)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^TX)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$\hat{y}_{\log} = X\hat{\beta} = .204(1.485)$ $\quad \hat{\beta} = (X^TX)^{-1}X^Ty$

$\hat{y}_{\log} = .30294$

$\hat{y}_{\log} \pm (t^*)(SE) \Rightarrow$ e^ each endpoint

10.894148

$.30294 \pm (2.004)(5.43620)$

$= (-7.86 \;\; , e^{.196148})$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

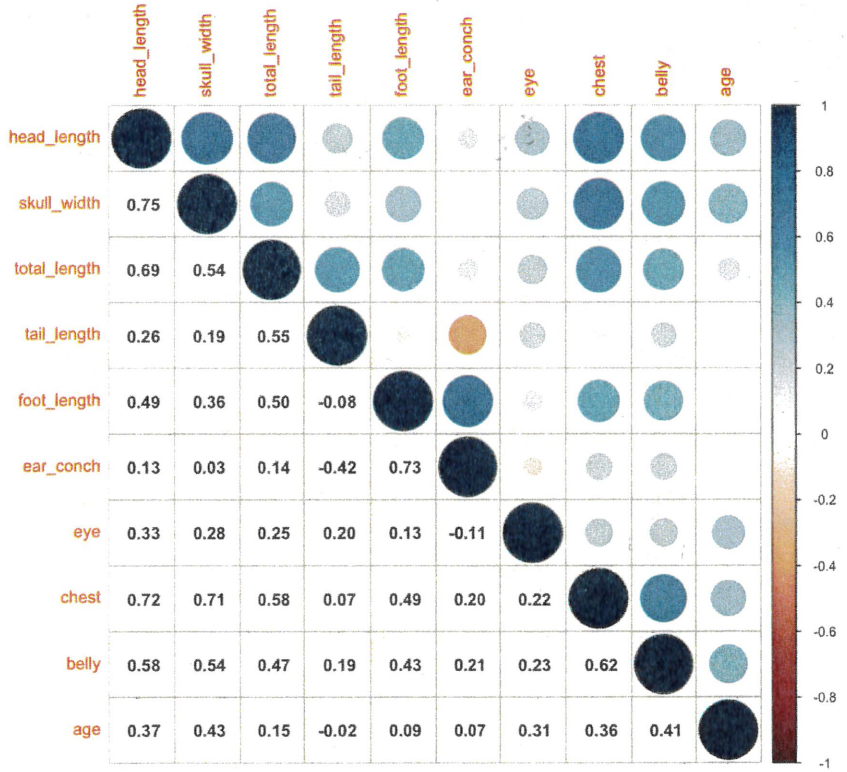

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$\hat{\beta} = (.16559)$

$.16559 \pm (2.01)(1.665)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

.20983

For each +1cm increase in belly girth in cm, holding all other factors constant, we are 95% confident that possom age will have an increase/change within the range (-.023, .443).

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

This "contradiction" is driven by high multi-collinearity, since they (most) all are connected to each other, they don't independently have a significant effect.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The best subsets algorithm using AIC looks at lower values to search through all predictor combinations, this is different from forward selection, where predictors cannot be rotated out to be optimized.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

D Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

D Use the models fitted to the training data to get predictions for the test data

END OF EXAM

Scratch Paper

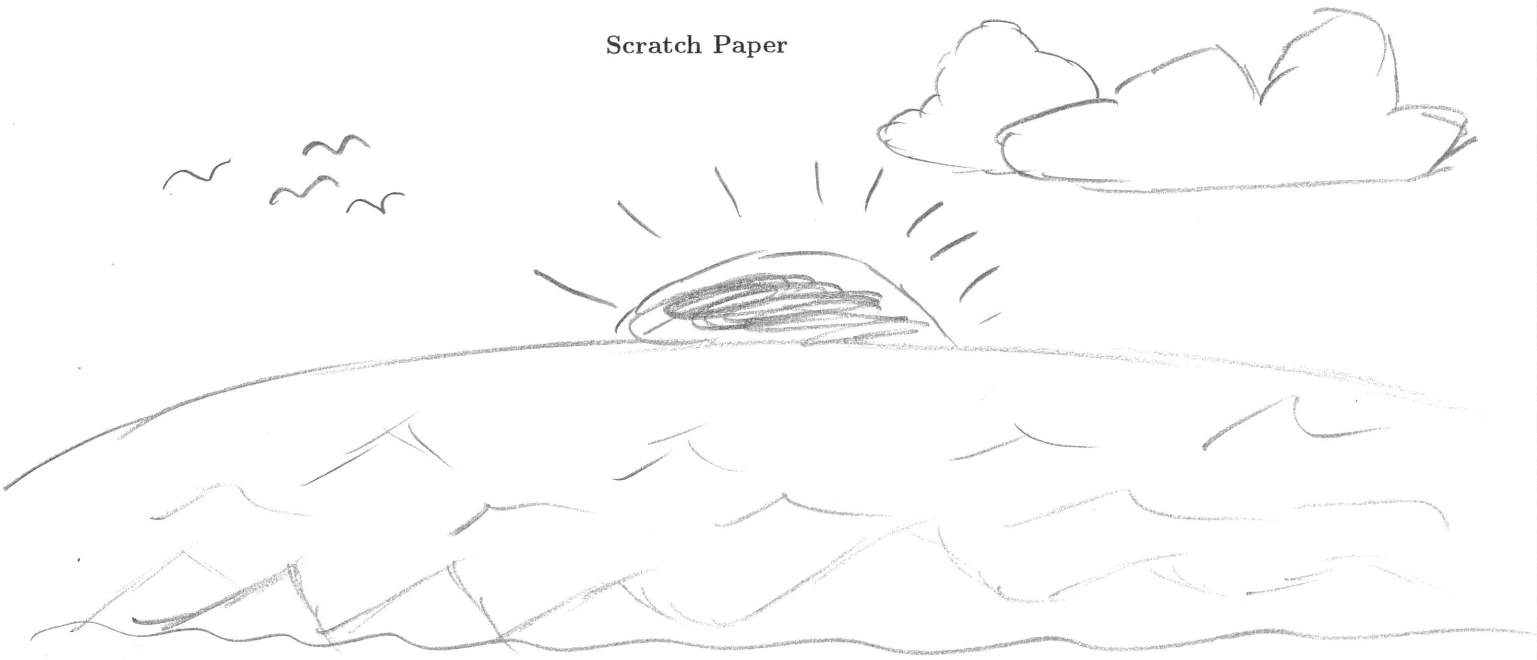

Vibes after completion.

A Gift for attempting to read my handwriting. :-)

Scratch Paper

STAT 330-001

40142034 44

Walker, Andrew

513610279 awalke33 480-309-2855

Test: 640402 Form: A
Entry: 03:13pm 10/21/2026
Employee: Abigail J.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

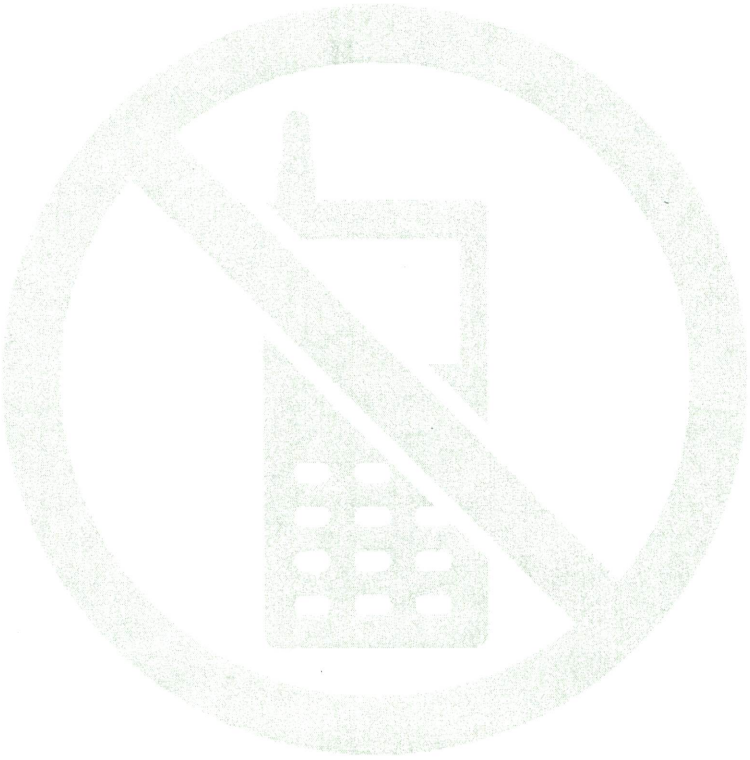

40142034
NS
044

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are reffered as "shrinkage" methods as LASSO and ridge regression shrink the coefficients of the model by either an L1 method or L2 method. L1 allows lasso to shrink some of them to zero while ridge uses L2 to get them very small but keep all that have some influence

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\frac{2.4^2}{4}\left(\frac{0.25}{1-0.25}\right) = \frac{5.76}{4}\left(\frac{1}{3}\right)$ 0.48

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

$\frac{4}{100} = 0.04$

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is the estimated coefficient of a predictor of the sample

β is the fixed coefficient of the population

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least Squares uses estimate regression coefficients to find the best value of $\hat{\beta}$ that minimizes the coefficient.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

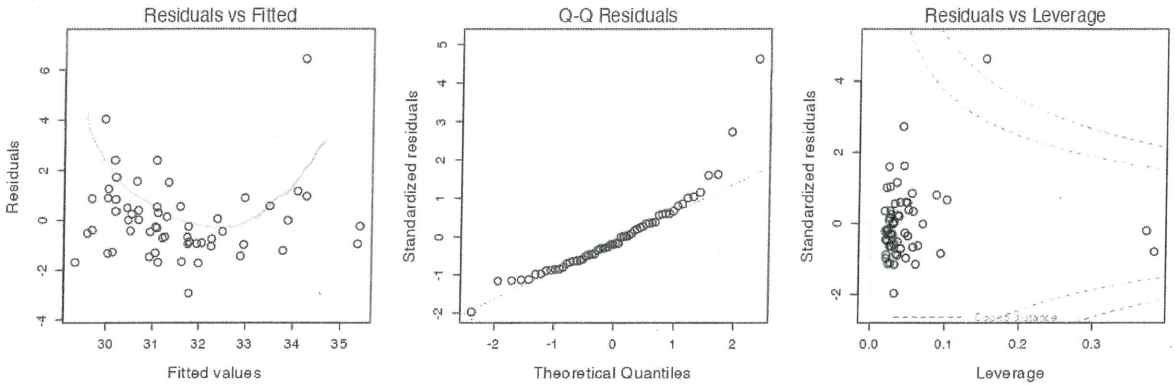

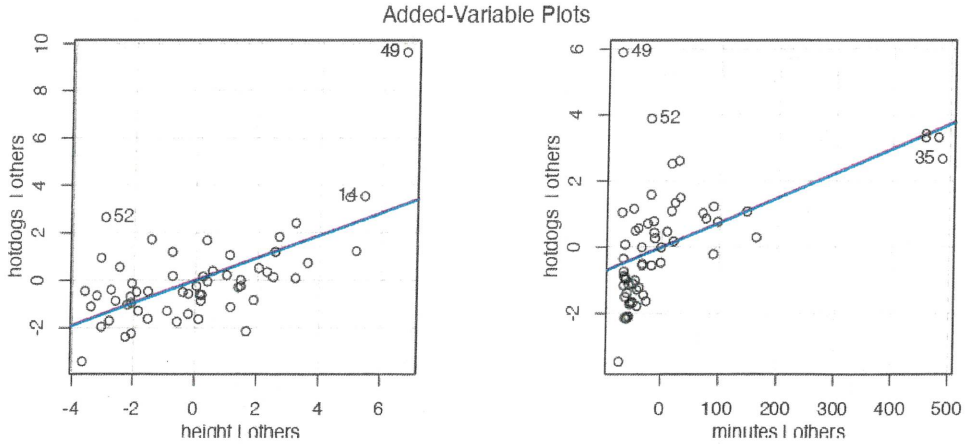

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = 0.472892$

Holding minutes constant, a one inch increase in contestants height is associated with an increase of 0.4729 in number of hot dogs eaten, on average.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

$0.007412 + 3.963(0.001871) = 0.014826$

$0.007412 - 3.963(0.001871) = -0.000002$

$3.963(0.001871) = 0.007414$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

Not quite By looking at the residual vs fitted and Added-variable plots we can see that there is a slight U shape or curve in the residual vs fitted which could violate linearity and the minutes AV plot doesn't show a lot of linearity as well suggesting a violation.

(B) There is independence between each observation

yes There seems to be a general independence between observations but it would be worth looking into the 4 largest counties must were selected from to see if there is any more evidence that it could be independent

(C) The errors are normally distributed NO

NO There is a lot of evidence this distribution is not normal. The shapiro wilk test has a very low p-value which is supported by the deviation off the line in the QQ plot at the tails.

(D) The errors have constant variance (homoscedasticity) yes

yes this is generally supported by the data as the Breusch-Pagan test has a p value of 0.071 which is higher than $\alpha = 0.05$ but it is close enough to still watch out for it and the residual vs fitted plot shows slight funneling

(E) There are no unduly influential points

NO Observation is a very influential point as we see it is an outlier in AV plot for height and also for minutes. We could look into this more by calculating cooks distance for this observation

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

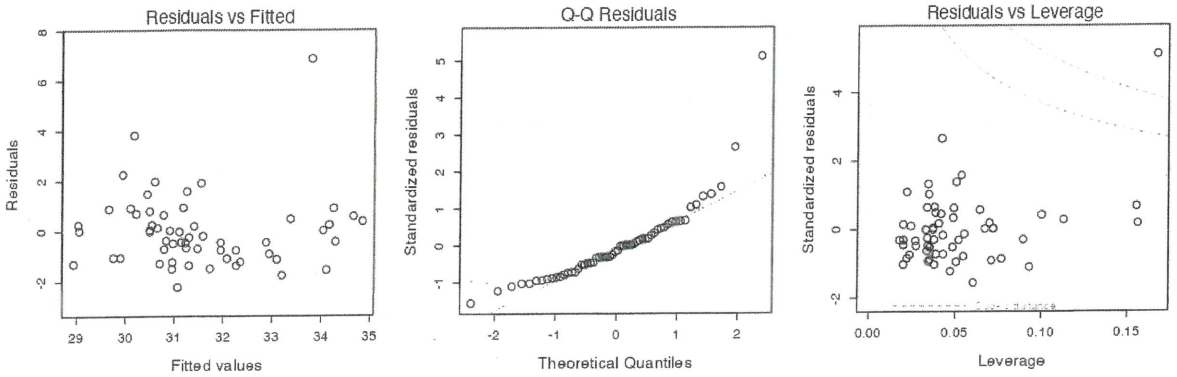

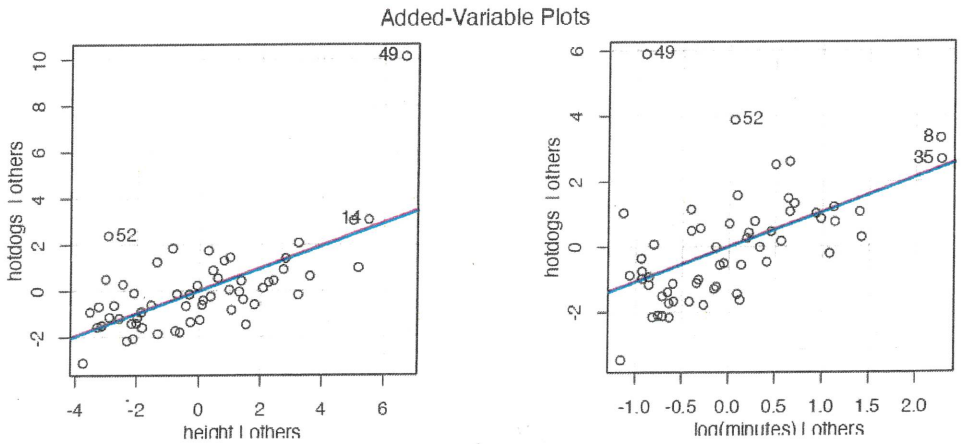

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

The researcher is over reaching in their claim. Even though model 1 has a p value to support the claim, there are other factors to look into when selecting a model like the other LINE assumptions and CVRMSE as well.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages:
Smaller CVRMSE
better linearity in AV plot minutes
Slightly better normality p-value, still low

Disadvantages:
Low BP test p-value so violation of equal variance

Question 16: State the estimated value of β_2 and provide an interpretation in context

$B_2 = 1.06702$

Holding height constant, a 1% increase in minutes trained is associated with a 0.0106702 increase in hot dogs eaten on average.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that $\log(79) = 4.369$. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$$\text{hotdogs} = B_0 + B_1(68) + B_2 \log(79)$$

$$\text{hotdogs} = -6.48834 + 0.48834(68) + 1.06702(4.369)$$

$$\text{hotdogs} = -6.48834 + 33.20712 + 4.66181038$$

$$\boxed{\text{hotdogs} = 31.38}$$

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error? $n-p-1$

$58-p-1$

$58-2-1 = 58-3 = 55$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$PI = \hat{Y} \pm t^* s\sqrt{1+h_0}$

34.36 ↑

upper $PI = 31.38 + 2.004(1.485) = 31.38 + 2.97594$

lower $PI = 31.38 - 2.004(1.485) = 31.38 - 2.97594$

↓ 28.40

$\boxed{PI(28.40, 34.36)}$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

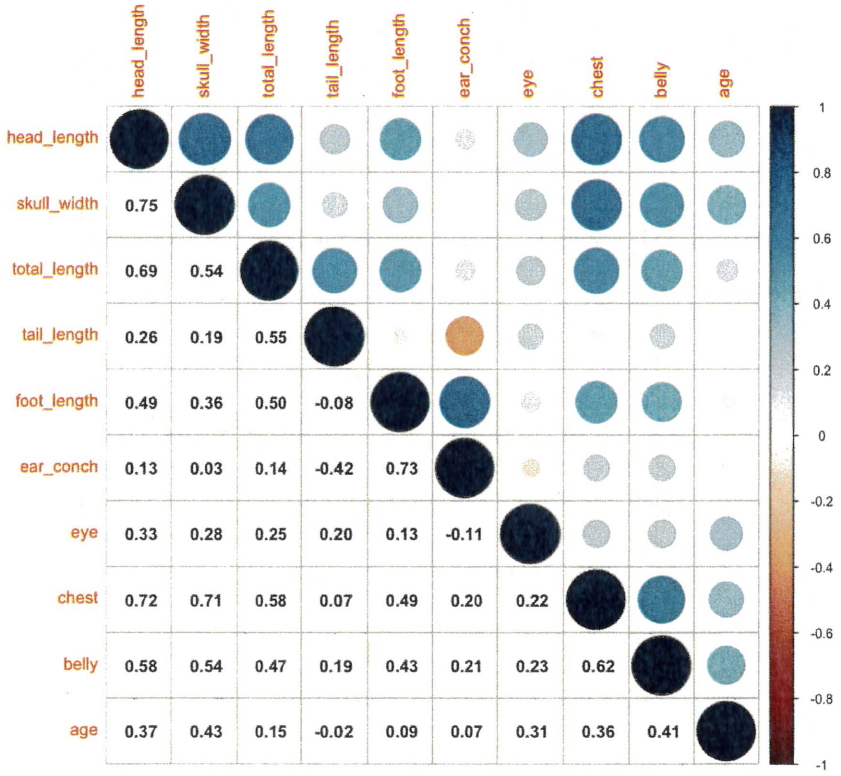

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 2.01(0.10599)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

Holding all other predictors constant, we are 95% confident that a 1 cm increase in belly girth is associated with a change in age between a decrease by 0.023 to an increase of 0.443.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

What is driving this contradiction could be the multicollinearity between head length and skull width, also tail length and total length. Individually they don't show significant effect but together they do.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best Subsets is an exhaustive algorithm that looks at all possible models and by using AIC as its metric, it will select the model and its variables with the lowest AIC.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

metrics penalize differently so they select different models

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

C Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 43

Long, Gracie Ann

525555856 glong25 9013262021

Test: 640355 Form: A Wadsworth, Spencer Gordon

Entry: 03:04pm 10/21/2026 Checked out:

Employee: Christian S.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

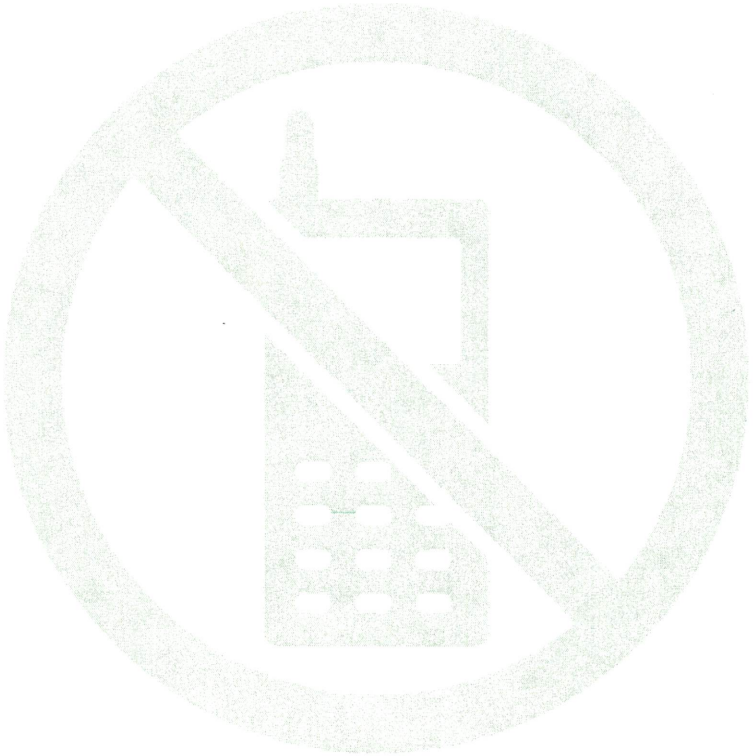

.007417 ± 1.995(.001871)

(.00368, .0111)

NS

40142034

043

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

Lasso and ridge regression use λ to shrink the coefficients. Lasso can set coefficients to zero using absolute value and ridge minimizes the coefficients using squared values but never sets them equal to zero.

$\frac{4}{100} = .04$

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\frac{2.4^2}{3+1}\left(\frac{.25}{1-.25}\right) = .48$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^TX)^{-1}X^TY$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is an estimate and can vary from sampling variability. β is a fitted value that stays the same throughout sampling.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least squares is used to estimate regression coefficients by minimizing SSE from finding a $\hat{\beta}$ that makes SSE as small as possible.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

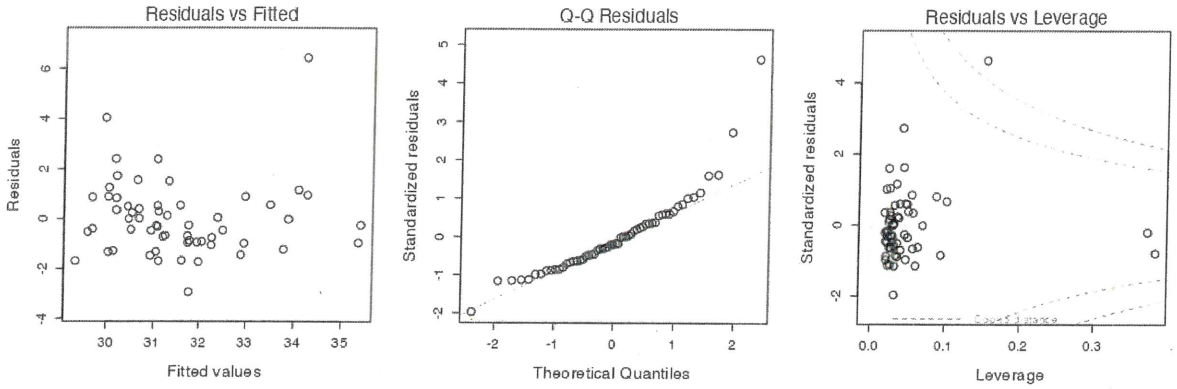

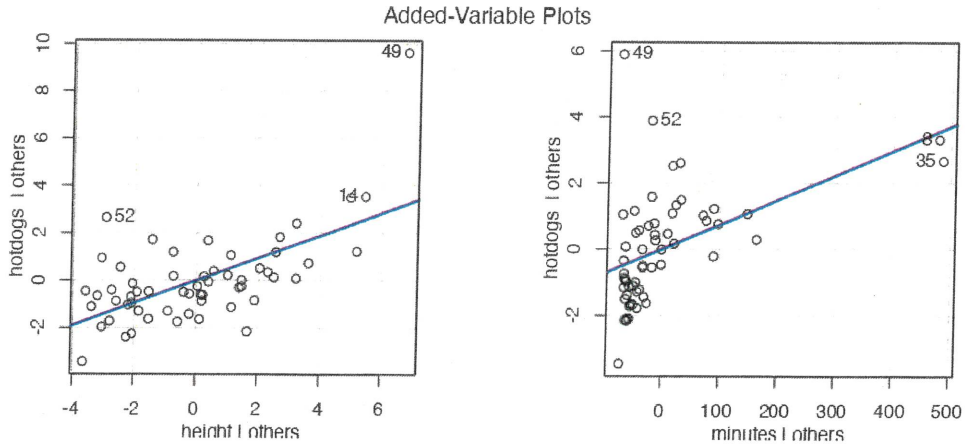

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\beta_1 = 0.473$

Holding all other predictors constant, a one inch increase in height is associated with an expected change of 0.473 hotdogs consumed.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The residual vs. fitted plot shows mostly centered around zero but a possible funnel shape is present, so there is possible violation of linearity.

(B) There is independence between each observation

Since the data comes from the 4 largest counties in the state, that does not appear to be random and so independence could be violated.

(C) The errors are normally distributed

With a tiny Shapiro-wilk value (< .05) and a QQ plot that follows the line pretty well, the errors appear to be normally distributed.

(D) The errors have constant variance (homoscedasticity)

With Breusch Pagan value larger than .05 and the scale-location plot, there appears to be unequal variance.

(E) There are no unduly influential points

With a Cook's distance that exceeds 0.069, there is a possibility for high influence observations.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

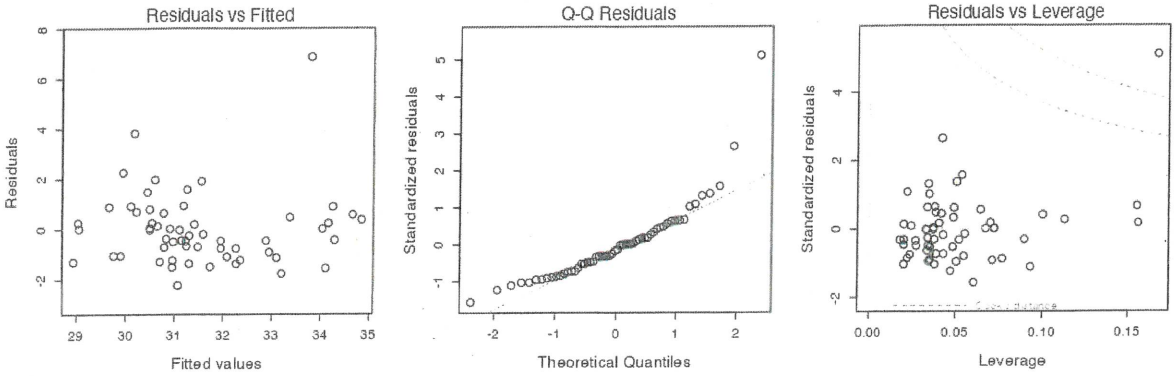

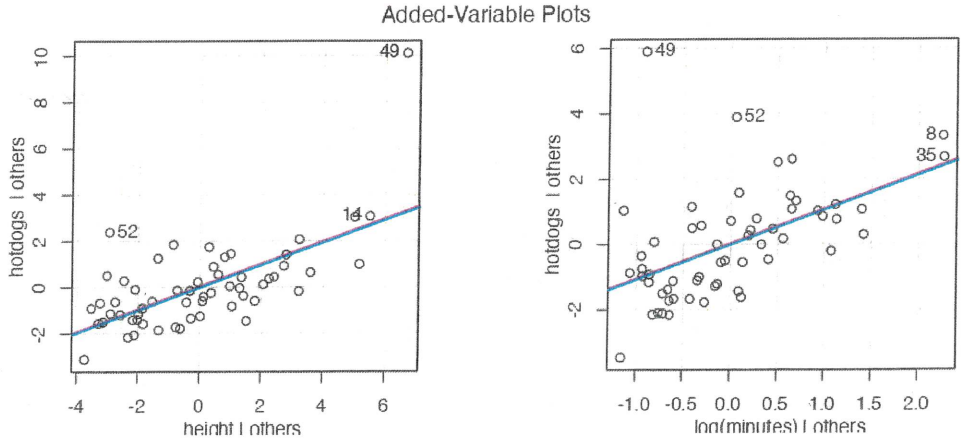

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

Since model 2 now has a Breusch Pagan value less than .05, this suggests that there now is equal variance. This statement makes model 2 more aligned with the LINE assumptions and therefore fits the data better after the transformation.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Linearity and equal variance are more strongly supported in model 2, which was less supported in model 1. Using a transformation on x can make interpretations more difficult along with the possibility of calculation error.

Question 16: State the estimated value of β_2 and provide an interpretation in context

$B_2 = 1.067$

Holding all other predictors constant, a 1% increase in minutes is associated with an expected change of 0.01067 hotdogs consumed.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

hotdogs = $-6.056 + 0.488(68) + 1.067\log(79)$

$= 31.79$ hotdogs consumed

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

$n-p-1 = 58-2-1 = 55$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$10 \pm 2.004 \times 1.485 = (7.024, 12.9759)$

We are 95% confident that the actual number of hotdogs consumed in 10 minutes by my best friend will fall between 7.024 and 12.9759 hotdogs consumed.

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

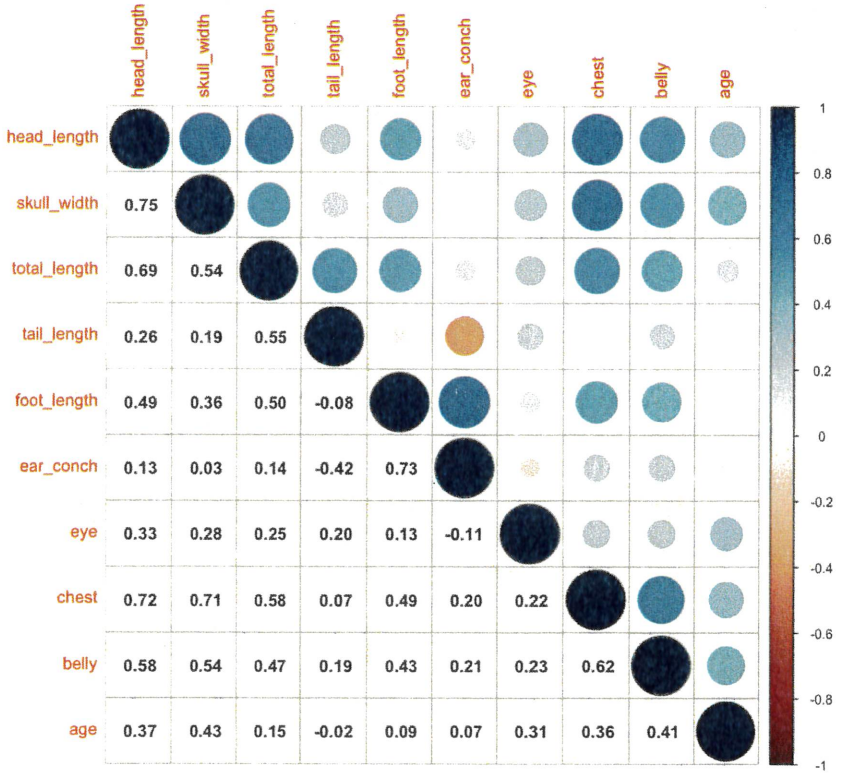

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 2.01 \times 0.10599$

$0.16559 \pm 2.01(0.10599)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that the true median for the size of belly width is between -0.023 and 0.443 centimeters.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

Each possum grows at a different rate and some may be larger at a younger age and larger than a senior possum.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The best subsets algorithm is a method used to explore combinations to decide which models to fit. AIC is a metric criterion that is used to evaluate models against each other and figure out how good the model is.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

C Use the models fitted to the training data to get predictions for the test data

END OF EXAM