STAT 330-001

40142034 42

Givens, Andrew William

659505364 andrewwg 804-774-1328

Test: 640264 Form: A Wadsworth, Spencer Gordon

Entry: 02:34pm 10/21/2026 Checked out:

Employee: Ty B.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

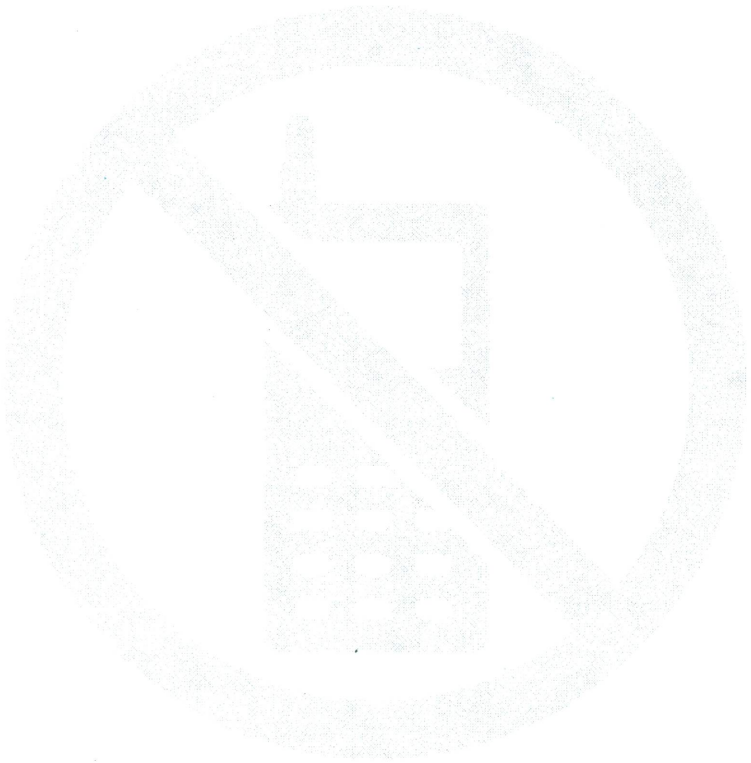

40142034

042

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They shrink the coefficients by applying a scalar λ. The bigger λ is, the greater the shrink. Ridge uses L2 which involves squaring whereas Lasso uses L1.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\frac{5.76}{4} \cdot \left(\frac{.25}{.75}\right) = .48 > .04$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$(X^TX)^{-1}X^TY$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β exists but is unknown to us; it's part of the theoretical model.
$\widehat{\beta}$ is our best estimate.

?

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Your model optimizes coefficients by targeting the minimum least squares.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

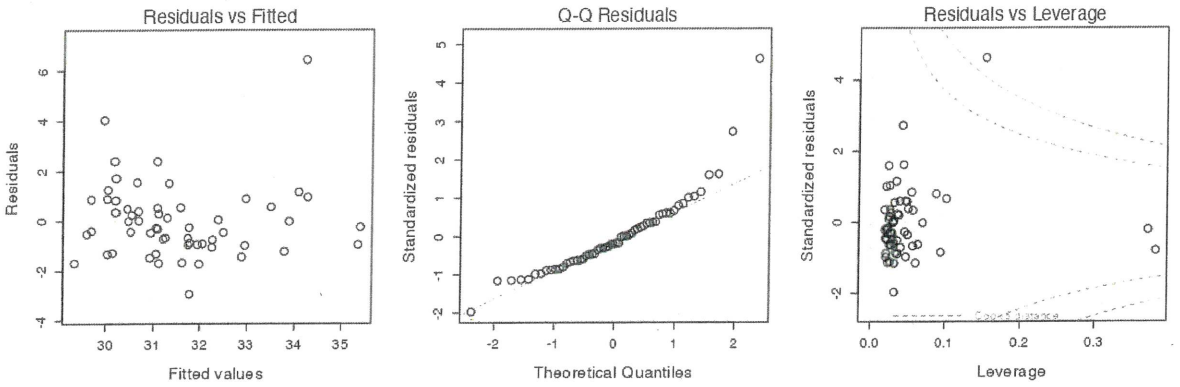

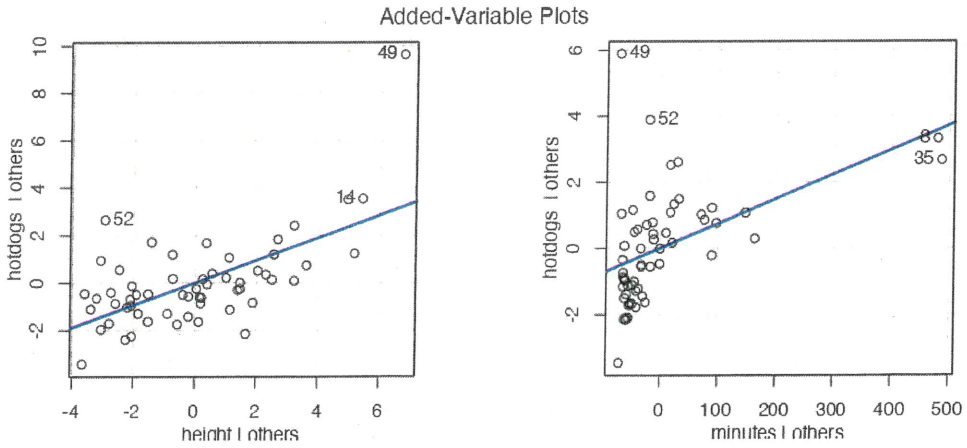

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

0.472892. holding minutes trained each week constant, a one inch increase in height is associated with an average increase of .472892 hot dogs eaten

Question 12: **(Circle all that apply)** Which of the following are true?

$CI = \hat{\beta} \pm t^* SE$

$.472892 \pm 5.885(0.080352)$

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

→ depends on how many decimal places...

0.00002 , 0.941576

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

looking at residuals vs fitted, it would seem there's an issue with linearity represented by the curvature present.

(B) There is independence between each observation

There could be an issue with independence. The data doesn't seem to have been collected randomly or from a wide sample, but most of the observations are from the same 4 largest counties in one state. Unknown if they're all from same state.

(C) The errors are normally distributed

The shapiro-wilk p value, 2.43e-8 suggests normality is violated. Looking at the Q-Q plot would suggest the departure from normal is mild w/ stronger issues on both tails (the right more so). may be more a right skew.

(D) The errors have constant variance (homoscedasticity)

Breusch-pagan p-value, 0.071 > 0.05, but only just w/ no scale-location plot, there's not enough evidence to say there's an issue w/ equal variance.

(E) There are no unduly influential points

Looking at residuals vs leverage, there do seem to be influential points of note. looking at the AV plots, observation 49 stands out.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

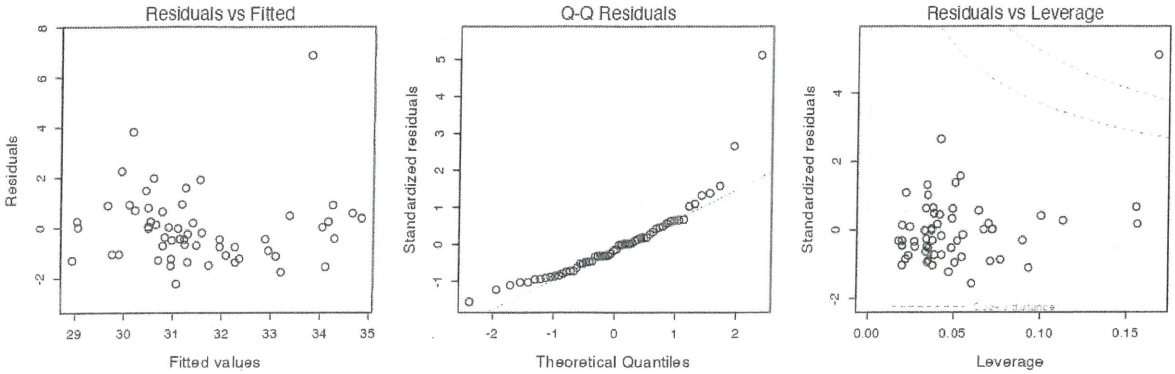

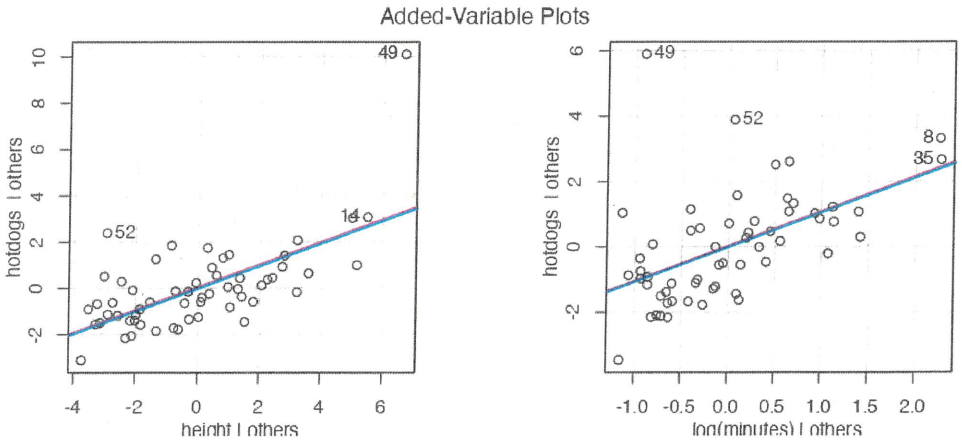

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

Clearly is a stretch. a change from .071 to 0.028 is statistically significant, however it isn't far below 0.05, whereas the p value for shapiro wilk improved by 3 decimal places. You also need to go through LINE in its entirety when a transformation takes place.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

To me the difference seems marginal. Equal variance decreased slightly, but normality improved by a larger factor. This is reflected in the Q-Q plot and SW p value. In the residuals vs leverage one point moves much closer to the curved lines.

Question 16: State the estimated value of β_2 and provide an interpretation in context

1.06702. Holding height constant, a 1% increase in minutes per week spent training is associated with an average increase of 0.01067 hot dogs eaten.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

Hotdogs = .48834(68) + 1.06702(4.369)

37.87 hot dogs

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55 df = n − p − 1

(B) 57 58 − 2 − 1

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

.204 ± 2.004(1.485)

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

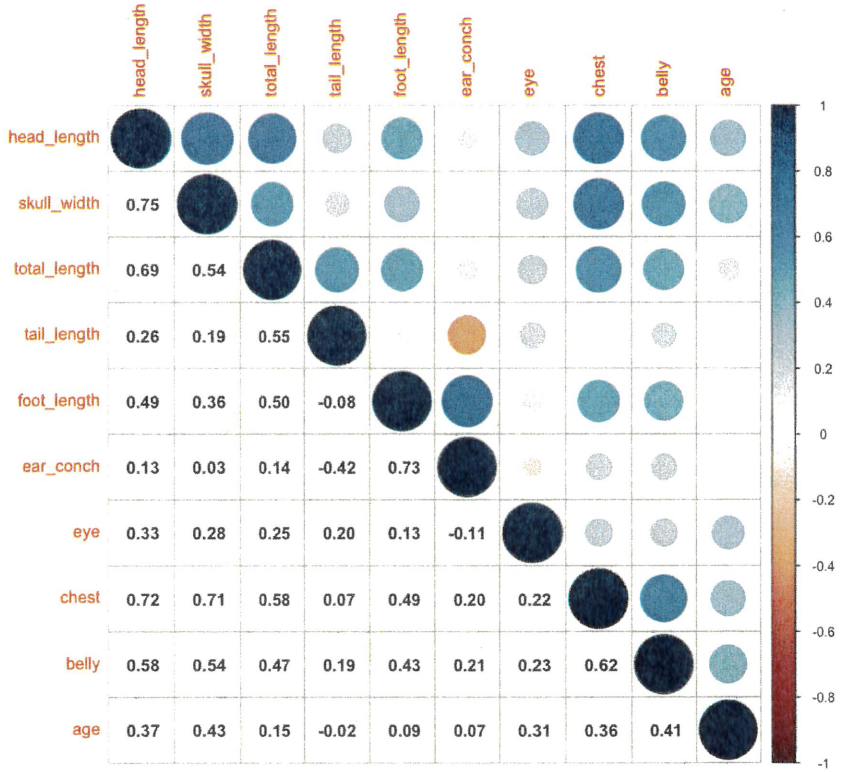

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

.16559 ± $t^* = 2.01$

$SE = .10599$

$.16559 \pm 2.01(.10599)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

It contains 0, which I think is a good thing. We're 95% confident that the true Beta for Belly is w/in 1 standard error of our estimate? The true beta does not move

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

Our F-test checks the combined effect of our predictors and if they're statistically significant. the t-test checks them individually w/o taking into account the other predictors

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

AIC is the metric used to grade the models produced by the algorithm best subsets. Adjusted R^2 applies smallest penalty for multiple predictors, then AIC, and then BIC and the harshest penalties applied. Whichever model has the lowest AIC score will be selected.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 41

Waites, Drake Webster

156955952 dwwaites 509-760-5955

Test: 640258 Form: A Wadsworth, Spencer Gordon

Entry: 02:33pm 10/21/2026 Checked out:

Employee: Lillie A.

SPECIAL INSTRUCTIONS
TC Calc Notes: 3x5 Dictionary Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034
041

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

Variance is being shrunk.
OLS chooses coefficients to minimize the sum of the squared residuals.
LASSO or Ridge regression also do this, but they add an extra penalty, either L1 or L2, to the coefficient selection. This penalty "shrinks" the coefficients, intentionally biasing them to reduce variance and can improve overall prediction.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is the vector of estimated coefficients

β is the vector of unknown true population parameters

$E(\hat{\beta}) = \beta$

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least squares estimates regression coefficients by minimizing the sum of the squared residuals.

Least squares estimates a line through the data, regression coefficients are chosen to minimize $y - \hat{y}$, the observed response values minus the estimated model values.

Each coefficient is assigned using a partial derivative interpretation, holding all else constant there is an estimated change in the response associated with each coefficient.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

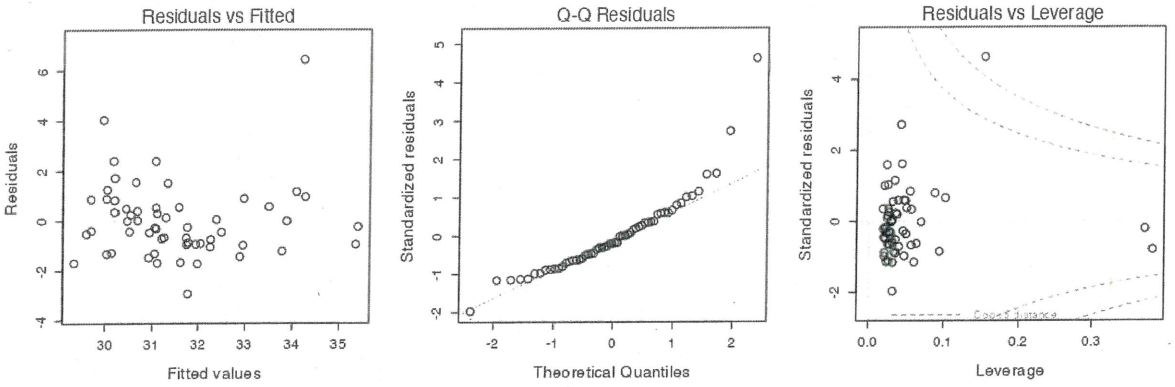

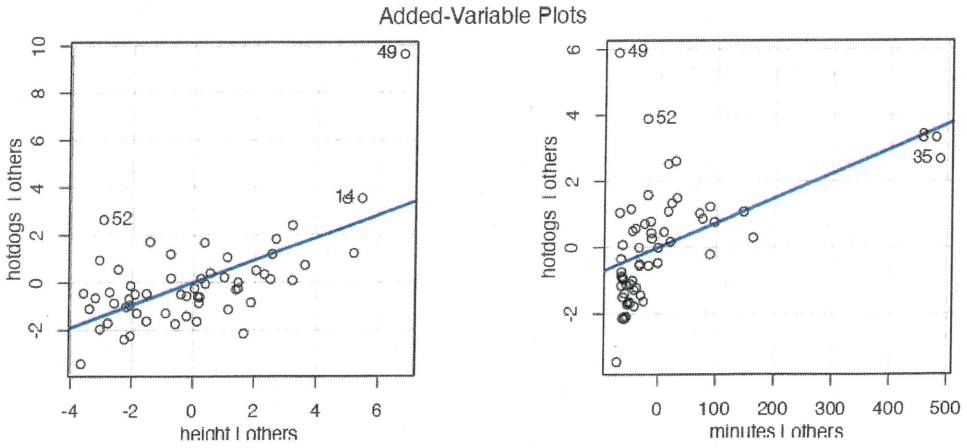

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = 0.472892$

On average, there is an associated .47 hotdogs increase in number of hotdogs eaten with a 1 inch increase in height, holding all else constant.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

Height and hot dogs eaten look to have a linear relationship. Minutes training and hot dogs eaten have only a roughly linear relationship; I think it's fair to say the relationship isn't very linear at all.

(B) There is independence between each observation

Independence seems a likely assumption, given how the data are collected. Because most of the observations are people from the same county, there could be some relationship there, but independence is a fair assumption to make here.

(C) The errors are normally distributed

Almost all of the data follows a normal distribution, however there is some deviation in the upper tail. The Shapiro Wilk test confirm this, based on the results of the test we would reject the null hypothesis that the errors follow a normal distribution. I would investigate those points and how the data was collected, and assume normality might be violated.

(D) The errors have constant variance (homoscedasticity)

The Breusch Pagan test has a p-value of .011, meaning we would fail to reject this test at the 5% and the 1% levels. Based on the residuals vs fitted plot and the test, I think this assumption is reasonably acceptable. The residuals vs fitted plot isn't perfect, but most of the data is fitted around 0 and appears to be randomly scattered, although there is some outliers.

(E) There are no unduly influential points

There are definitely some points with high influence. 1 point in particular sits outside the Cooks distance lines at the top of the graph. 2 more sit far to the right on the graph. I would investigate all 3 points.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

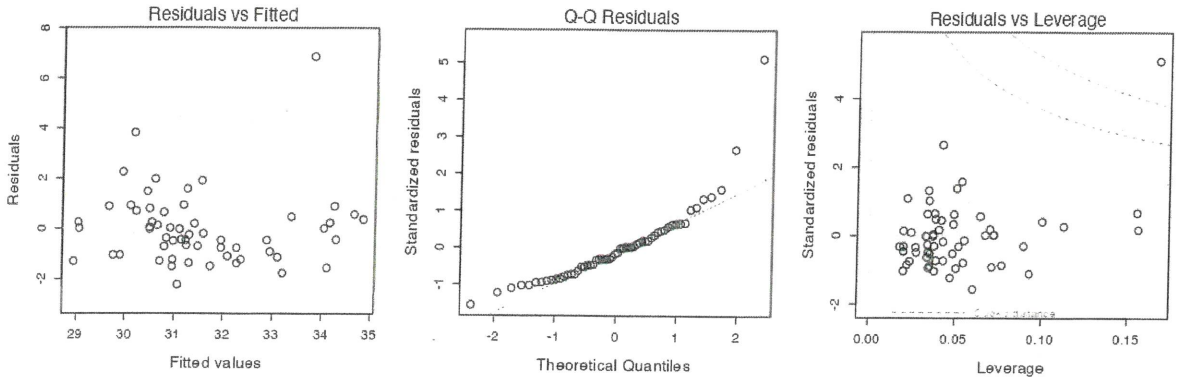

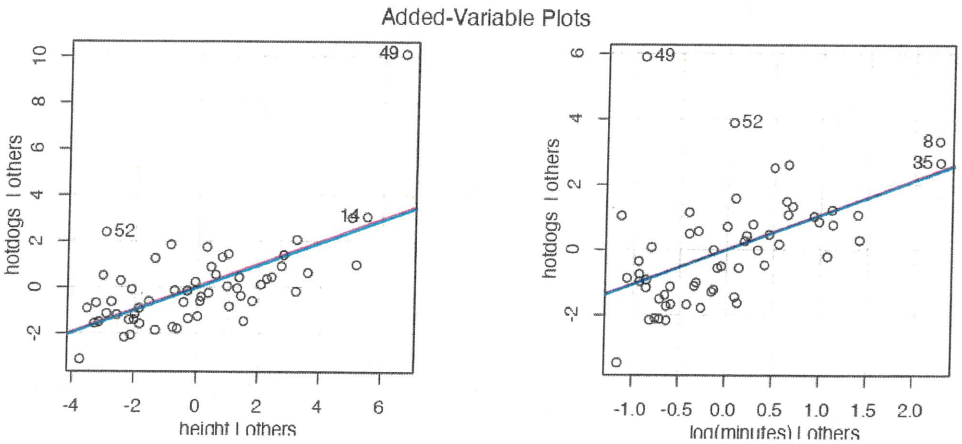

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

This is not necessarily true. It depends on what your goals are for the model. Model 1 may have lower SE and so will be better for inference

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages:
- Less influential points
- Better linear relationship between minutes and hot dogs eaten
- Lower CV RMSE; better out of sample prediction

Disadvantages
- In sample predictions by not be as good
- Inference may not be as stable, higher SE
- More difficult to interpret

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$

For a 1% increase in minutes training, there is an associated increase of 1.06702 hot dogs eaten, on average, holding all else constant.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$\hat{y}_{find} = -6.05589 + (.48834)(68) + (4.369)(1.06702) = 31.81304038$

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$PI_{95\%} = 31.81304038 \pm 2.004(1.1485) = [28.83710038, 34.78898038]$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

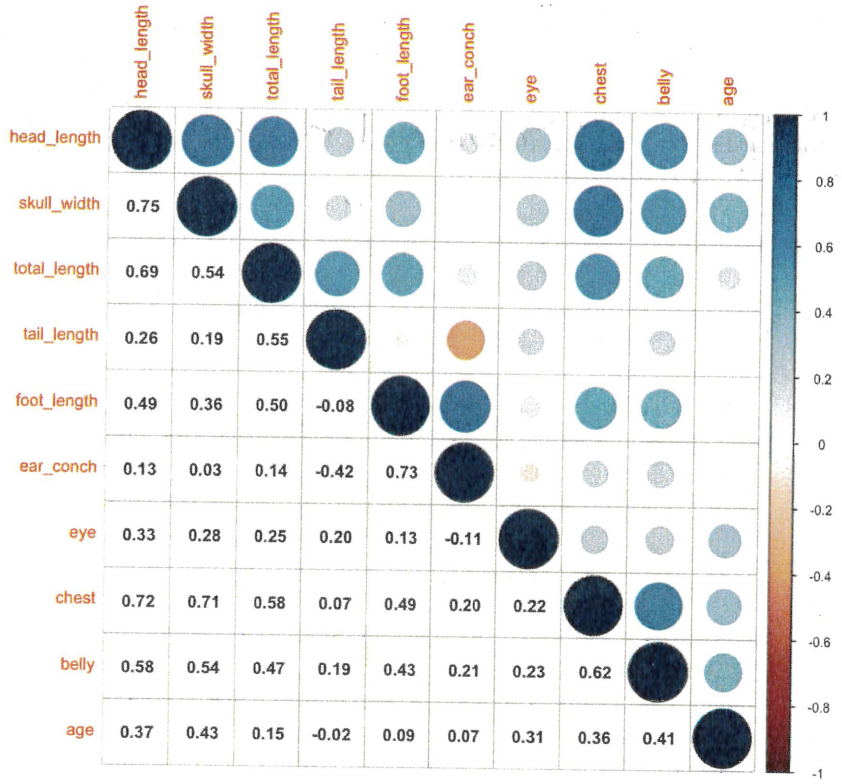

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

95% CI = 0.1659 ± 2.01√

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

Using repeated sampling, if we estimated the coefficient for belly width repeatedly, 95% of the estimates would fall into the interval (-0.023, 0.443).

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

The overall F-test is testing that all of the coefficients are equal to zero. The associated p-value is .0065, which means that we reject the null hypothesis that all the coefficients are zero.

The individual t-tests are looking at each coefficient individually, and none are significant at the 5% level, meaning none of the coefficients have a significant result that they are equal to zero, we fail to reject the hypothesis.

The tests are different: one is an overall test that says based on the observed data, it's unlikely all the coefficients equal zero, but it doesn't tell us which.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets tests every model: it takes every combination of predictors and computes its AIC. AIC is a selection metric, $AIC = n \log(SSE/n) + 2k$. Best subsets fits every model, then calculates the AIC for each one, and returns the model with the lowest AIC.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

B Average RMSE across all test-training splits and report

B Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

B Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 40

Batsaikhan, Nandintsetseg

401393027 bnandin1 7206629825

Test: 640237 Form: A Wadsworth, Spencer Gordon

Entry: 02:26pm 10/21/2026 Checked out:

Employee: Ty B.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

040

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

Lasso and ridge regression shrink correlation coefficient while some bias introduced but lowering variance

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$D_i = 0.48$ $\quad \frac{2.4^2}{4}\left(\frac{0.25}{0.75}\right) = 1.44(0.333) = 0.48$. $\quad \frac{4}{100} = 0.04$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is uknown parameter and it is population parameter.

$\hat{\beta}$ is observed value of our sample parameter which we can measure.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

$$SSE = (Y - Xb)^T (Y - Xb)$$

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

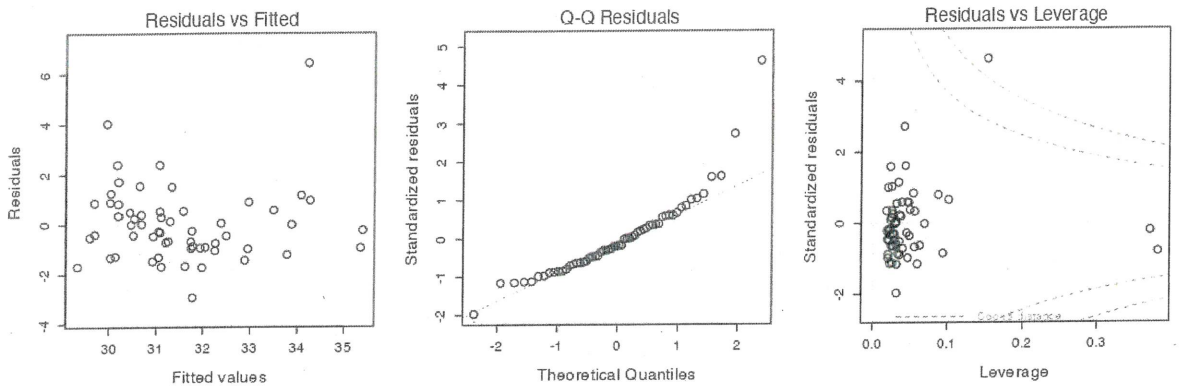

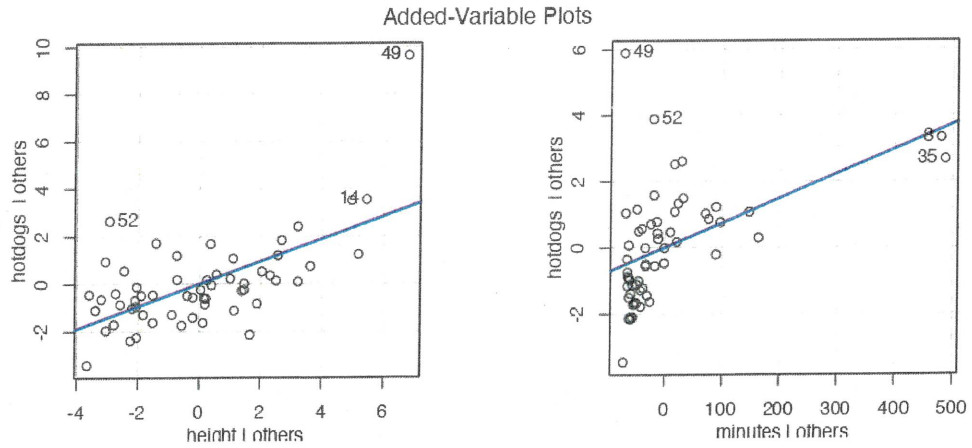

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

Holding minutes predictor constant, average hotdogs eating time can increase per hotdogs height.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

Given Residual vs Fitted graph, relationship between explanatory and response are ~~satisfied~~ linear assumption satisfied.

(B) There is independence between each observation

Yes sampling disign, competitive eater are individual people. So independse assumption satisfies.

(C) The errors are normally distributed

QQ plot of normality looks spreading in the lower and higher tail. Shapiro Wilck value is very low. further investigation on normality needed.

(D) The errors have constant variance (homoscedasticity)

BreuschPagan test has 0.071 > 0.05 value, So homocedacity ~~are~~ satisfies.

(E) There are no unduly influential points

No.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

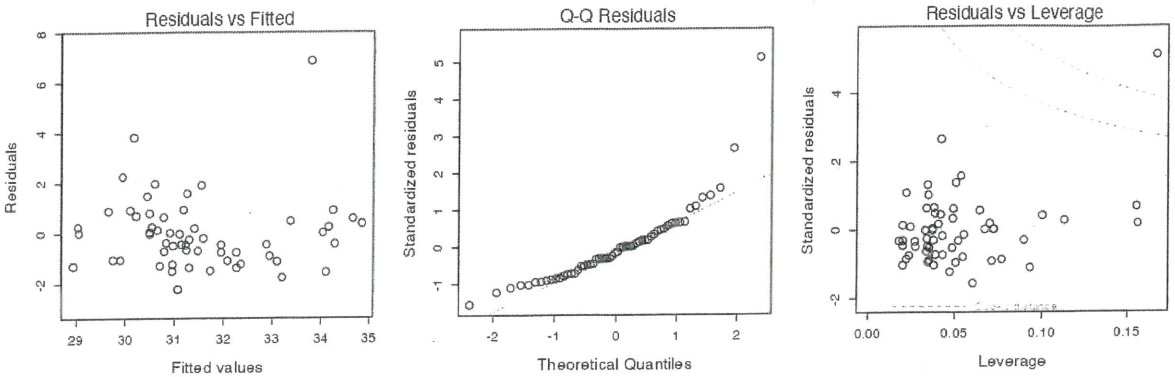

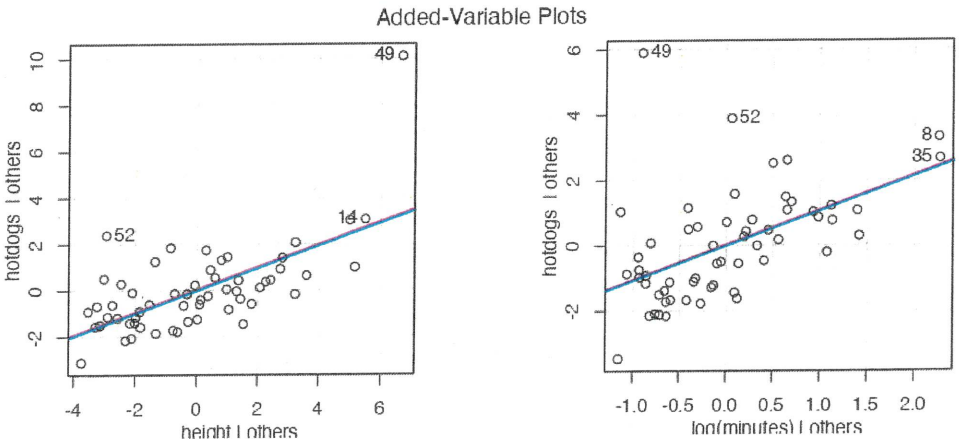

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

log transformed model 2 is more useful to proceed in this case.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

advantage → Can fit assumptions better.

disadadvantage → Can increase bias.

Question 16: State the estimated value of β_2 and provide an interpretation in context

B_2 is $e^{1.06702}$. Holding height in constant, average hotdog eating would increase by 1.06702 per minute.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 + 68(0.48834) + 4.369(1.06702) =$

$-6.05589 + 33.20 + 4.662 = 31.80$

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

$n-p-1 = 58-2-1 = 55$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

~~$31.80 \pm t^*(se) = 31.80 \pm 0.204$~~

$31.80 \pm 2.004 =$

~~$[33.804,$~~ $[29.796, 33.804]$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

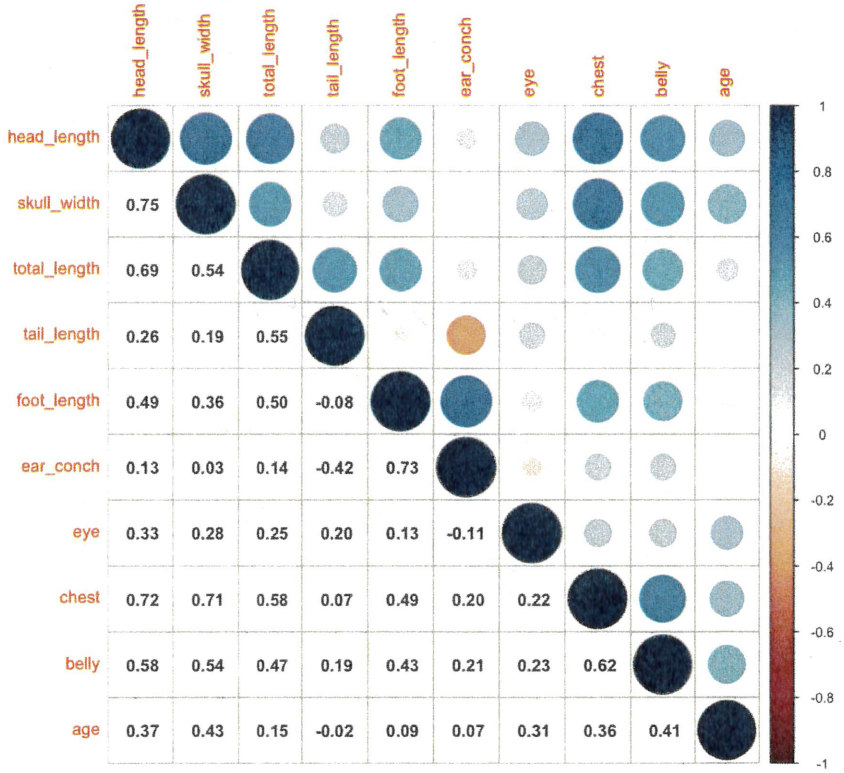

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.1659 \pm 2.01(SE)$

$0.1659 \pm 2.01(SE)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

Holding other variables constant, there is 95% probability that β_{belly} is in the interval $(-0.023, 0.443)$.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

F-test pvalue $0.0065 < 0.05$. So there is insufficient evidence that all predictors are linearly associated.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets algorithm selects all variables and chooses the best predictors.
Best subsets is algorithm method while AIC is metric score.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

B Average the predictions and record their value

D Choose the model with the largest RMSE

D Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

C Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

B Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 39

Zinn, Mason Robert

236667935 mrz24 408-722-8163

Test: 640186 Form: A Wadsworth, Spencer Gordon

Entry: 02:10pm 10/21/2026 Checked out:

Employee: Libby R.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034
039

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets.** Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

LASSO and Ridge have penalties which penalize the number of predictors used in the model. They shrink the $\hat{\beta}$'s for each predictor due to the penalties. Other than shrinking the $\hat{\beta}$'s toward zero, Ridge can reduce some to exactly zero, removing predictors from the model.

$(1.44)(.333) = 0.48$

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

"β" is the true population statistic while "$\hat{\beta}$" is our estimate of the true "β" using the data and a linear model to calculate it.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

If you fail "$SSE = (y - X\beta)^T(y - X\beta)$" and reduce you eventually get the normal equations and $\hat{\beta} = (X^T X)^{-1} X^T y$. Because of that minimizing our least squares gives us the best $\hat{\beta}$'s for our model.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

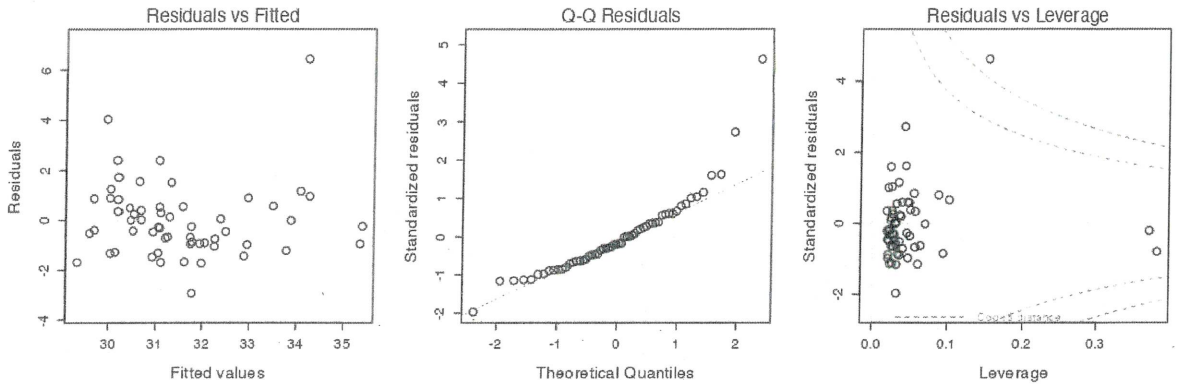

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = 0.472842$

holding all other predictors constant, a 1 inch increase in height is associated with an average increase of 0.47 hot dogs consumed.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0 (0.0031, 0.0109)

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

We get a random pattern from residual vs fitted which is good. For the added variable plots height vs hotdogs is great while minutes vs. hotdogs isn't super linear. Transformation might be able to fix that by reducing outlier weights.

(B) There is independence between each observation

I think its safe to assume independence because the nature of the contest is very individualized so you can't impact the other contestants.

(C) The errors are normally distributed

From Q-Q It appears to be fairly normal with some right skew. looking at Shapiro wilk test that is confirmed by small p-value. It would be good to see if a transformation would help resolve that.

(D) The errors have constant variance (homoscedasticity)

From residual vs. fitted there is a slight slight funnel shape and from BP test we get a non significant result. When looking at residual vs. leverage there appears to be an overly influencial point which should be investigated.

(E) There are no unduly influential points

As mentioned above from residual vs leverage there appears to be an overly influencial point which should be investigated. point number 49 most likely

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

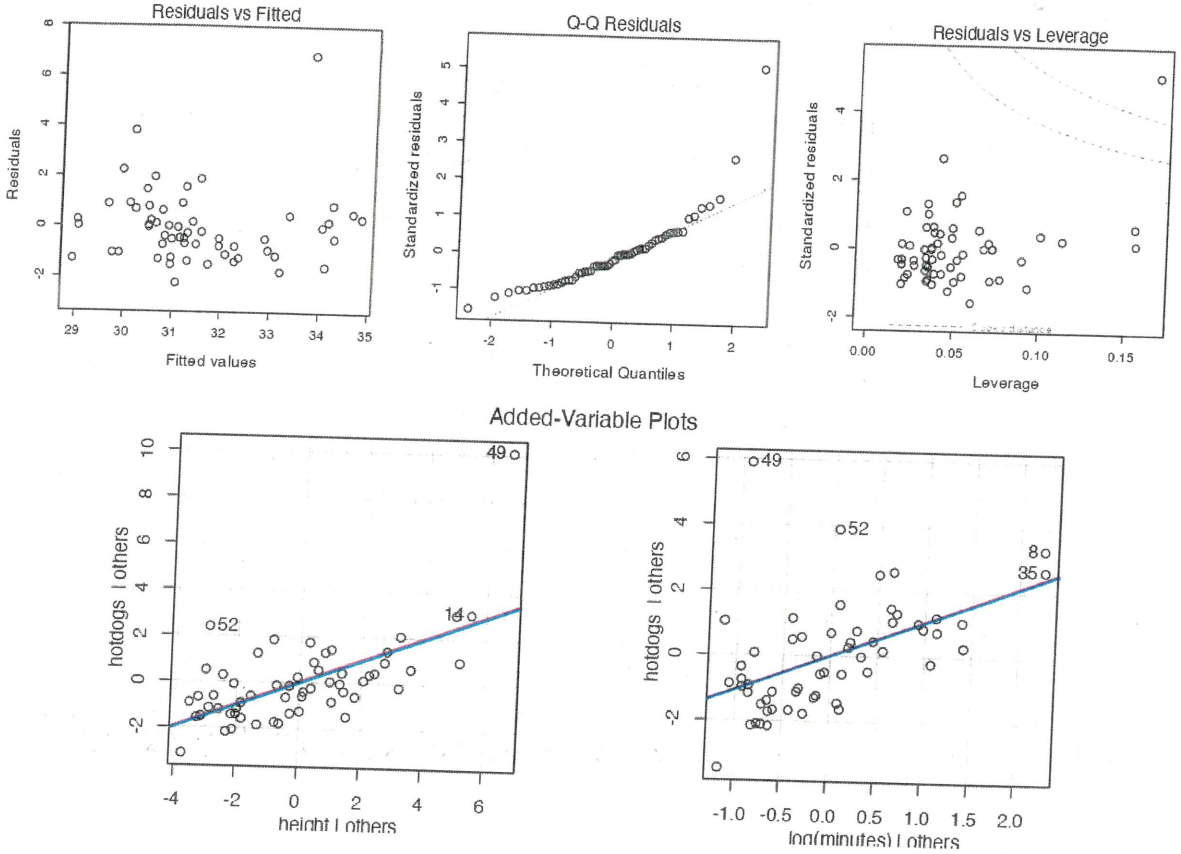

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

The difference in p-values is only 0.05 which is so small so that should not be the only reason you choose model 2. You need to look at the residual vs. fitted graphs as well before choosing. As for being superior, model (1) may be better for confidence/prediction intervals, t and F tests because it has better equal variance, but model (2) will have a model that makes more sense and has better coefficient estimates because it is more linear.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

The advantages are that it is a lot more linear as we can see from added variable plot minutes vs. hotdogs. Another advantage is that it is more normal. It still has a small p-value for shapiro wilks but it is x1000 times bigger than model 2's Shapiro wilk test.

Some disadvantages are that it made the errors very more. We get a lower BP p-value and the funnel is more exaggerated in resid vs. fitted plot. Also from residual vs. leverage it looks like the leverage of the outlier increased as it is even further behind cook's distance lines.

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$ $\quad C = 5$ $\quad \ln(1+\frac{5}{100}) = 0.048 \cdot 1.06702 = 0.052$

Holding all other predictors constant a 5% increase in minutes training for hotdog contest, is associated with an average increase of 0.052 hotdogs consumed

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that $\log(79) = 4.369$. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

hotdogs $= -6.05589 + 0.118(\text{height}) + 1.067(\text{mins})$

$= -6.05589 + 32.64 + 4.66$

$= 31.24$

We would predict they eat 31.24 hotdogs in 10 minutes

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

$n - p - 1$

$58 - 2 - 1 = 55$

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31.24 \pm 1.485(2.004) = (28.26, 34.21)$

95% prediction interval $= (28.26, 34.21)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

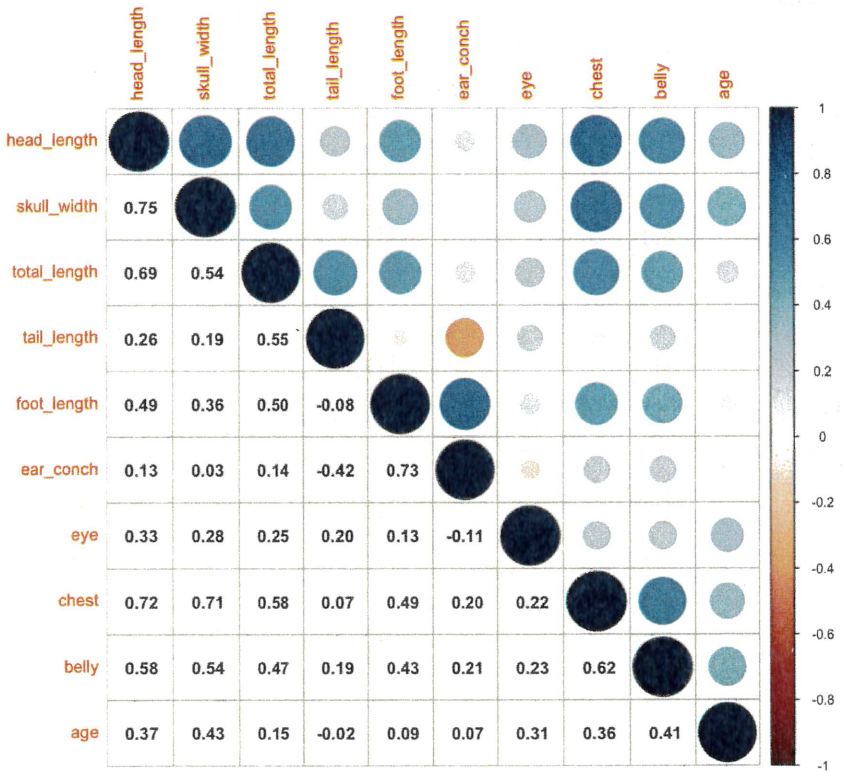

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 2.01(SE\sqrt{x^*(X^TX)^{-1}x^*})$

SE? SE=0.01554

$0.16559 \pm 2.01 * 0.01554$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

holding all other predictors constant, we are 95% confident that a 1cm increase in belly girth is associated with an average increase from -0.023 to 0.443 years of age for a possum.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

the overall test is comparing each one to the other so you are more likely to see that at least one $\hat{\beta}_j$ is different while the individual test compares the individual $\hat{\beta}_j$ to the model one at a time not capturing all $\hat{\beta}_j$'s at once.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

AIC penalizes the number of predictors with "2k" which helps simple models score higher. It also chooses models that are better suited for prediction rather than inference.

$AIC = -2\ln(\hat{L}) + 2k$

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

L Average RMSE across all test-training splits and report

D Average the predictions and record their value

C Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 38

Memmott, Tanner Johnson

881026140 tmemmot5 712-266-5510

Test: 640181 Form: A Wadsworth, Spencer Gordon

Entry: 02:09pm 10/21/2026 Checked out:

Employee: Lillie A.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

Tanner Memmott NS

40142034

038

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets.** Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are called shrinkage methods because they shrink the values of the estimated coefficients closer to 0. They do this by using a lambda value and penalizing variance in the models creating a bias in the model to give more accurate predictions.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\frac{5.76}{4}\left(\frac{.25}{.75}\right) \rightarrow 1.44\left(\frac{.25}{.75}\right) = .48$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is an unknown constant and a true population parameter.
$\hat{\beta}$ is a coefficient we derive from fitting a model to sample data.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Ordinary least squares estimates regression coefficients by creating a linear model that estimates and minimizes the sum of squared residuals.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

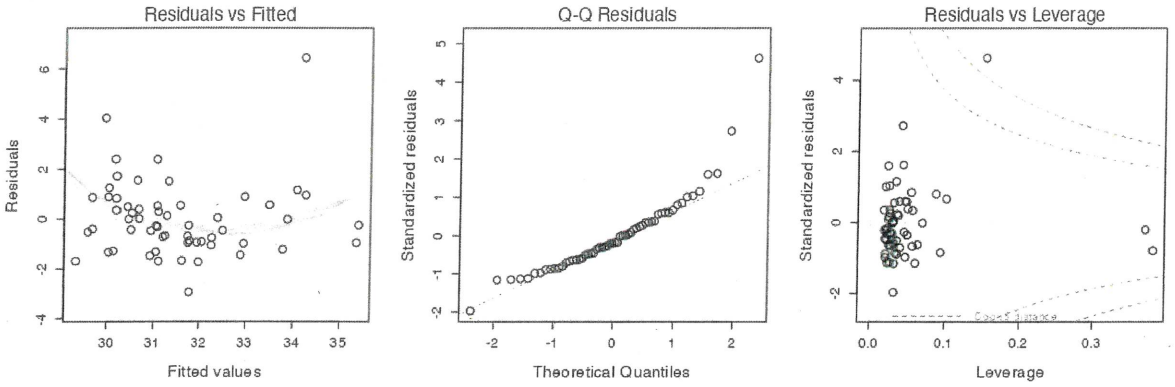

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 \approx .4729$

Holding minutes constant, a one-inch increase in height relates to an increase of .4729 hot dogs eaten on average.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

In our fitted vs. residuals plot, we see a slight curve in the data which could show non-linearity. In the added-variable plot for minutes we see a bigger curve which shows this data does not satisfy linearity very well.

(B) There is independence between each observation

Independence is questionable in this model. All of the contestants are from the same state, and even though they are mostly from the 4 largest counties, there would be reason to believe each contestant is from similar demographics.

(C) The errors are normally distributed

In the Q-Q plot we see a significant tail and s-shape. Combining this with a shapiro-wilk test who's p-value is much less than $\alpha = .05$, we can say the errors are not normally distributed.

(D) The errors have constant variance (homoscedasticity)

There is not a very distinct funnel shape in our fitted vs. residual plot. A Breusch-Pagan test gives a p-value of .071 which is greater than $\alpha = .05$. It seems like the equal variance assumption is satisfied here.

(E) There are no unduly influential points

Looking at the residual vs. leverage plot and the added-variable plots, Observation 49 is a point that could have some influence. We would want to further investigate that observation.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

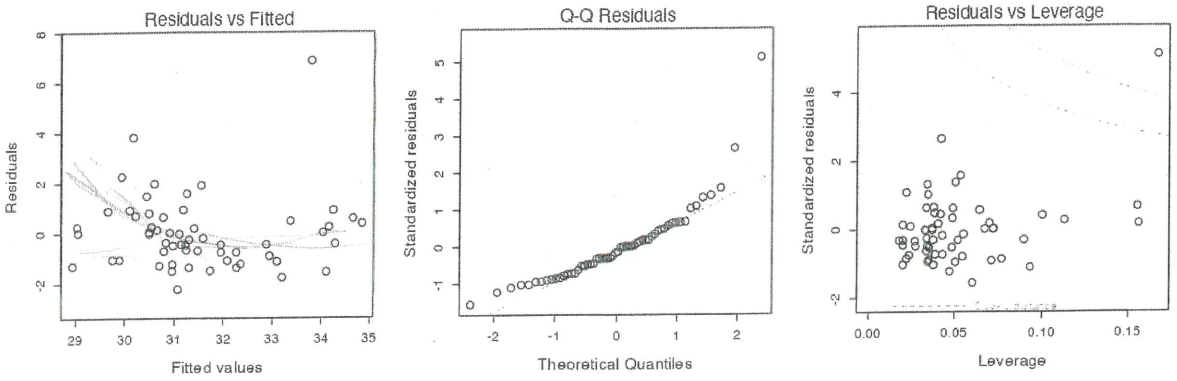

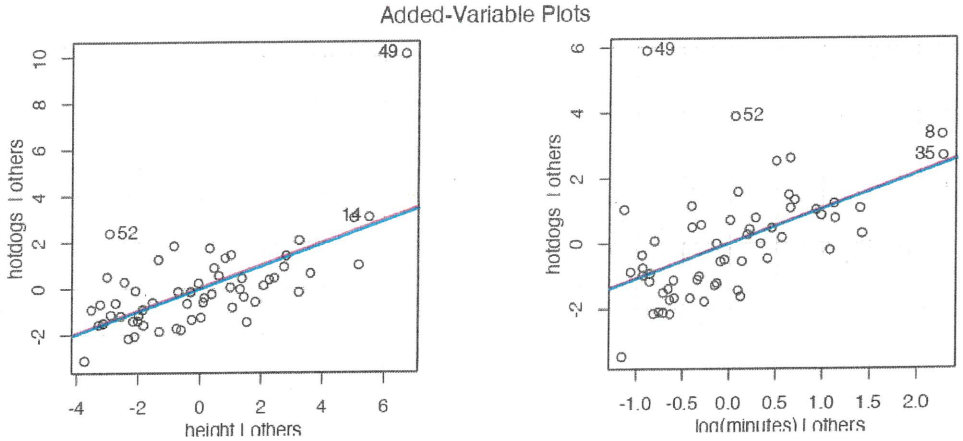

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

This is an overstatement when comparing models. A Breusch-Pagan test only shows evidence that the errors are not of equal variance. We should investigate our other assumptions and compare RMSEs before making that strong of a conclusion.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

When looking at linearity we see that minutes satisfies the assumption much better and the curve on the residuals vs. fitted plot is less severe. Model 2 also gives a lower CV RMSE compared to model 1. However, we did obtain a significant p-value in our Breusch-Pagan test at $\alpha = .05$ which model 1 does not. We also transformed our minutes predictor which changes our interpretations of the coefficient away from the simple unit increase.

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 \approx 1.06702$

A 6.7% increase in minutes trained, relates to a .03 increase in hot dogs eaten on average.

$\beta_2 \log\left(1+\frac{6}{100}\right)$?

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that $\log(79) = 4.369$. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.06689 + .488834(68) + 1.06702(4.369)$

31.81 Hot dogs

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31.81 \pm 2.004 \cdot (.204 + 1.485)$

$(28.353, 35.198)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

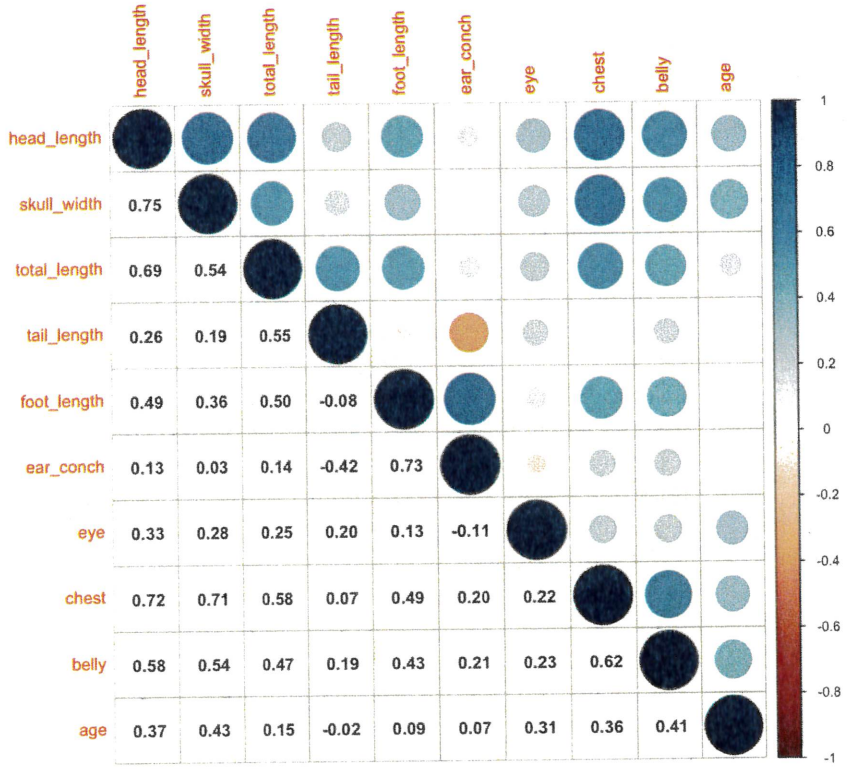

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$.16589 \pm 2.01 \times .10599$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that the true coefficient for belly girth is between -0.023 and 0.443.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

These significance tests have different hypothesis and only give us evidence. The f-test shows us there is evidence that when accounting for all the predictors together, at least one is not equal to zero. Our individual coefficient p-values just give us lack of evidence that the single coefficient is not 0. These do not give us concrete statements, they just test different things and give evidence or a lack of evidence.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The best subsets algorithm moves through all of the possible subsets of predictors, then will select the subset of predictors with the lowest AIC. It will calculate the AIC of each subset to do so.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM