STAT 330-001

40142034 6

Haglund, Quinlin

159850864 qhaglund 860-772-9239

Test: 636298 Form: A
Entry: 04:11pm 10/19/2026
Employee: Ashlyn T.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly.
Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

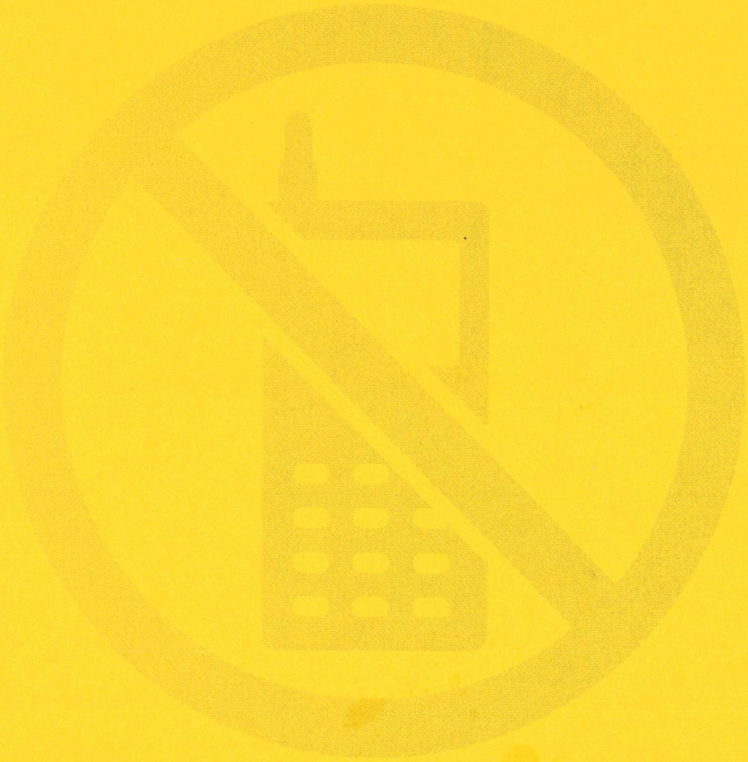

NS

40142034

006

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE

└ log transforms

- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression. → Line? but now influence too? so no

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They shrink the coefficients, thus they are regression called shrinkage methods. They introduce bias to reduce variance, in doing so they "shrink" the regression coefficients. LASSO can set some coefficients to zero, ridge just shrinks everything smoothly towards zero.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$D_i = \frac{2.4^2}{4} \cdot \frac{.25}{.75}$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is a theoretical vector of unknown population parameters, while $\hat{\beta}$ is the vector of coefficients found by conducting MLR → used to produce fitted values.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least squares is the sum of all the squared residuals, (difference between fitted & actual value)

Minimizing this (that sum) by manipulating regression coefficients gives use the coefficients that best describe the behavior & effects of predictors (usually) in the dataset.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

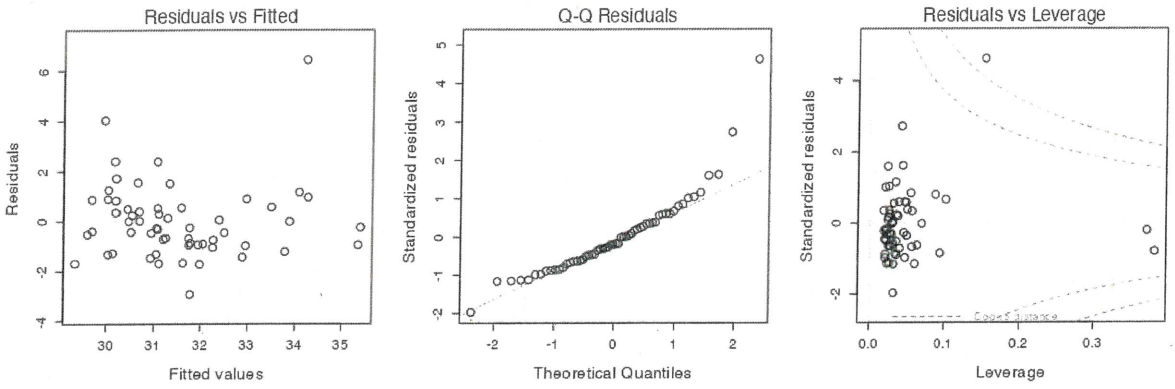

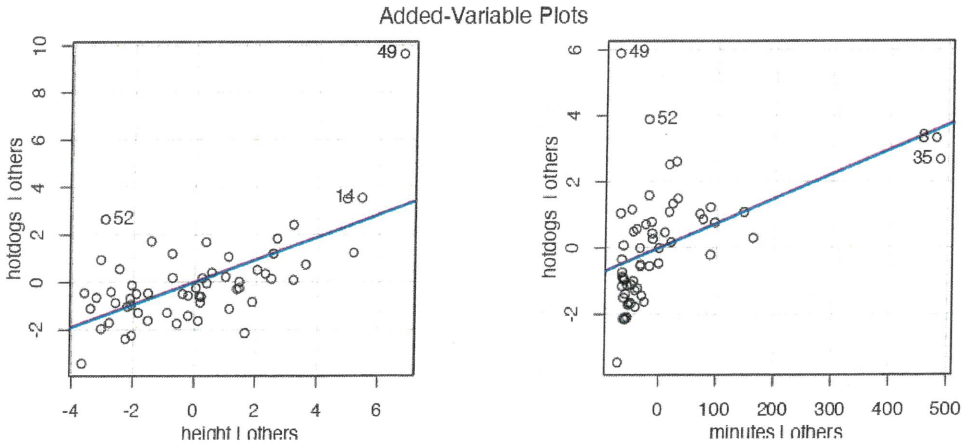

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = .47889$ A one inch change in height is associated with a change of .47589 hotdogs (consumed in one 10-minute sitting) on average, holding all other predictors constant.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$ ($\hat{\beta}_1 = .4729$)

I had to get a t* from question 14.

(Increase in height associated w/ increased hot dog consumption)

$.007417 \pm .00187 \cdot 2.004$

↖ am I supposed to have a t_{55} distribution memorized?!

↗ what?

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

No. Minutes looks remarkably suspect in the added variable plot. It is roughly linear for height though

(B) There is independence between each observation

Mostly true. These guys probably don't travel together, & one man's height doesn't influence another. However, competitive pressure might influence how many hot dogs they eat.

(C) The errors are normally distributed

No. Q-Q plot is not the right shape for a normal distribution, & Wilks (Shapiro) P-val is ~0, thus rejecting normality.

(D) The errors have constant variance (homoscedasticity)

Mostly. Residuals vs. fitted plot narrows ~~considerably~~ from left to right, but Breusch-Pagan test fails to reject null

(E) There are no unduly influential points

False. One point in Residuals vs leverage plot that is way out there (outside cooks cutoffs).

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

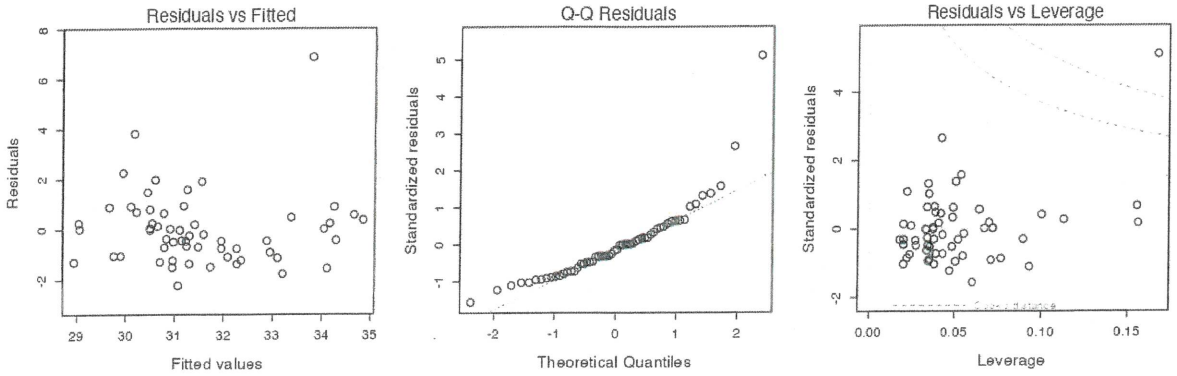

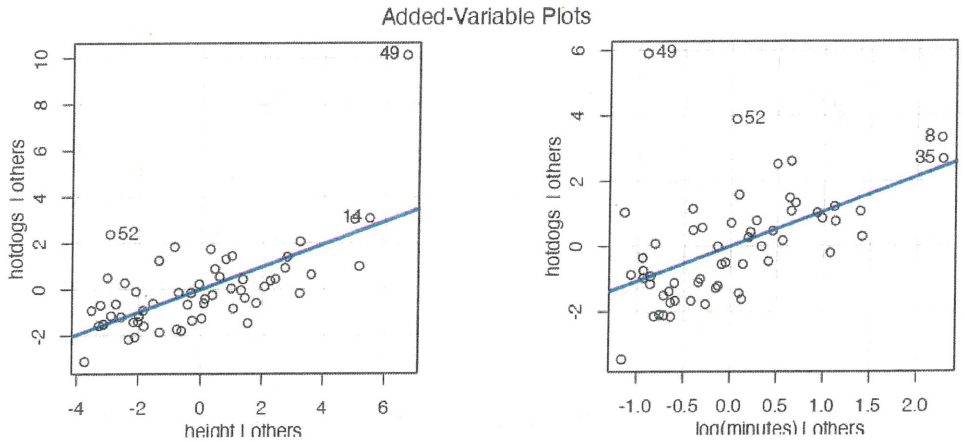

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

It's not that simple. Look at the AV plots, w/ the log transformation, we have linear relationships. Also, the CV RMSE is better than model 1, so predictive power is better.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Model 2 Advantages	Disadvantages
Better prediction (lower CVRMSE)	More complex interpretation
Slightly more normal	Error variance is unequal
↳ larger SW p-value	↳ significant BP
AV plots are both linear now	Still have one high influential observation

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$ a 10% increase in the time spent training (in minutes) is associated with a .101 increase in hotdogs eaten in 10 minutes, holding height (our other predictor) & any other predictors constant.

$\hat{\beta}_2 \cdot \log\left(1+\frac{c}{100}\right)$

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

Hotdogs = -6.05589 + .48834 · height + 1.06702 · log(min)

68 79

4.369

31.52 = 31.52 hotdogs in 10 minutes

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

$n-p-1$

$58-2-1 = 55$

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1+x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

used this t^* in question 12

$X^{*T}\hat{\beta} \pm t^*_{55,.975} \cdot s\sqrt{1+x^*(X^\top X)^{-1}x^*}$

91

31.52

$2.004 \cdot 1.485$

2.97594

Upper PI = 34.49594

Lower PI = 28.54406

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

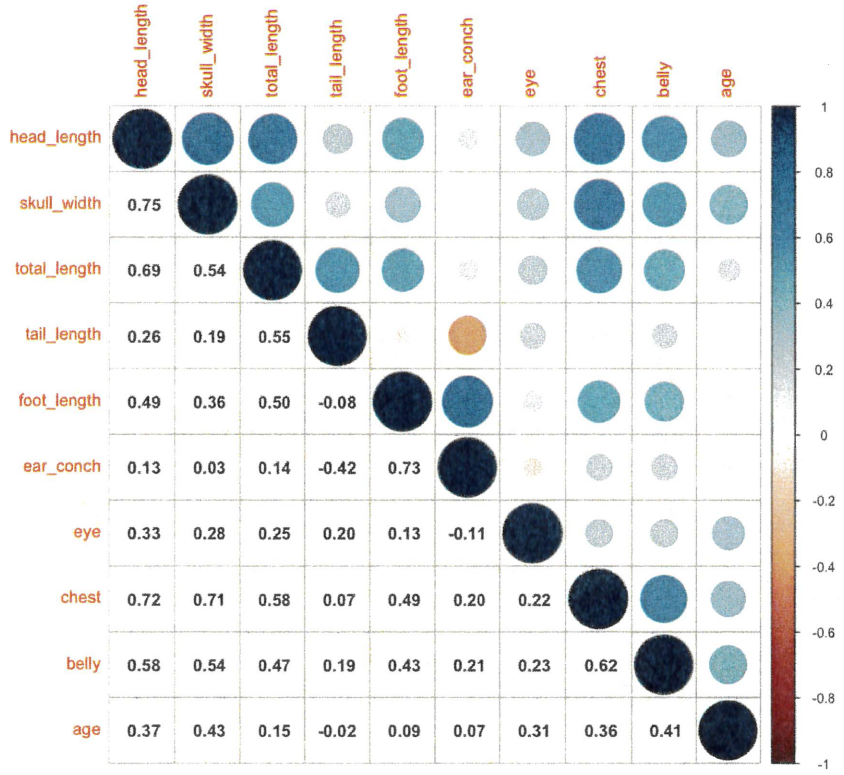

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$.1655 \pm 2.01 * .1055$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that the true population value for bellywidth effect is between -.023 & .443.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

Multiplicity! We have lots of variables We are performing lots of tests so the F-test on model significance is going to have a lower p-value than any one coefficient

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets goes through every possible model construction (from one variable to all nine) & every possible combination → 2^p measuring the AIC of each model. It selects the model with the smallest AIC as its "best subset" or best candidate model.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 5

Santacruz Concha, Jorge Emiliano

239004630 mxesco 385-221-8630

Test: 635556 Form: A Wadsworth, Spencer Gordon

Entry: 01:44pm 10/19/2026 Checked out:

Employee: Nathan M.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

005

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

LASSO and ridge shrink regression coefficient toward 0 by adding a penalty to the least squares objective. Ridge uses an L2 penalty, while LASSO uses an L1 penalty, can set some coefficient exactly to 0.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

4/100 = 0.04

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$ $D_i = 0.48$

(Circle one answer) Which conclusion is most appropriate? $D_i = 0.2$

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is the true coefficient, while $\hat{\beta}$ its estimate calculated from sample data.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least squares finds their regression coefficient that minimize the sum of squared differences between observed and predicted values

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

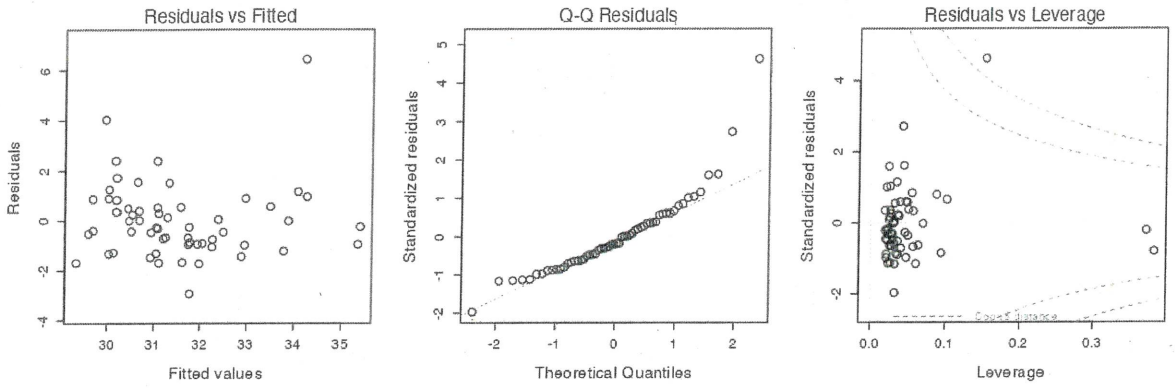

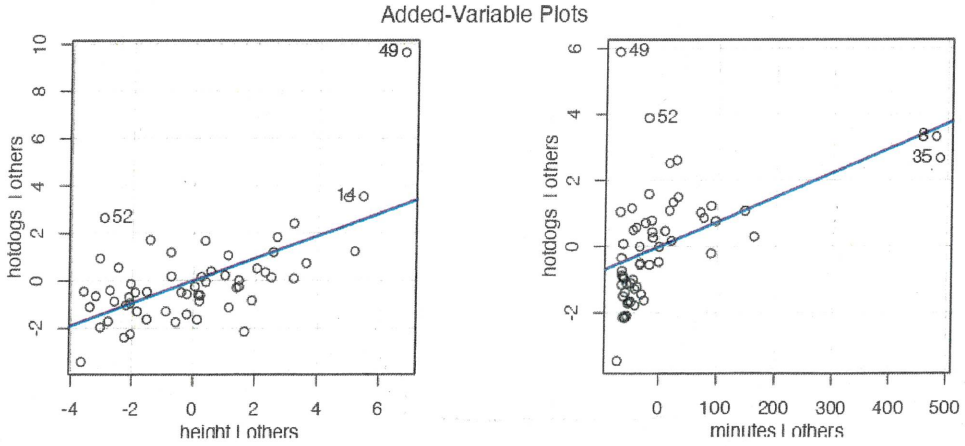

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

For each additional inch of height, the expected number of hot dogs eaten increases by $\hat{\beta}_1 = 0.472892$, holding training minutes constant

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The residuals vs fitted plots show a possible pattern, so linearity is questionable.

(B) There is independence between each observation

Independance seems rasonable because each observation is from from different constant, but it is not guaranteed

(C) The errors are normally distributed

The errors are normaly not supported. Shapiro-Wilk $p = 2.43 \times 10^{-8} < 0.05$. The Q-Q plot shows strang deviations.

(D) The errors have constant variance (homoscedasticity)

Constant variance seems razonable Breusch-Pagan $p = 0.71 < 0.05$. Although it's not guarenteed.

(E) There are no unduly influential points

The residuals vs leverage plot, shows a few points with high leverage and large residuals, suggesting possible influencial observations. These should be checked using Cook's distance.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

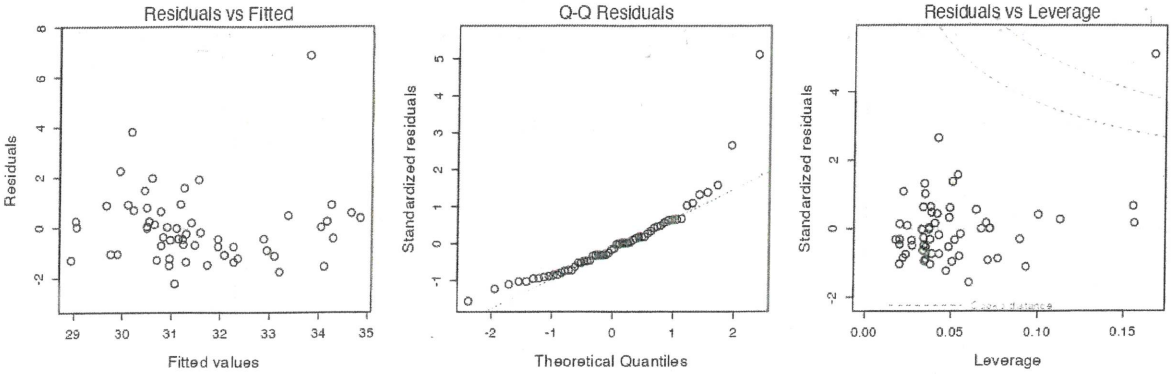

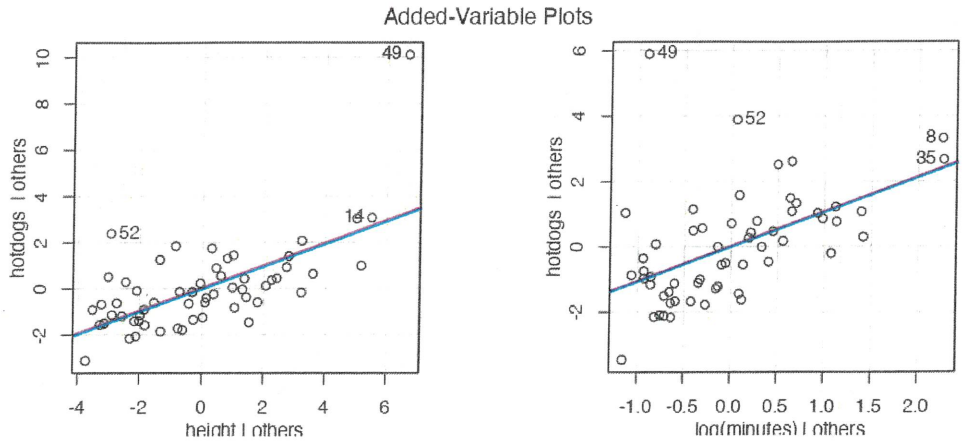

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I disagree. Model (1) has less evidence of heteroscedasticity ($0.071 > 0.05$), but model (2) has a lower CV RMSE 1.667 vs 1.846. We should consider prediction accuracy, normality, and residual plots before choosing a model.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Model (2) has better predictive accuracy CV RMSE 1.667 vs 1.846 and less evidence against normality. However, it shows evidence of heteroscedasticity $p = 0.028$, while model (1) does not $p = 0.071$. Both models have normality problems.

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$

Holding height constant, a 1% increase in training minutes is associated with approximately 0.01067 more hotdogs eaten in 10 minutes.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 + 0.48834(68) + 1.06702(4.369) = \boxed{31.81}$

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$ $\quad s\sqrt{1+x^*(X^\top X)^{-1}x^*} = 1.485$

$t^*_{df,0.975} = 2.004$ 95% consume in 10 minutes

$31.81 \pm 2.004(1.485) = \boxed{(28.83, 34.79)}$ hotdogs

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

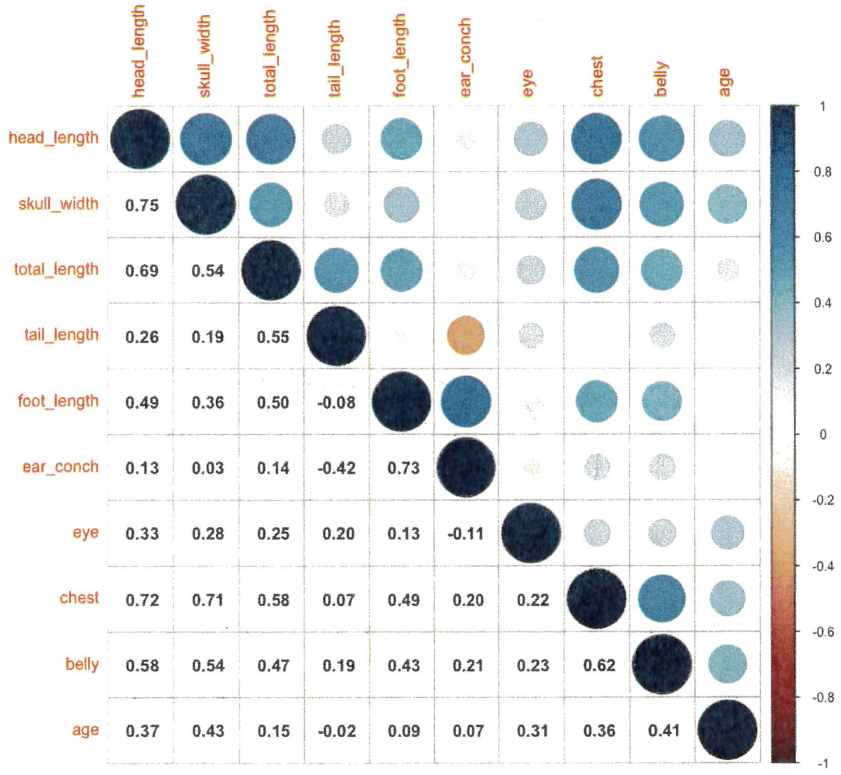

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

Given 0.16559 and 0.10599 $t^* = 2.01$

$0.16559 \pm 2.01 * 0.10599 = (-0.04745, 0.37863)$

$(-0.04745, 0.37863)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that each additional cm of belly width is associated with a change between (-0.023, 0.443) cm in skull width, holding other variables constant.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

Multicollinearity may be causing this. The predictors are highly correlated, making individual coefficience less significant, even though the overall model is statistically significant.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets evaluates all possible combinations of predictors, calculate the AIC for each model, and select the model with the lowest AIC

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

Scratch Paper

$$9 = D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right)$$

a) $D_i = 48$ $4/n = 0.04$

b) $D_i = 48$ $4/n = 0.04$

c) $D_i = 48$ 0.50

d) $D_i = 0.2$ $4/n = 0.4$

e) $D_i = 0.2$ $4/n = 0.4$

f) $D_i = 0.2$ 0.50

g)

Part II

58-3=55

Scratch Paper

STAT 330-001

40142034 4

Davis, Matthew Divers

404666721 mdovis 203-464-5660

Test: 634962 Form: A Wadsworth, Spencer Gordon

Entry: 08:44am 10/19/2026 Checked out:

Employee: Annabelle H.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034

NS

004

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are shrinkage methods because they are shrinking the $\hat{\beta}$ coefficients. It is changing λ coefficients repeatedly to find a value with the lowest error. Then you can see where the best coefficients are.

$\frac{2.4^2}{3+1} = 1.44$ $\frac{0.25}{1-0.25} = \frac{1}{3}$

$1.44 \cdot \frac{1}{3} = 0.48$

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$? I'm not sure here but I know its matrix algebra

$(X^TX)^{-1}$ [illegible] $= \widehat{\beta}$

or

lm(y ~ $X_1 + X_2 + \cdots + X_n$, data) ← R function to get the model

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\widehat{\beta}$ is the estimated coefficient values where β is the actual coefficient values

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

The goal of least squares is to minimize the least squares error, so when you get the minimum value of sum of least squares you have the optimal $\widehat{\beta}$ values

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

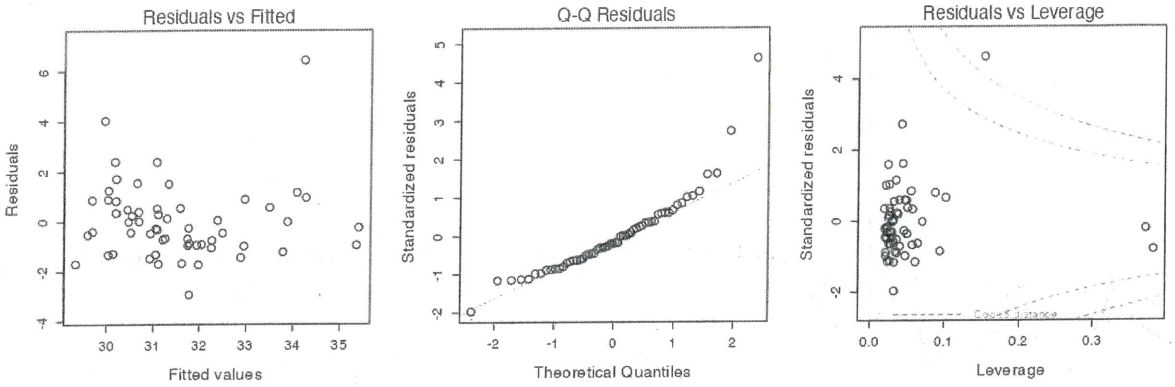

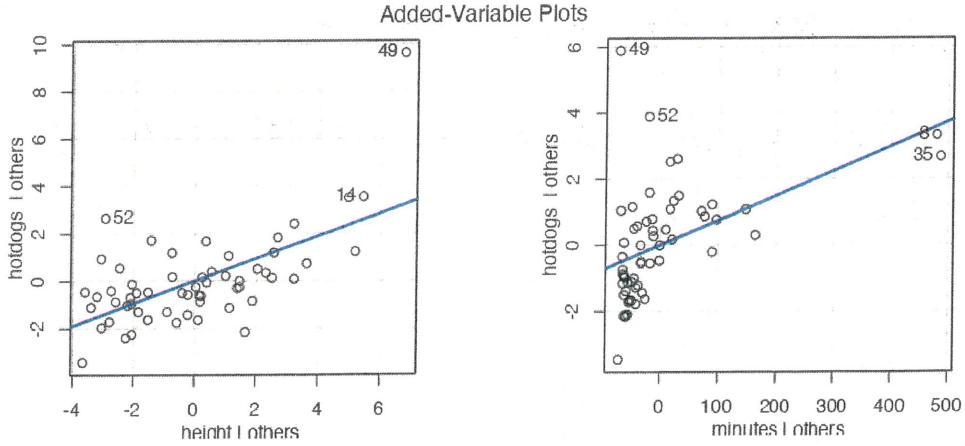

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

$0.007417 \pm 3.963 \cdot 0.001871 = 0.000002227,$
0.014831773

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = 0.472892$ This means for each inch taller someone is that is associated with a 0.47 hot dog ~~per~~ in 10 minute increase.

Question 12: **(Circle all ~~that~~ apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The added variable plots look like the the explanatory variables and response variables are not linear.

(B) There is independence between each observation

They seem independent.

(C) The errors are normally distributed

No. Not normally distributed (Q. Q plot. ~~Residual fitted plot~~).

(D) The errors have constant variance (homoscedasticity)

No, Residual and fitted plot

(E) There are no unduly influential points

No, Residuals and leverage plot.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

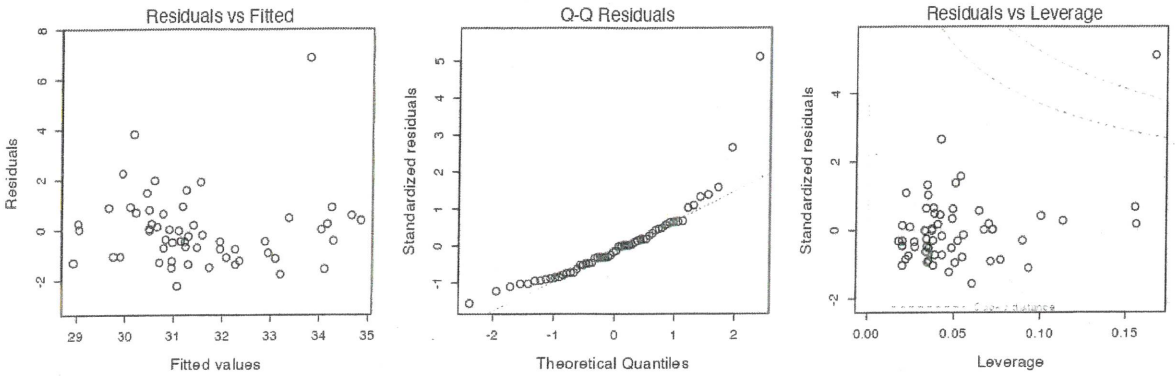

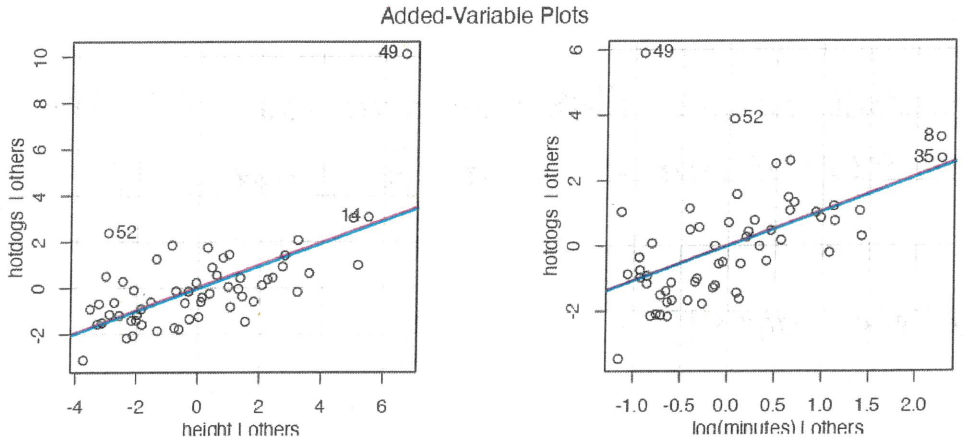

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would disagree. The BP test does not account for issues in linearity which is what model (2) is trying to fix. Also a BP test doesn't binarily decide if something is good or bad.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

It looks like the response/explanatory variables now have a linear relationship, but the error variance is still screwed up. Also, sometimes a log transform makes it harder to explain out comes.

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$

For each log(mins) increase the number of hot dogs increases by one.

I forget the real definition for this one.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 \quad \hat{\beta}_1 \, 68 + \hat{\beta}_2 \, 4.369 = 34.448$ hotdogs

$\hat{\beta}_1$: 0.48834 $\quad$ $\hat{\beta}_2$: 1.6702

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

31.47 , 37.42 Prediction interval

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

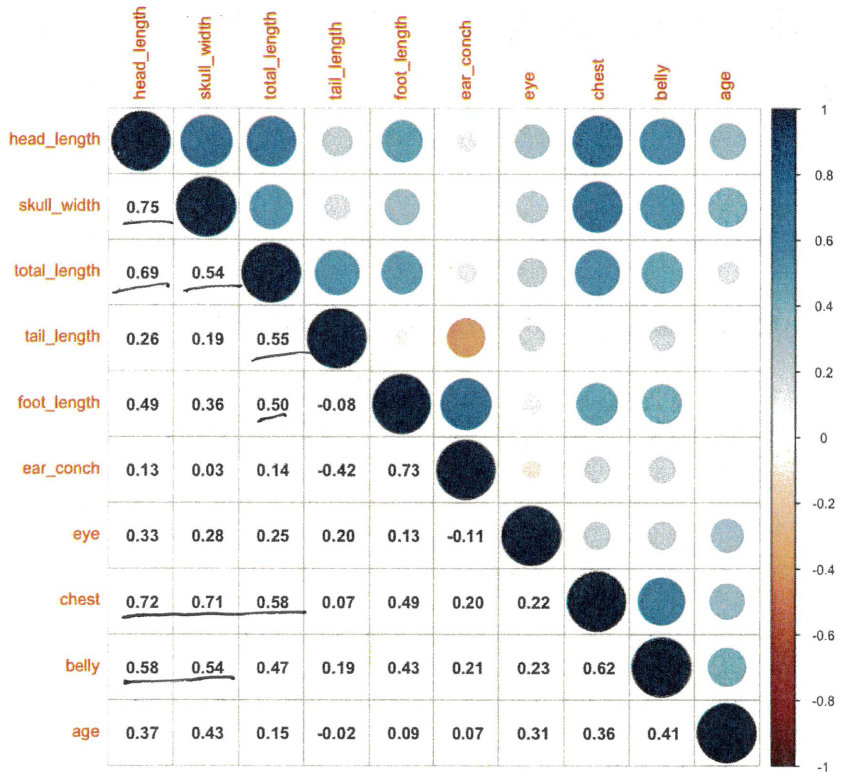

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 2.01 \cdot 0.10599$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

Because 0 is in the 95% CI that tells us that the true coefficient could reasonably be 0.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

This is telling us that together the predictors have a high likelihood of being significant. One reason this could be happening is multicolinearity.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

AIC recursively adds or subtracts variables to find the best model.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

B Average the predictions and record their value

D Choose the model with the largest RMSE

A Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

B Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

C Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 3

Wright, Jason

846461801 jasontw0 763-777-1458

Test: 634679 Form: A
Entry: 12:58pm 10/17/2026
Employee: Oscar J.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

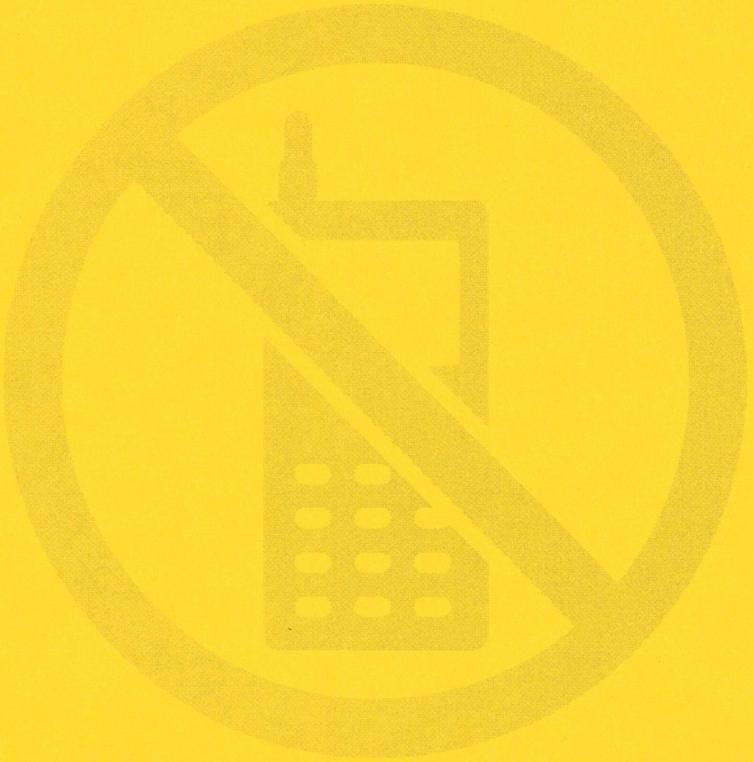

40142034

NS

003

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

Jason Wright

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

linearity of

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

LASSO and ridge regression are able to produce a sizable reduction in variance by introducing usually a smaller amount of bias. They both add a penalty to the OLS amount which leads to the lessening of coefficient amounts or sometimes a coefficient gets set to zero, eliminating a predictor. Because of this, variance is decreased.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is the actual/unknown value of the coefficient/slope for the population, where as $\hat{\beta}$ is just the estimate of the slope based on the sample.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

What Least squares (SSE(β)) does is find the value for β for which $\sum_{i=1}^{n} e_i^2$ is the smallest. By squaring the errors and finding which value of β makes the sum of those errors the smallest, least squares is able to provide a good estimate for the regression coefficient.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

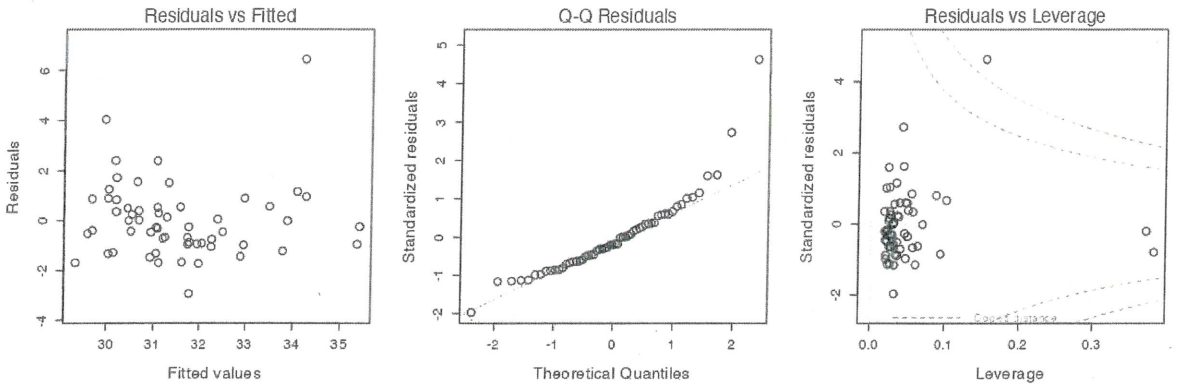

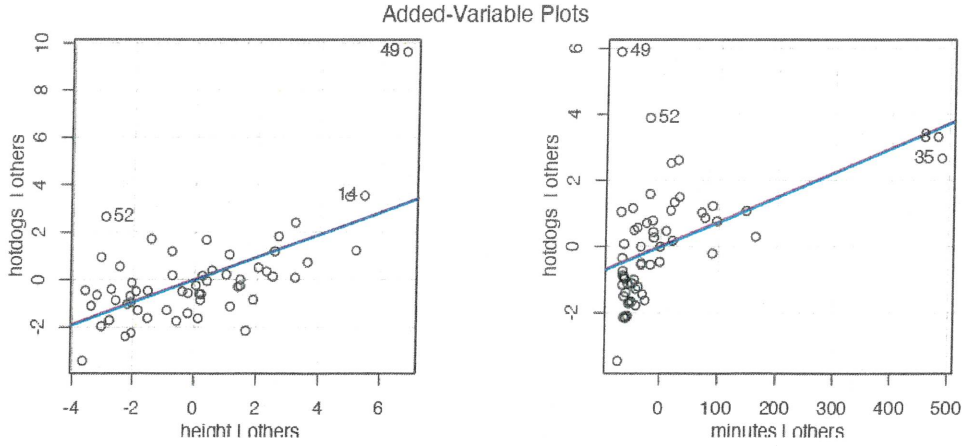

Coefficient	Estimate	Std. Error	t value	Pr($>\|t\|$)
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

0.472871

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

Estimated Value of β_1 : 0.472892

Interpretation :

Holding 'minutes' constant, an increase of 1 inch in the contestant's height is associated with an estimated increase of 0.472892 hot dogs on the average total number of hot dogs consumed in a single 10 minute sitting.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

$t^* \times SE = 0.472872$

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The relationships are most likely linear. The dots on the residuals v. fitted plot appear to follow a horizontal pattern, and the added-variable plot for height has the dots following a linear trend as well. The only non-linear evidence is the roughly non-linear look of the dots on the added-variable plot for minutes.

(B) There is independence between each observation

Yes, we can assume that each contestant's eating period wasn't affected by the others' and that each contestant was given a similar eating experience; so the errors are independent of each other.

(C) The errors are normally distributed

There does appear to be some slight evidence for non-normality, but I'm going to assume normality. The Shapiro-Wilk test did result in a small p-value and there is some slight curving on the ends of the Q-Q plot line, but for the most part, the dots on the Q-Q plot very closely stick to the line.

(D) The errors have constant variance (homoscedasticity)

There is some evidence for non-constant variance. The residual v. fitted plot shows some slight fanning and non-equal spread of the dots, and the Breusch-Pagan test p-value of 0.071 doesn't provide strong evidence that error variance is constant.

(E) There are no unduly influential points

There does appear to be a few potential unduly influential points when looking at the plots. The 3 dots, in particular, on sort of the top-right and bottom-left residuals v. leverage plot are sign of potentially strong influence.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

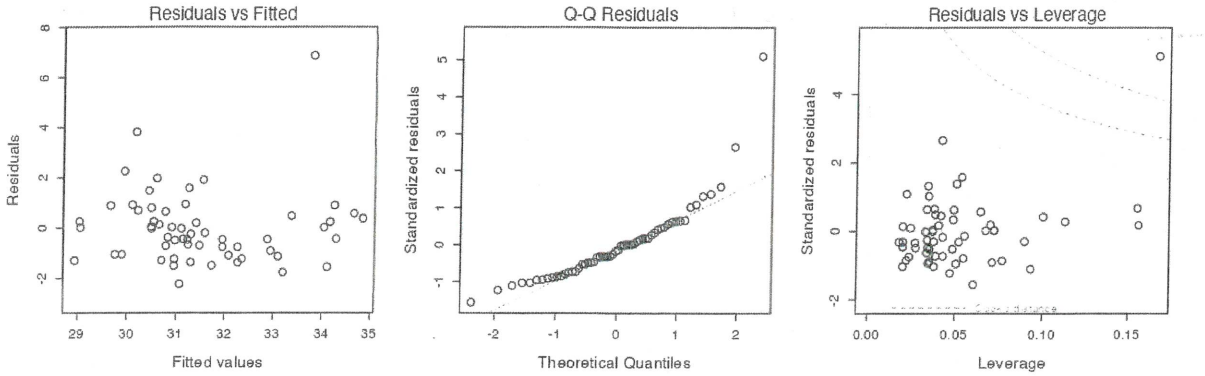

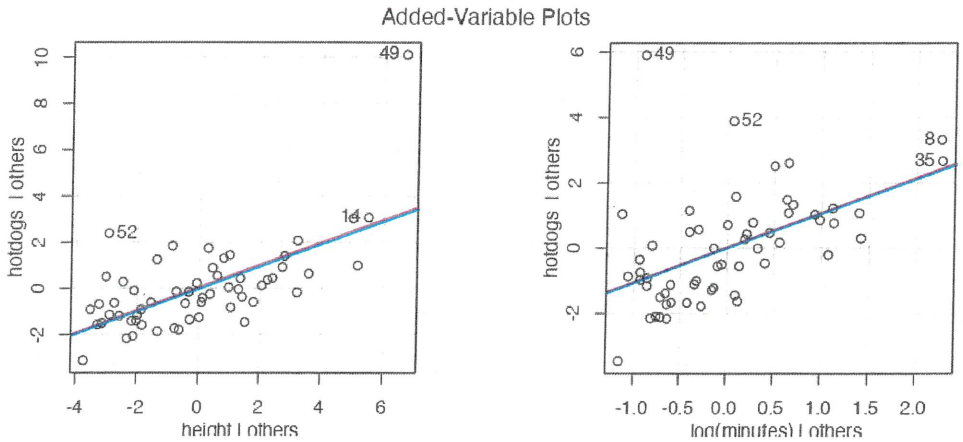

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would say that's not necessarily true. More information and analysis is required than just looking at the B-P p-values. The difference in the p-values isn't super large. For equal variance assumption comparisons between models, the residuals v. fitted plots + the scale-location plots should also be taken into consideration. Looking at the other L.I.N.E. assumptions along with other results like R^2 should also be noted when deciding which model is "better" than the other.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages: The level-log transformation led to a more linear-looking added-variable for 'minutes'. The p-value for the B-P test is also lower.

Disadvantages: The S-W test p-value is a bit larger. Its residuals v. leverage plot also shows some perhaps more influential points.

Question 16: State the estimated value of β_2 and provide an interpretation in context

Estimated value of β_2 : 1.06702

Interpretation:

Holding 'height' constant, a 1% increase in minutes is associated with an estimated 0.0106702 increase in the average total number of hot dogs consumed in the single 10 minute sitting.

31.81304038

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

33.20712

$\hat{y} = -6.05589 + 0.48834(68) + 1.06702(4.369) \approx 31.813$ hot dogs

↓

(or ≈ 31 hot dogs)

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

$n - p - 1$

$58 \quad 2$

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$\hat{y} \pm t^* \times SE$

$t^* \cdot SE = 2.004 \cdot 1.485 = 2.97594$

$31.813 + 2.97594 \approx 34.789$

$31.813 - 2.97594 \approx 28.837$

→ 95% Prediction Interval:

(28.837, 34.789)

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

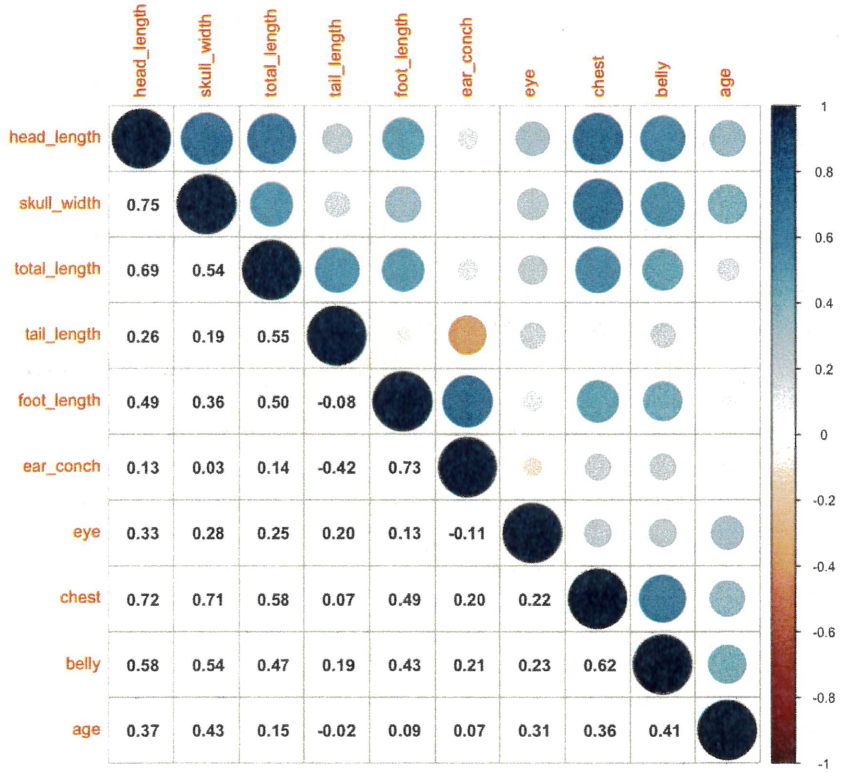

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

estimate $\pm$ t^* $\times$ SE

0.16559 2.01 0.10599

$0.16559 \pm 2.01 \times 0.10599$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

Holding all other predictors constant, we are 95% confident that the true value of the belly width coefficient is between -0.023 and 0.443.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

The null hypothesis for the F-test is that all the coefficients equal zero, and the alternative hypothesis is that atleast one slope isn't zero. So when we reject the null, we consider the predictors together and conclude that atleast one slope isn't zero; we are not saying that all the predictor's coefficients individually have a significant effect.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The "best subsets" algorithm, in general, goes through all possible combinations of the predictors when choosing a model. AIC is a metric used to choose which of these model options is best. It selects which predictor combination to use through the formula "$n\log(SSE/n) + 2k$". AIC is most useful when trying to predict and when you're ok having a larger model.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

B Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 2

Anderson, Tyler Del

651205673 tylerda2 801-699-3978

Test: 633271 Form: A Wadsworth, Spencer Gordon

Entry: 04:18pm 10/16/2026 Checked out:

Employee: Isaac C.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034
NS
002

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are adding bias in order to shrink the λ. We shrink coefficients to zero/close to zero to improve our interpretations

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$D_i = \frac{5.76}{4}\left(\frac{0.25}{0.75}\right)$

$D_i = 1.44\ (0.333)$

$D_i = 0.4799999$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is to a fixed set of data

$\hat{\beta}$ can be applied sample to sample

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

We can use least squares to find which SSE is the smallest in order to give us the desired OLS.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

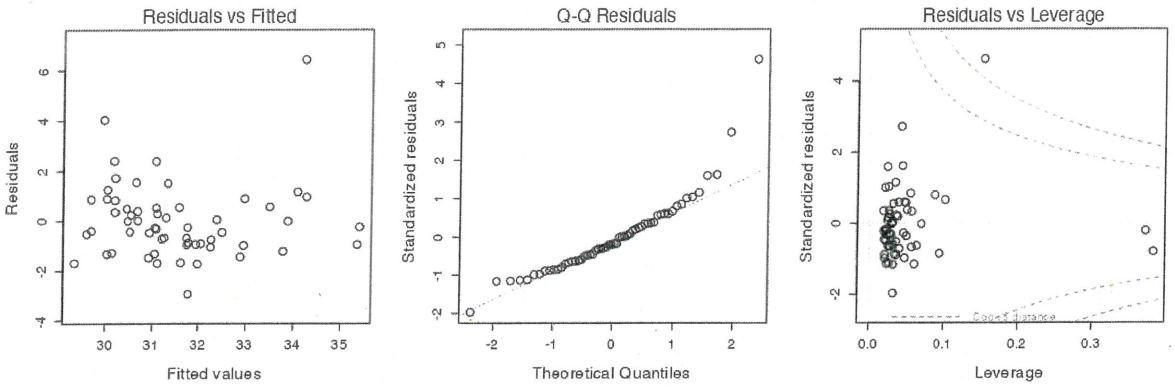

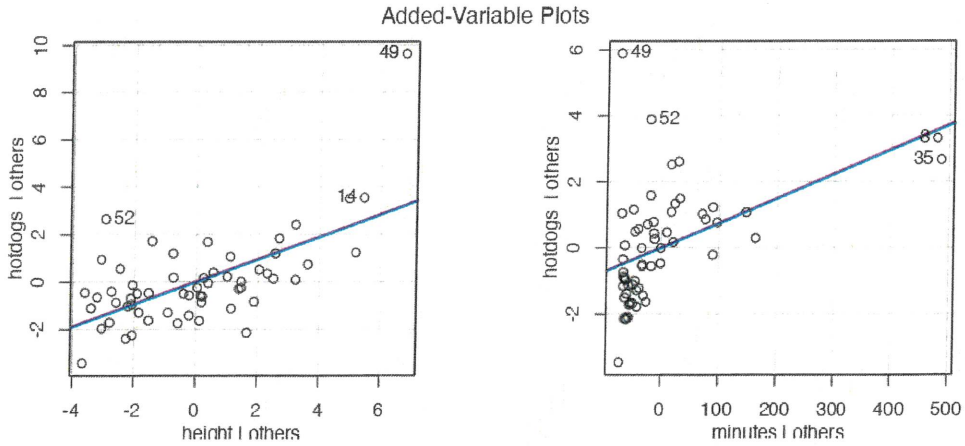

Coefficient	Estimate	Std. Error	t value	Pr($>\|t\|$)
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

B_1 = height, the estimate is 0.4172892.

For every unit change in minutes, height increases by .4172892 as an estimate holding everything else constant

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

$\hat{B_j} \pm (SE(\hat{B_j}))$

$\hat{B_j} \pm 0.237(5.46771(\hat{B_j}))$

-1.295392 ± 0.167863

0.067417 ± 0.000055

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

While there are some data points scattered around 0, I would say the residual graph raises concern because a line would be curved I don't believe it's very linear.

(B) There is independence between each observation

There is independence. They selected random enters that weren't affected by other people.

(C) The errors are normally distributed

The Q-Q plot shows a pretty solid line followed until the end of the graph where it curves Because the Shapro-Wilk test shows a p-value of $2.43e^{-8}$ (0.0008) I would say it's not normal

(D) The errors have constant variance (homoscedasticity)

the residual vs leverage shows a bunched up spot of data, while this isn't a funnel shape it could be considered not constant but the Breusch-Pagan Test p-value of 0.071 implies it probably is equal variance.

(E) There are no unduly influential points

There are some points that seem to be outliers like (49) that could be considered to be removed

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

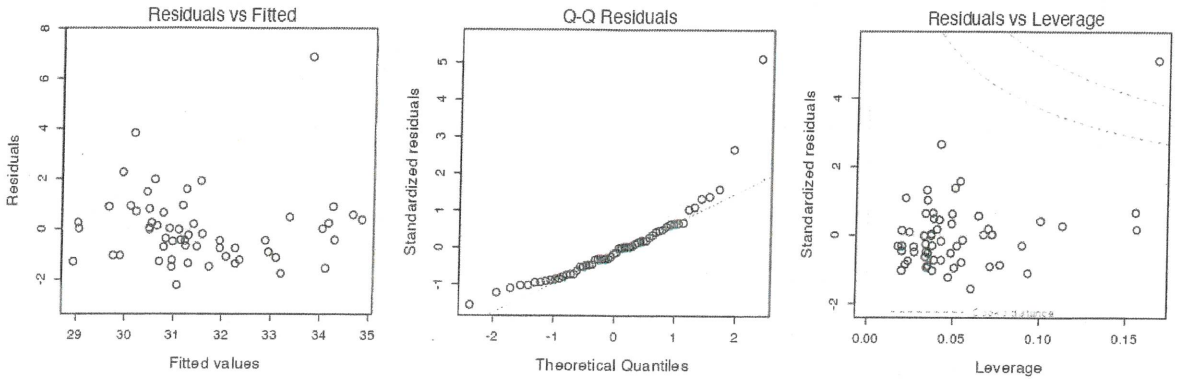

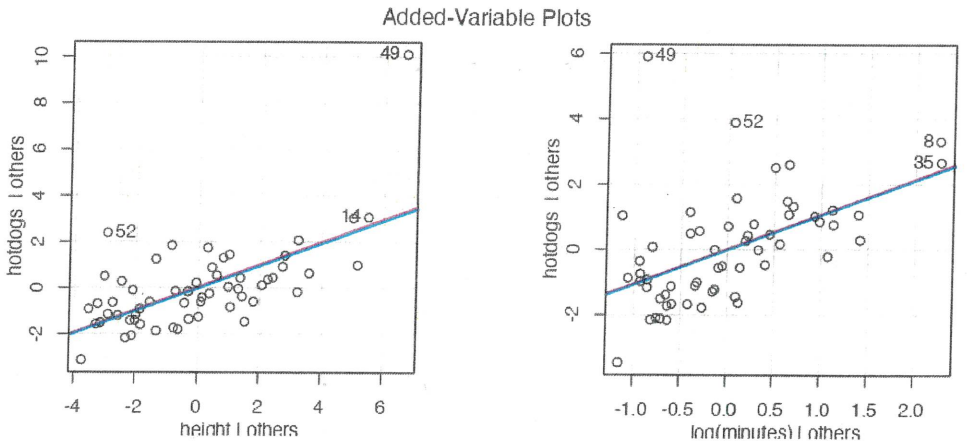

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

While the p-value of model 1 may be higher, you shouldn't conclude that because of one p-value it is superior. You should consider things like RMSE and number of predictors to decide

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Model 2

advantages:
- lower RMSE
- higher Shapiro-Wilk p-value (more normal)

disadvantages:
- lower P-value
- not equal variance

Question 16: State the estimated value of β_2 and provide an interpretation in context

B_2 log(minutes) is equal to 1.06702. As there is a one-unit increase in minutes there is a median change of 1.06702 B_2 log(minutes)

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$\log(79) = 4.369$
$B_2 = 79$ minutes
$B_1 = 68$ inches

$e^{4.369} = 78.965$ hot dogs

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

$df = n - p - 1$
$58 - 2 - 1 = 55$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$t = 2.004$ $\quad 1.485 \pm 2.004(0.204(1.485))$

$1.485 - 0.6071 = 0.8779$
$1.485 + 0.6071 = 2.0921$

$(0.8779, 2.0921)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

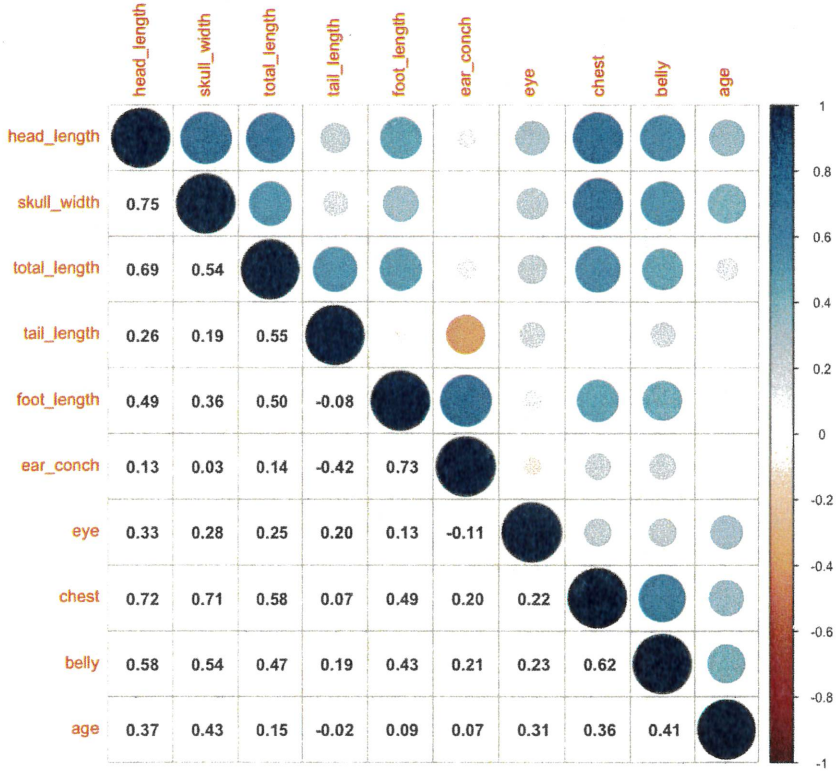

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

t= 2.01
SE= .10599
est = 0.16599

$0.16599 \pm 2.01(.10599(6.16599))$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident the mean for belly width will be within the range of (-0.023, 0.443)

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

We may assume with our p-values that there is no multi-collinearity but by just relying on p-values we may miss things like linearity or independence

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The best subsets algorithm determines which model is examined and the AIC model can be scored the highest as a selection metric seemingly making it the best model to use. AIC is not complicated as a model and depends on the penalty added when extra variables are introduced to see A's value.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

B Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

C Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

D Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

D Use the models fitted to the training data to get predictions for the test data

END OF EXAM