STAT 330-001

40142034 1

Jorgensen, Tyler S Thorne

120497550 tjorgy04 801-708-2915

Test: 632340 Form: A
Entry: 12:52pm 10/16/2026
Employee: Luke Prestley M.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

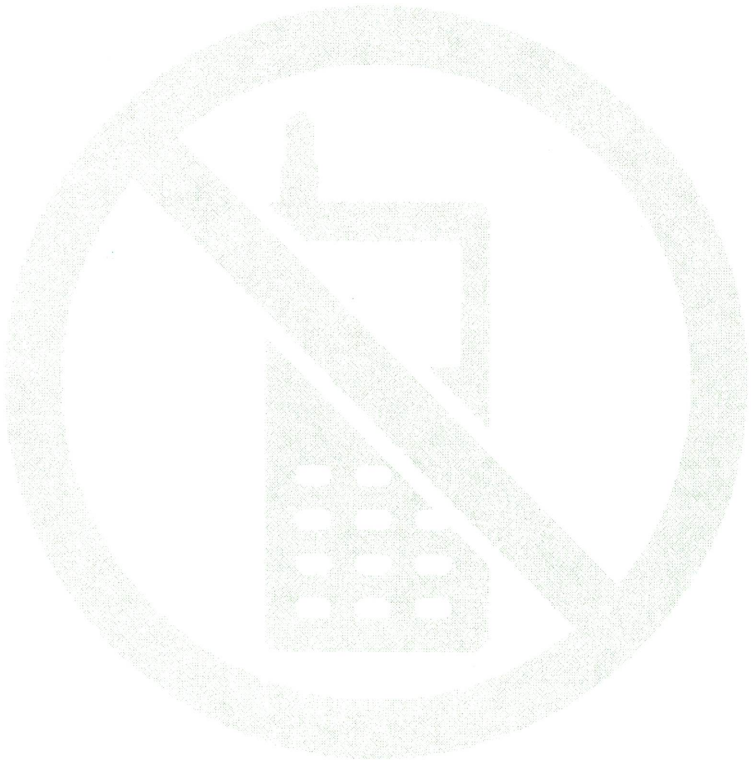

40142034
NS
001

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

Lasso and ridge are sometimes referred to as "shrinkage" methods because the coefficients are being shrunk towards zero. This is done by adding either $\lambda\sum|\beta_i|$ or $\lambda\sum\beta_i^2$ to the sum of squares error which is being minimized. Doing so punishes large coefficients more introducing some bias, but removing some variance.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is our estimate of β. β is the true population coefficients which we can not truly know, $\hat{\beta}$ is our estimate of what β is based on the sample date that we have

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

least squares is used to estimate regression coefficients, by finding the $\hat{\beta}$ estimated coefficient vector where the sum of the residuals squared is the smallest it can possibly be. This is accomplished by doing the Matrix multiplication $(X^T X)^{-1} X^T Y$

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

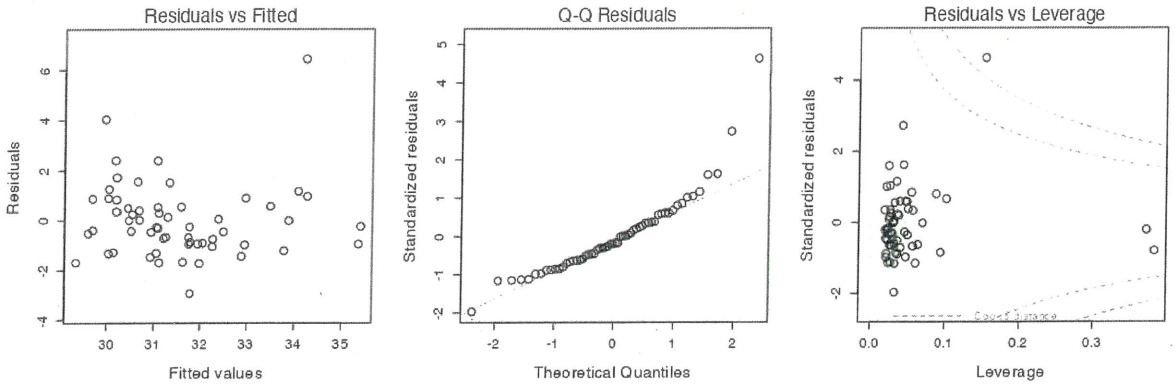

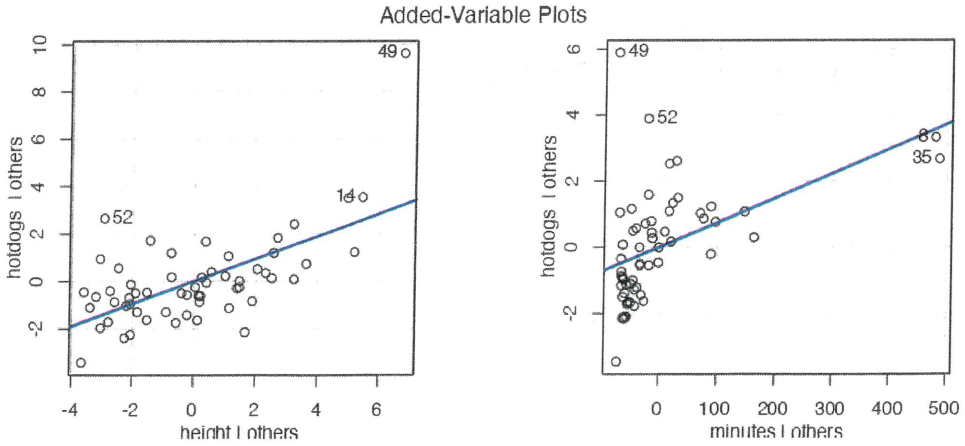

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

.472892

holding all else constant you would expect a eater that is an inch taller to eat on average .472892 more hotdogs in the 10 minutes

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

Looking at the residual vs fitted plot there are some concerns with there being a seeming bend in the data and large negative outliers. The minutes added variable plot is especially concerning and does not appear very linear at all. The height added variable plot looks more linear, with a large outlier.

(B) There is independence between each observation

There good reason to believe that independent, but there could be some independence problems such as people training in groups, or who they are competing next too affecting their scale.

(C) The errors are normally distributed

Most of the Q-Q plot looks promising and normally distributed. The big outliers on the right end are worth investigation though

(D) The errors have constant variance (homoscedasticity)

In the residual vs fitted plot there seems there might be a bit of a funnel with the variance being greater on the left

(E) There are no unduly influential points

the residuals vs leverage plot seems to indicate there is one unduly influential point in the top right of that graph. Based on the other plots I assume that is observation 49

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

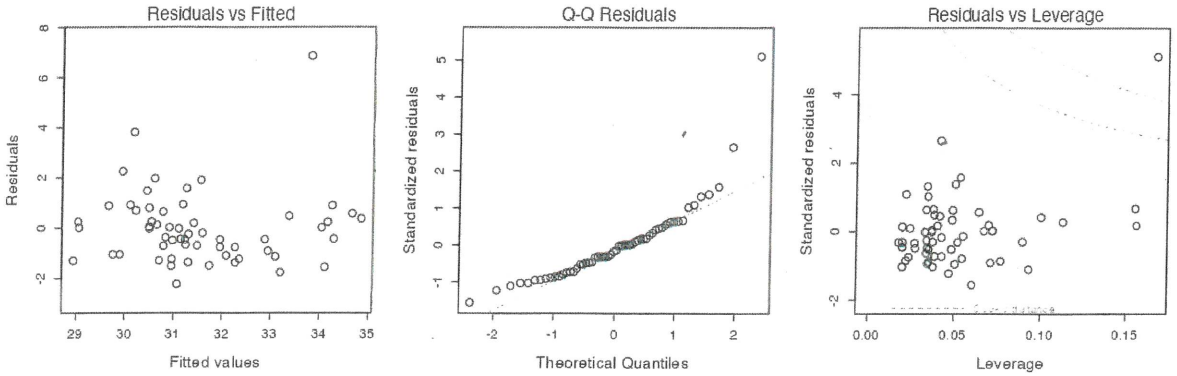

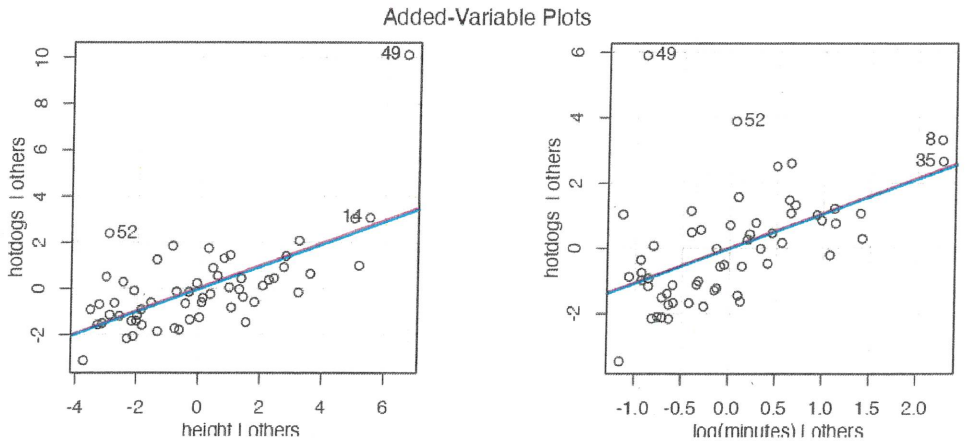

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would say that that is only one measure of the accuracy of the model and the difference between .071 and .028 is not particularly large.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Model 2 seems to have largely solved the linearity problem previously had with the minutes variable. It is now a bit more difficult to interpret after the transformation a x% change in minutes now correlates to a $1.0670^{\log(1+\frac{x}{100})}$ change in the mean # of hot dogs eaten for someone with that height and minutes.

Question 16: State the estimated value of β_2 and provide an interpretation in context

The estimated value is 1.0670 this means that holding all else constand a x% change in minutes correlates with $1.0670^{\log(1+\frac{x}{100})}$ increase or decrease in the mean number of hot dogs eaten by someone with these minutes and that height.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that $\log(79) = 4.369$. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 + .048834(68) + 1.06702(4.369)$

$= 31.81304038$

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31.81304038 + 1.485(2.004) = 34.78898038$

$-1.485(2.004) = 28.83710038$

95% prediction interval:

$28.83710038 \leq y_i \leq 34.78898038$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

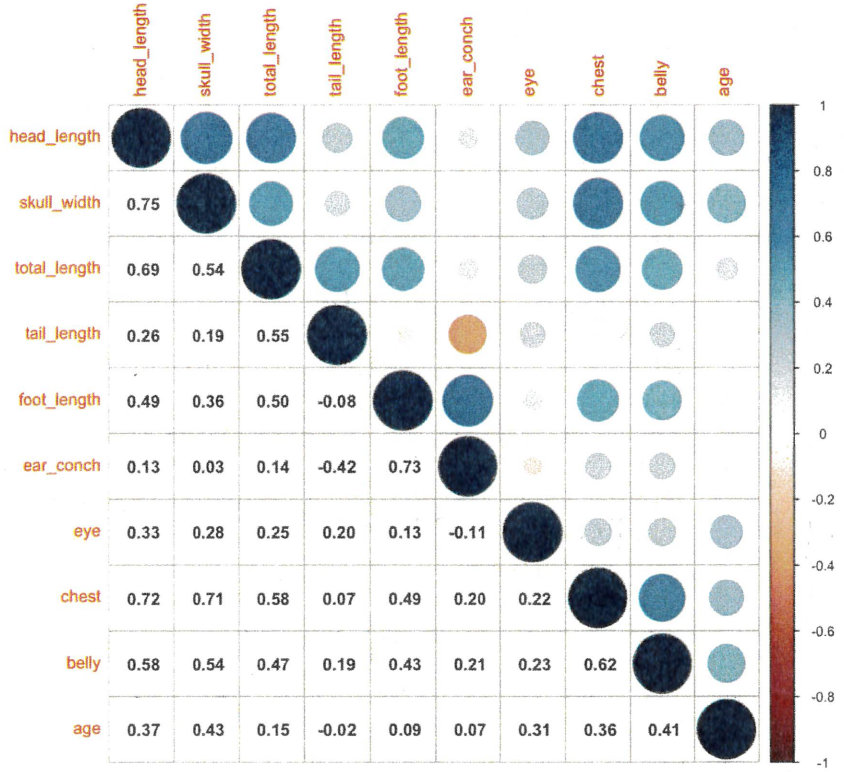

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$.16559 \pm .10599(2.01)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

If we were to repeat this same process of measuring possums and creating confidence intervals 95% of those intervals would contain the true coefficient for belly width's impact on a possum's age.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

multicollinearity is driving this contradiction. In the model each of those coefficients are found by holding all else equal, which is problematic, because these other predictors are strongly correlated with one another. The F test instead simply checks if any single predictor alone has significant correlation with age.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The best subsets algorithm compares the calculated AIC for every possible combination of predictors, then selects the combination with the smallest AIC. Because AIC is used the algorithm will favor larger models than it would if BIC was used and the model will be better fit for prediction.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

D Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 12

Warner Jr, Kevin Scott

544946884 kswjr360 385-522-9143

Test: 637433 Form: A Wadsworth, Spencer Gordon

Entry: 12:47pm 10/20/2026 Checked out:

Employee: Hannah d.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly.
Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034

NS

012

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are referred to as shrinkage models because they use a penalty (found from CV RMSEs) that help us decide which predictors need to be shrunk in their beta hat values or in the case of LASSO, get shrunk to 0. They shrink to simplify & more accurately predict from our model.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\left(\frac{2.4^2}{4}\right)\left(\frac{0.25}{0.75}\right)(1.44)(1/3)$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is our statistic, our estimated coefficient.
β is the parameter, theoretical perfect coefficient.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

We take our X, Y & β matrices* & use some calculus to find our minimum sums of squares error. It produces a 1x1 value that we then use for our coefficient estimate.

*and use the matrix $(Y_i - X\beta)^T(Y - X\beta)$ for our $SSE(\beta)$

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

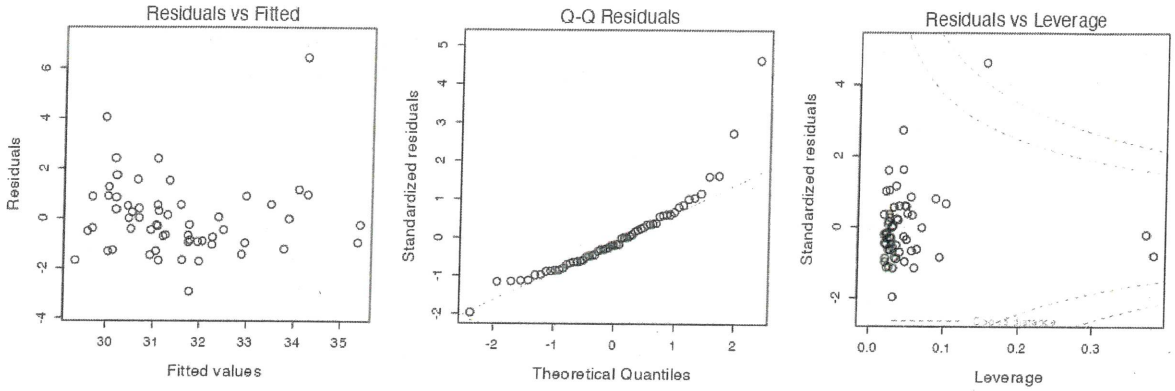

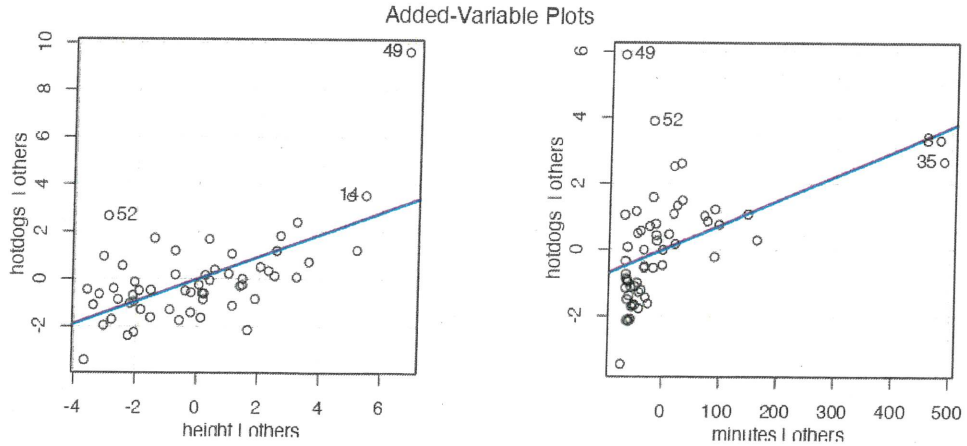

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

estimated value of β_1 is 0.472892. Leaving minutes constant, each one inch increase in the eater's height is associated with a 0.472892 increase in # of hotdogs they can eat in a ten minute interval, on average.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

$\vec{\beta}_1 =$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The av plots show good shape for height but for minutes, the cluster is concerning.

(B) There is independence between each observation

Yes, the eaters are measured individually and one eater's performance should not have too much of an effect on the others

(C) The errors are normally distributed

There is a normality problem as indicated by our Shapiro-Wilk test and extreme curvature at the end of our Q-Q plot.

(D) The errors have constant variance (homoscedasticity)

Yes, most dots don't vary too much as indicated by the plot & the BP-test does not point out too much evidence that the variance depends only on the predictors

(E) There are no unduly influential points

There are 3 values that are really influential*, so that statement is false

*2 with high leverage & 1 with a high standardized residual

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

level-log

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

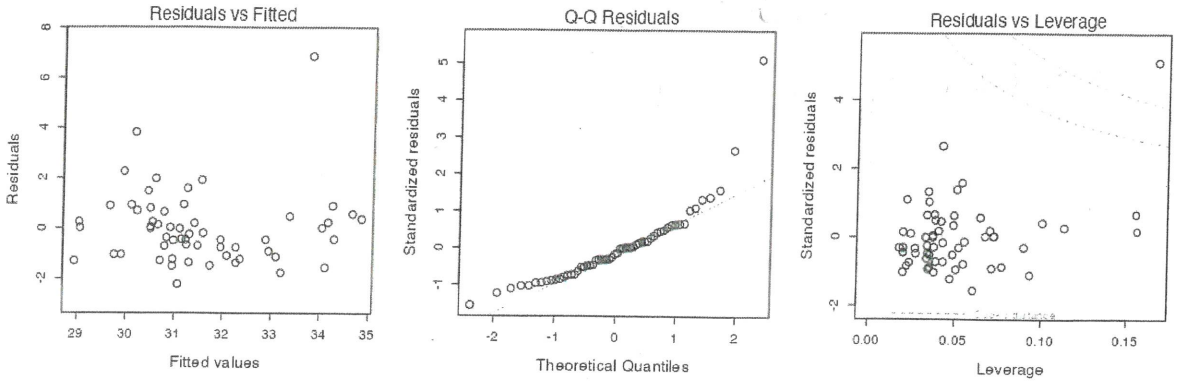

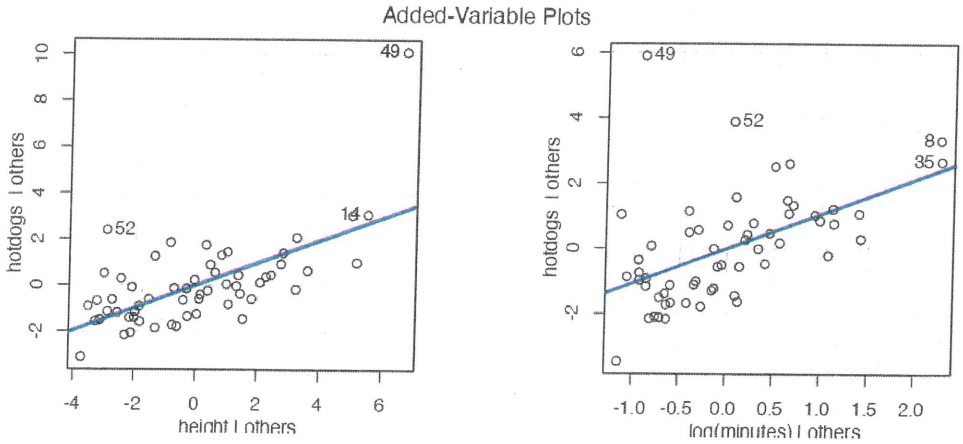

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

Although there is more evidence that error variance is constant in model one, it does not serve as the only deciding factor in determining which model is better. It is important to also look at our residuals versus fitted graph & investigate why we got a lower p-value

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages: CV RMSE is lower, residuals follow a slightly more normal distribution

Disadvantages: Equal variance dropped, more points have much higher leverage, making them more influential

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$

Holding height constant, we observe that a 1-to increase in minutes* is associated with a 0.012 increase in # of hot dogs eaten in the 10 minute interval, on average.

$\ln(1.01) \cdot 1.06702$

*dedicated to training per week

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

hotdawgs = $-6.05589 + 0.148631(68) + \log(79) =$

$-6.05589 + 33.207 + 4.369 = 31.52$ hotdogs eaten in 10 minutes

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error? $n-p-1 = 58-2-1$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$PI = 31.52 \pm 1.485 = (30.035, 33.005)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

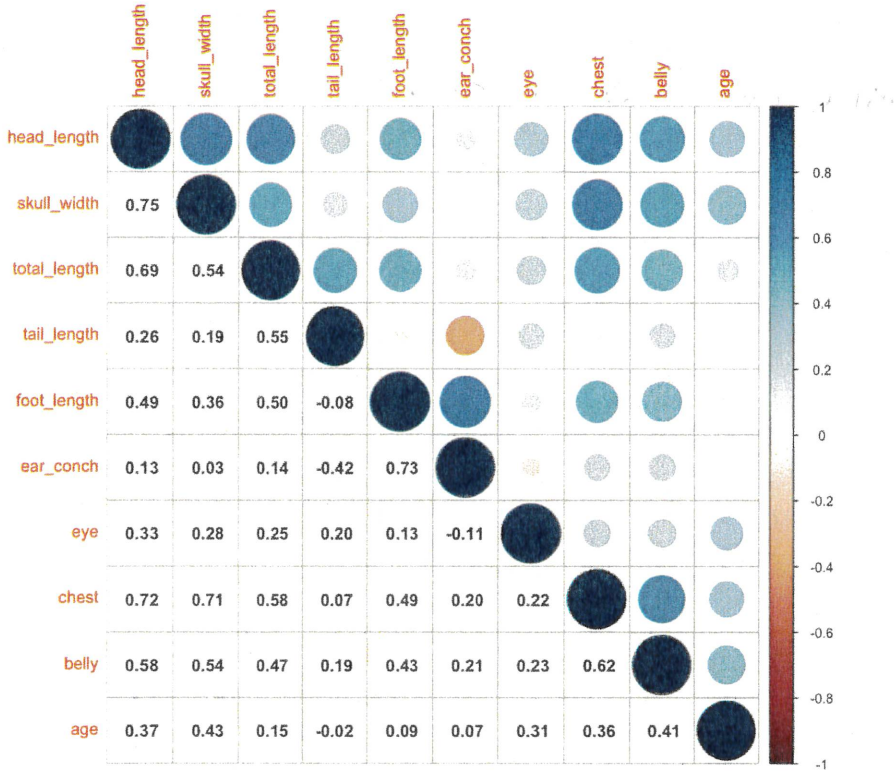

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm SE \cdot 2.01$

$0.10599 \quad 1.6559 \pm 0.2130$

$(1.4429, 1.8689)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

If we were to construct more confidence intervals, 95% of them would contain the true coefficient on belly width between (-0.023, 0.443) centimeters.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

The F-test only shows evidence that at least one of the model coefficients is not 0, which is different from testing each individual predictors' significance to the model. A low p-value on an F-test is not a guarantee for low p-values in t-tests

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets goes through every single possible model, in this case there are 512 different possible models. AIC helps pick which one of these models is the best by telling us which model minimizes our expected out of sample prediction error.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

B Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

C Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 17

Johnson, Aiden

403177698 aj12861 385-225-7505

Test: 637944 Form: A
Entry: 03:17pm 10/20/2026
Employee: Katie W.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

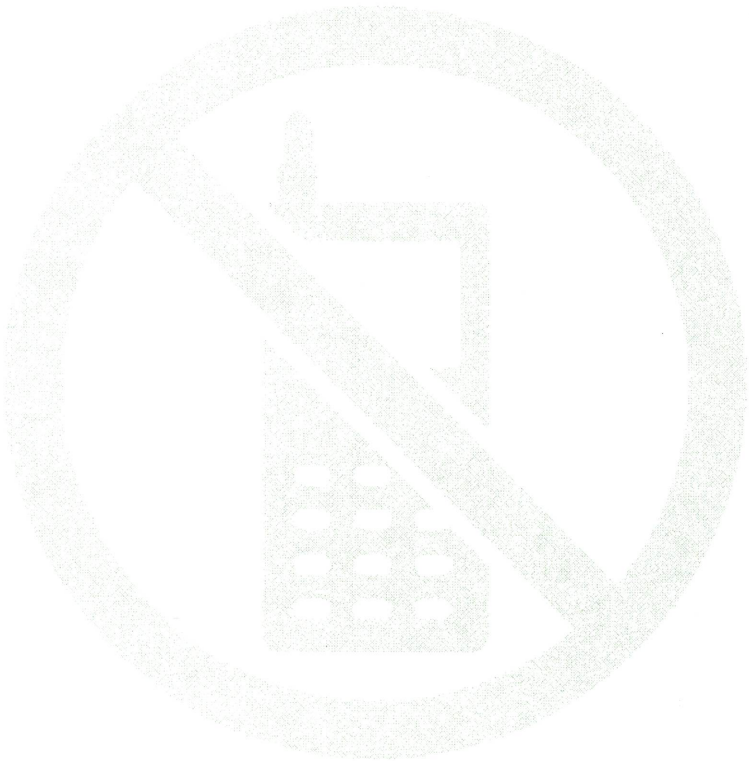

NS

40142034

017

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

LASSO and ridge regression are sometimes referred to as "shrinkage" methods because they shrink the estimated effects of the predictors, the $\hat{\beta}$'s in the model. This is done by adding a penalty term to the SSE: ($\lambda\sum_{j=1}^{p}\beta_j^2$ for Ridge Regression and $\lambda\sum_{j=1}^{p}|\beta_j|$ for LASSO) and finding the values for which the SSE is minimized.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

[(A)] $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β represents the true, exact values of the regression coefficients while $\hat{\beta}$ represents the estimated values of the regression coefficients.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least squares can be used to estimate regression coefficients by finding the values of $\hat{\beta}_i$ that minimize the sum of each estimated residual squared. This process ensures that no two residuals cancel each other out (due to the squaring) and that the estimation is representative of the data (due to the minimization of the total sum).

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

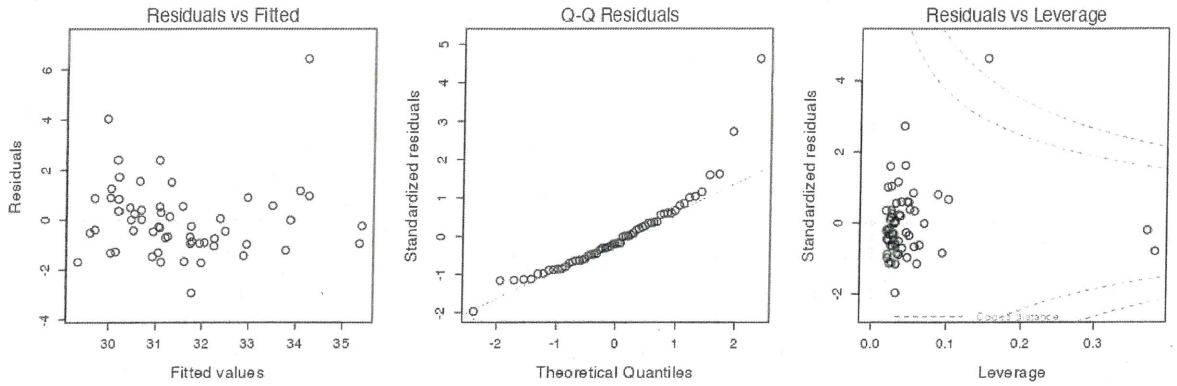

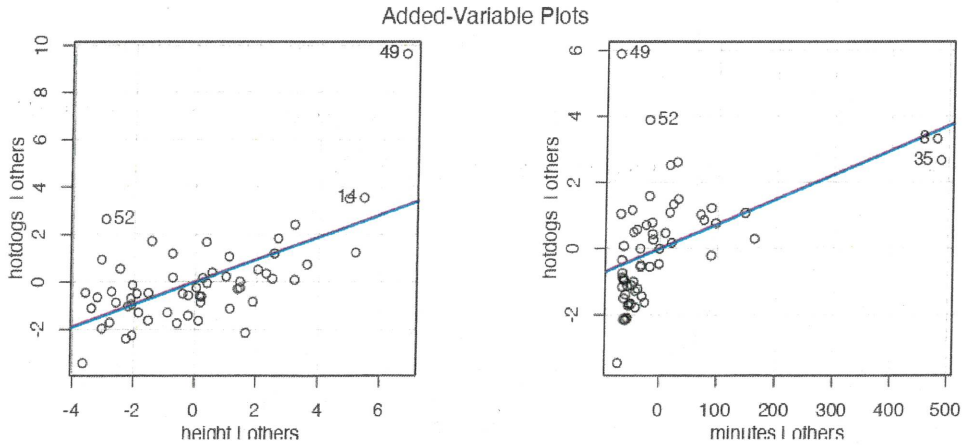

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = 0.472892$. Holding all else constant, a 1-inch increase in contestant height is associated with an average increase of 0.472892 in hot dogs consumed.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

I would say that the relationship between height and consumption is linear but that the relationship between minutes and consumption is not. The AV-plot for height shows a consistently followed trend, but the AV-plot for minutes shows no clear trend.

(B) There is independence between each observation

I would say that the independence assumption is not met because the contestants were not selected from a random sample.

(C) The errors are normally distributed

I would say that the evidence does not suggest that the errors are normally distributed because there is clear deviation on the left and right sides of the Q-Q plot, and the P-value obtained from the Shapiro-wilks test is very small, leading me to believe that the errors are not normally distributed.

(D) The errors have constant variance (homoscedasticity)

The Residuals vs fitted plot shows some slight funneling as fitted value increases, but it isn't anything overly alarming. The P-value from the Breusch-Pagan test also suggests the same conclusion. I feel comfortable with the equal variance assumption.

(E) There are no unduly influential points

The Residuals vs leverage plot shows that there is one very influential point with a significant combination of leverage and residual size, resulting in a large Cooke's distance.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

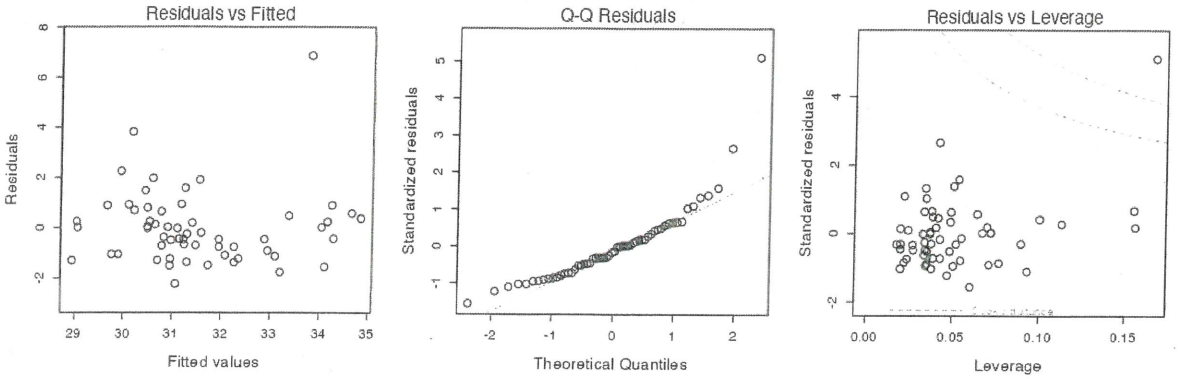

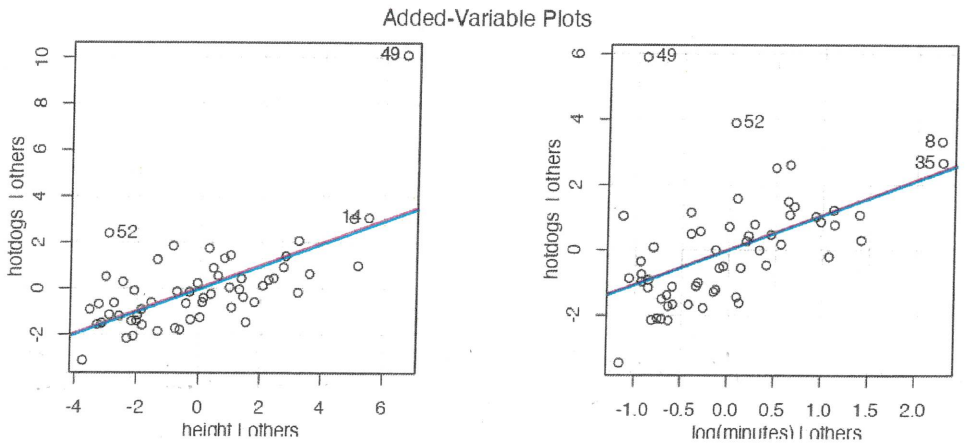

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would respond hesitantly because one test's result does not make or break a model. While the results of the Breusch-Pagan test for the models are important, there are many other aspects of each model that should be compared before concluding which model is better.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

The biggest advantage that model 2 provides over model 1 is that the second predictor has a much clearer linear association with the response variable in model 2 than it does in model 1.
one disadvantage that model 2 has in comparison to model 1 is that the result of the Breusch-Pagan test suggests that the variance in the errors is not constant, leading me to believe that that assumption may be violated in model 2 when it is not in model 1.

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702$. A 1% increase in minutes spent training per week is associated with a $1.06702 \ln(1.01)$ increase in total hot dogs consumed, approximately a $c/100$ increase in hot dogs eaten for a 1% increase in minutes spent training per week.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that $\log(79) = 4.369$. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$HD = 0.46834(68) + 1.06202\ln(79) \approx 37.869$ hot dogs consumed in 10 minutes

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

n-p-1

58-2-1 = 55

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$37.869 \pm 1.485 \cdot 2.004 = (34.893, 40.845)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

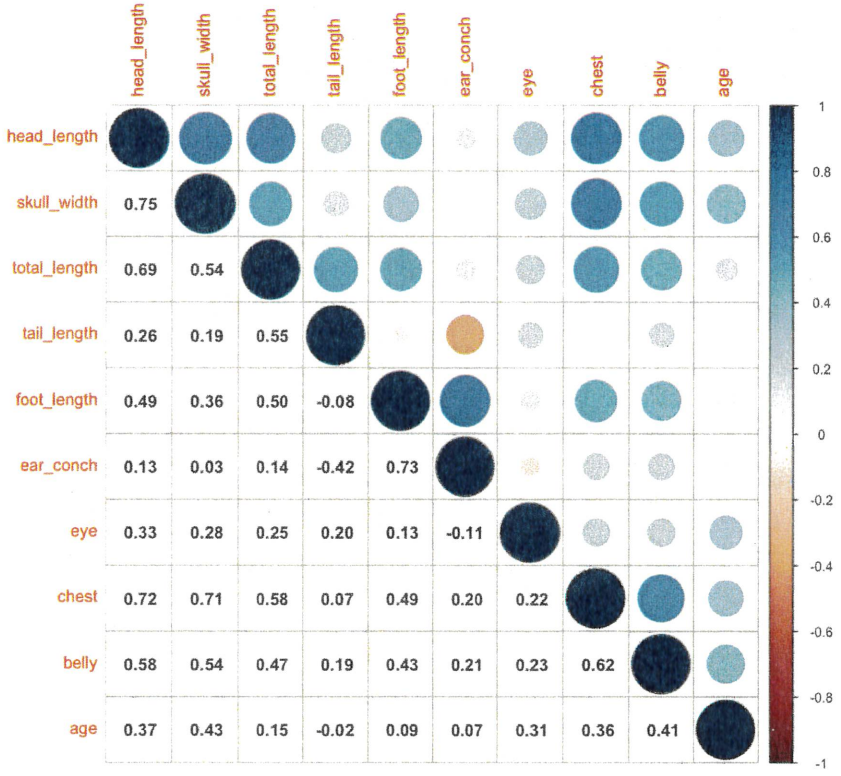

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 0.10599 \cdot 2.01$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We have 95% confidence that the true value of β_{belly} is between -0.023 and 0.443

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

While the association between one individual predictor and the response may not be statistically significant, the association between many predictors and the response is. This is likely because a combination of predictors is needed to show any significant association with the response.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The "best subsets" algorithm using AIC as its criterion performs variable selection by calculating the AIC of every possible model and choosing the best one. AIC favors models where small signals are vital and loosing them would be costly.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 16

Taylor, Brian

314785077 bjazzfan 916-755-9178

Test: 637877 Form: A
Entry: 03:01pm 10/20/2026
Employee: Katie W.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

Brian Taylor

40142034

016

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

The coefficients are being shrunk so that the important and impactful coefficients remain. This accomplished by getting the λ from cross validation and and shrinking the B_j's either by taking the absolute value or squaring it.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is a fixed parameter while $\hat{\beta}$ is an estimate. $\hat{\beta}$ can change depending on the data while β can't.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least Squares uses RSS that minimizes the sum of squares you differentiate, make it equal to 0, which then gives you $X^T X\beta = X^T Y$

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

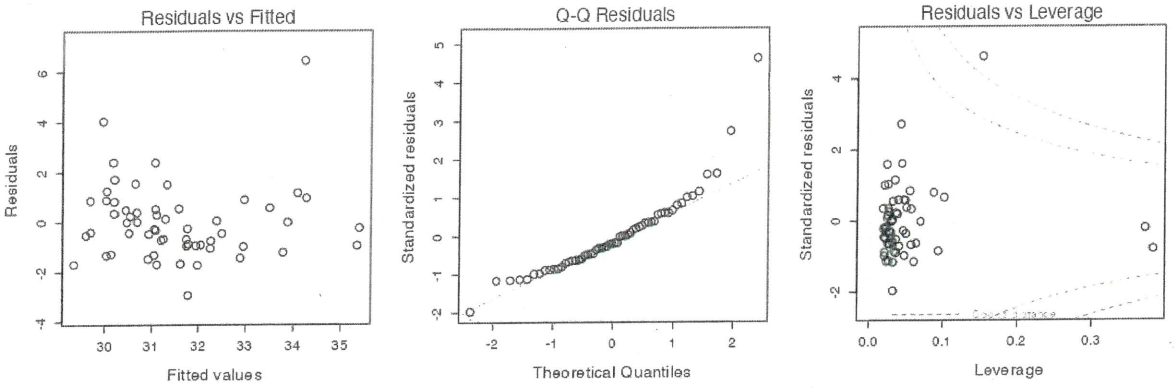

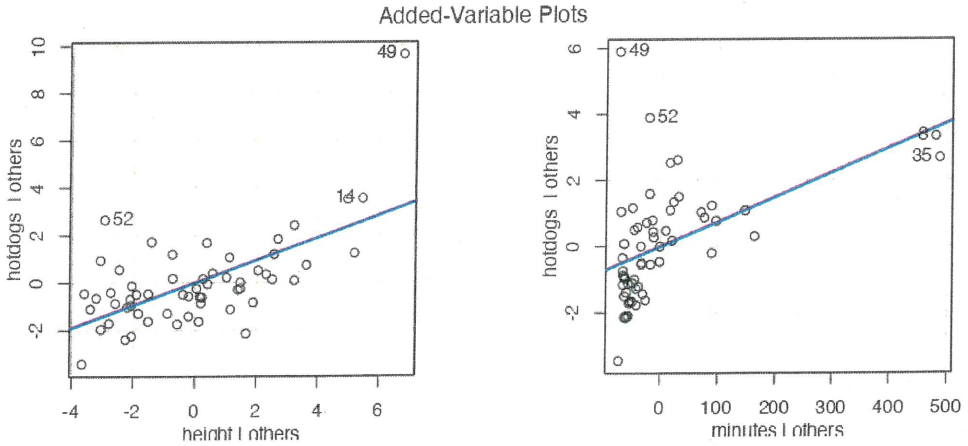

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

Estimate $B_1 = -0.472892$. Holding Minutes constant, for every hot dog eaten, there is associated with a .47 [illegible] increase in height.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

Yes, its linear. There's no major gap in the residual vs fitted plot and the added-variable plots ASSUMPTION satisfied.

(B) There is independence between each observation

It doesn't specifically say that it is. Each eater in the contest was recorded. convience samples aren't the best for independence. ASSUMPTION is not satisfied, but you can still use this in some circumstances

(C) The errors are normally distributed

There is conform- since the tails of the Q-Q plot go up and some more than others. ASSUMPTION is not met our predictions will not be as accurate.

(D) The errors have constant variance (homoscedasticity)

ASSUMPTION is largely met. variance seems pretty constant throughout the residual vs fitted plot.

(E) There are no unduly influential points

There are some points notably point 49 is way further out on both added variable plots. There's one point that's flagged for cooks Distance and I assume it's that one. It's worth investigating

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

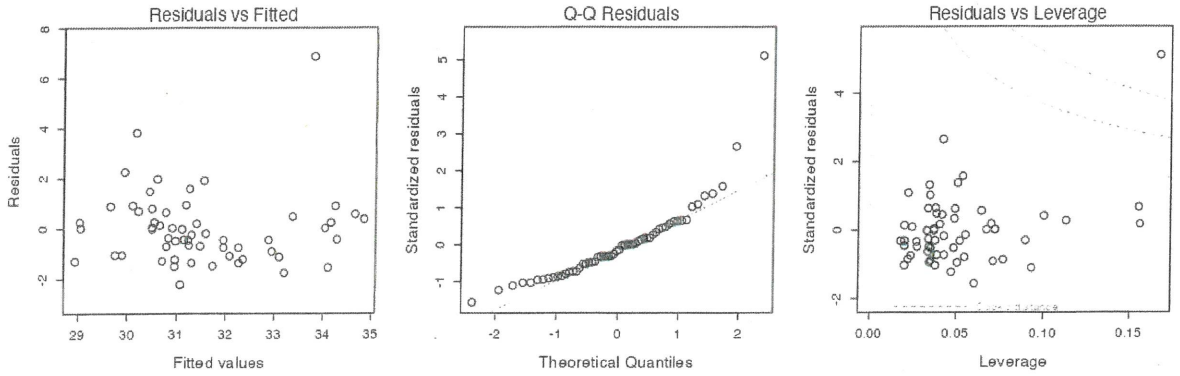

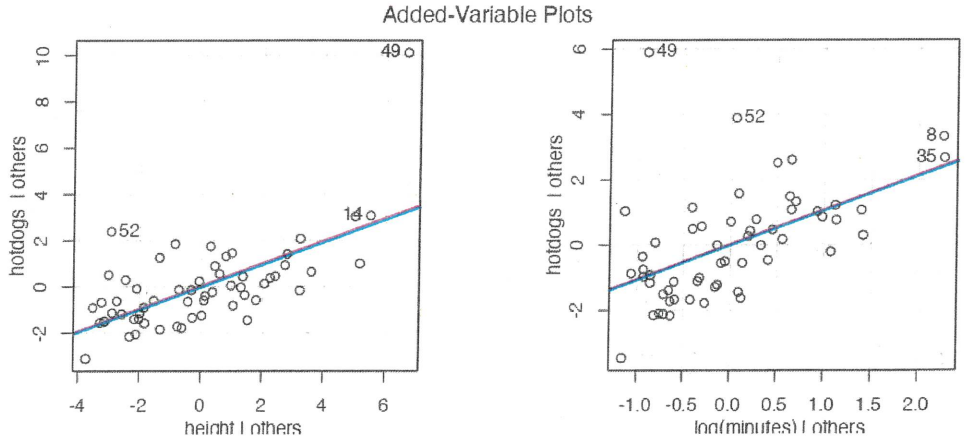

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

That's not necessarily true. models with a smaller sample size can be influenced by small or even normal deviations.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

The main advantage is that there's a smaller CV-RMSE than in model 1. However, the Shapiro-wilk and Breush-Pagan Pvalues are more significant. Since points have more meaningful leverage and now the equal variance assumption is more questionable.

Question 16: State the estimated value of β_2 and provide an interpretation in context

Holding height constant, with a 1% increase in hotdogs eaten gives (1.06702 (log (1+1/100) ≈ .00461 increase in in Minutes spent per week in preparation

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-4884168 + 4.369 = 37.57612$

33.20712

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$37.57612 \pm 2.004(1.485)$ $\pm .204$

2.850901944

$(34.725, 40.43)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

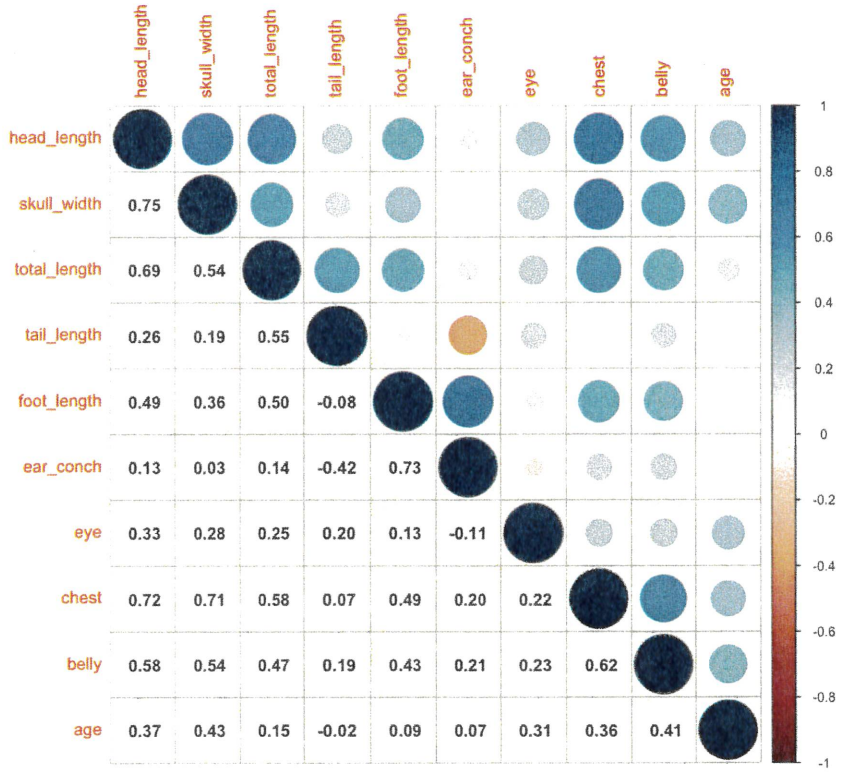

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$.1659 \pm 2.01(.1059)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident the true population coefficient of Belly girth on age is (-.023, .443)

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

When we add predictors, the p-value of an F test gets smaller because its testing if at least one is not 0.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The same split mean models are compared on the identical best subsets and paired differences remove the split-to-split noise which helps AIC perform variable selection. It works to find the model that best explains all the necessary important variables.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

B Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

C Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 15

Jorgensen, Zachary

180263508 zajorgen 801-376-9605

Test: 637871 Form: A Wadsworth, Spencer Gordon

Entry: 03:00pm 10/20/2026 Checked out:

Employee: Jacob M.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

015

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS ✗

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are shrinkage methods because they reduce/shrink the coefficients in the model. They do this by adding penalties to the SSE calculation, $SSE + \lambda \Sigma B^2$ for ridge and $SSE + \lambda \Sigma |B|$ for LASSO

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\frac{2.4^2}{4}\left(\frac{.25}{.75}\right) = .48$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is the coefficient calculated based on sample data, β is the true population coefficient. Its theoretical

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

You want to find the $\hat{\beta}$ values that minimize the SSE. To do this, start w/ the SSE equation and expand it. You can take a derivative in terms of β to find the critical value. Since the SSE equation is concave up, this critical value is a minimum

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

outside cook's D

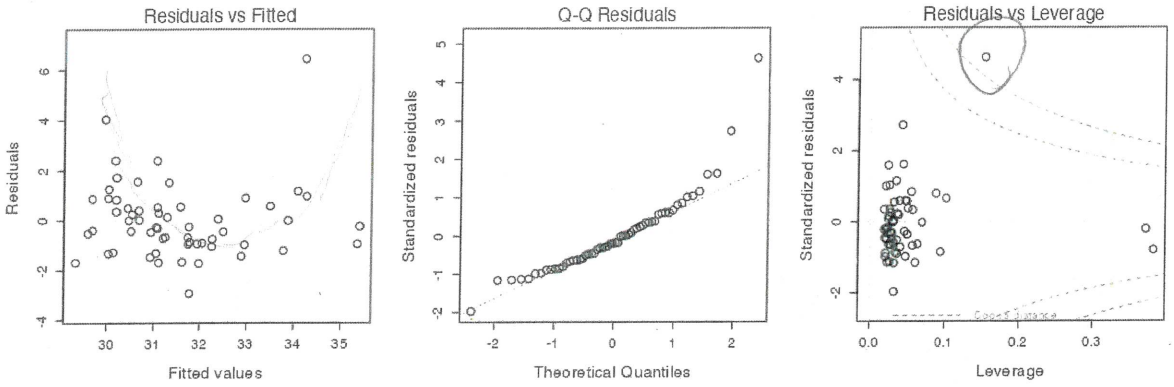

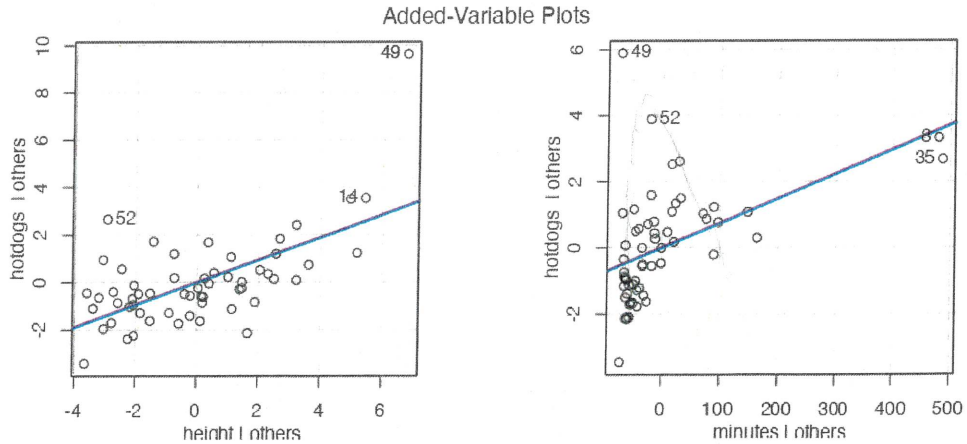

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = .472892$ holding all other predictors constant, we predict that for each additional inch tall a contestant is, they will be able to eat .473 more hotdogs on average

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

This assumption isn't well-satisfied. the resid vs fitted plot shows curvature and the 2nd AV plot is clearly non-linear

(B) There is independence between each observation

If the contestants watched each other eat the hotdogs for 10 mins then I'd say these observations may not be independent, since someone's competitiveness could inflate the values. since I don't know if that was the case, I'd say independence likely isn't violated here

(C) The errors are normally distributed

The QQ plot shows substantial skew and the Shap-Wilk test has a small p-value so I'd say normality is likely violated

(D) The errors have constant variance (homoscedasticity)

The resid v leverage plot may taper off a little on the right side but its generally fairly well spread. Due to this and a high p-value for the BP-test, I'd say equal variance likely isn't violated

(E) There are no unduly influential points

I'd say there is at least 1. It looks like observation 49 based on the Resid v leverage and av plots but its hard to say for sure

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

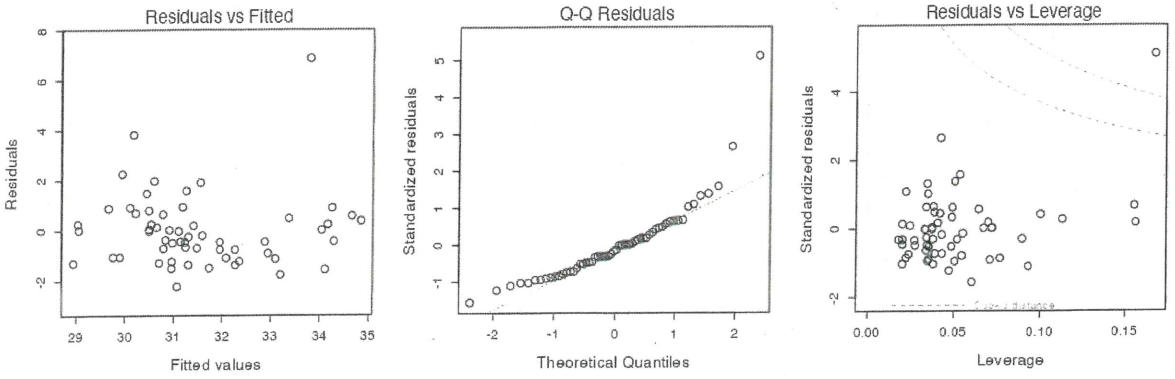

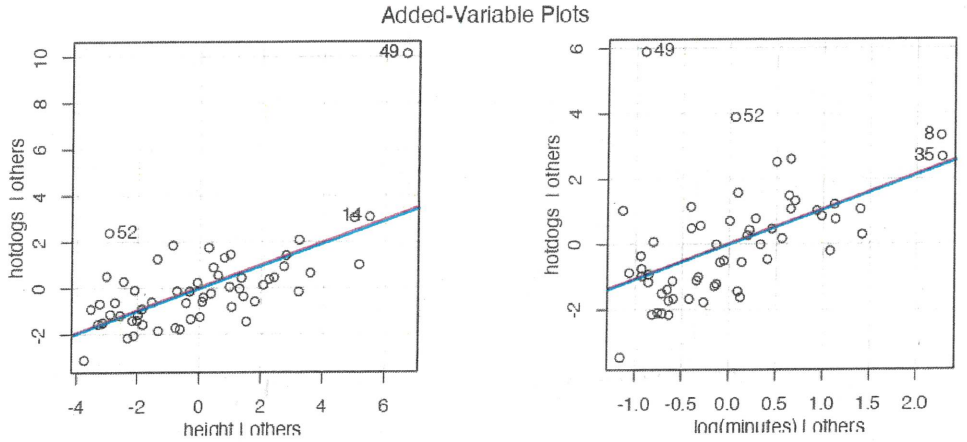

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

First, BP test results aren't hard lines, they are just helpful. the other plots aren't crazy so equal variance doesn't seem like a major issue in either model.

Second, there are way more important diagnostics than that one. the 2nd model has a better CV error meaning model 2 is likely better for prediction than model 1

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Pros: Better CV RMSE, possibly is lesser violation of normality assumption but it still has issues

Cons: Harder to interpret, closer to a violation of equal variance

Question 16: State the estimated value of β_2 and provide an interpretation in context

Holding all other predictors constant, we expect that for a 1.0° increase in minutes spent training, a contestant per week Will eat .011 more hot dogs in 10mins on average

$\beta_2 = 1.06702$

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 + .48834(68) + 1.06702(4.369) = 31.813$

hotdogs in 10 mins

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

$n-p-1 \quad n-3 = 55$

$n = 58$

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31.813 \pm 2.004(1.485)$

Pred. Int. $= (28.837, 34.789)$

hotdogs in 10 mins

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

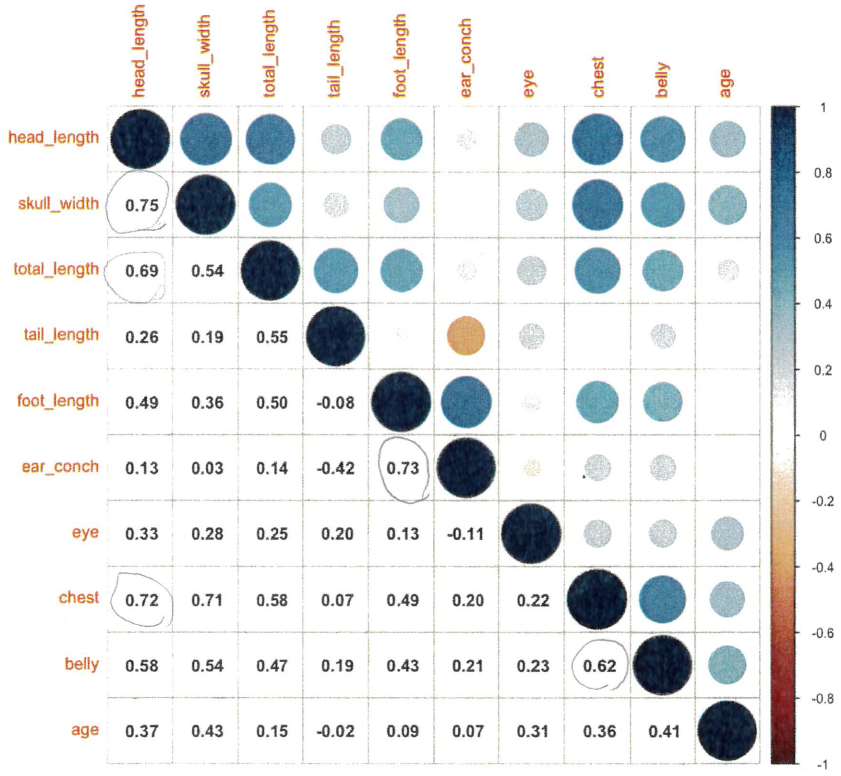

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$.16559 \pm 2.01(.10599)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that the true population β_{belly} value, when holding other predictors constant, is between $(-0.023, 0.443)$. this means there is not good evidence to believe that the β_{belly} is not $= 0$, or belly girth doesn't seem to be a good predictor of age accounting for other predictors

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

This can happen in a data set w/ multicollinearity, while this model doesn't show extreme VIF values, there a several variables w/ relatively high correlations. Perhaps there is less independent information here than we think, making it harder to find good predictors

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets is an algorithm that checks all possible models. 1st it looks at all one-predictor models, then two predictors and so on until its tested every possible combo. It calculates the AIC for each model and finds the exact combo of predictors that gives the lowest AIC.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM