STAT 330-001

40142034 14

Allen, Brighton Michael

367982208 allenb03 801-560-1424

Test: 637744 Form: A
Entry: 02:18pm 10/20/2026
Employee: Esperanza C.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

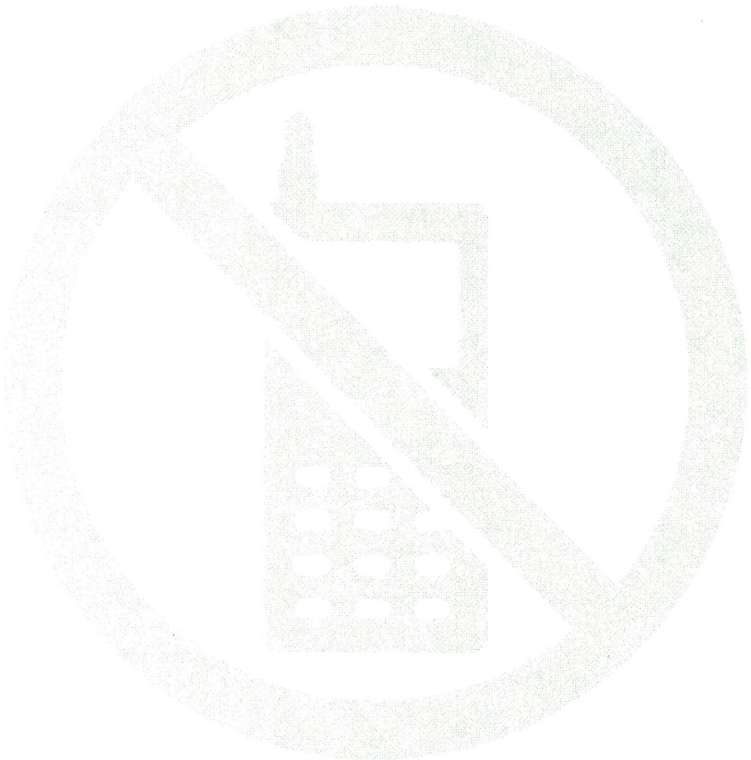

NS

40142034

014

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are-allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets.** Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: **(Circle one answer)** Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

$y \sim x_1 \ldots x_n$

$x_{n+1} \sim x_1 \ldots x_n$

(A) TRUE

(B) FALSE

Question 2: **(TRUE/FALSE)** If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: **(Circle one answer)** When adding an additional variable to a linear regression model, the model's R^2 ...

$R^2 = \frac{SSR}{SST}$

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: **(Circle all that apply)** Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

The coefficients are being shrunk. This happens by applying a penalty (λ) to either the sum of the coefficients squared (Ridge) or the sum of the absolute value of the coefficients (Lasso), all in addition to the least squares estimate of the coefficients. λ is multiplied by the "Ridge" sum or "Lasso" sum, and this will make the coefficients smaller.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$D_i = \frac{(2.4)^2}{4}\left(\frac{0.25}{0.75}\right) = 1.44\,(1/3) = 0.48$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$(X^TX)^{-1}X^Ty$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β is the true population coefficient (which we don't know), while $\hat{\beta}$ is our best estimate of the coefficient based on our data.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

It starts with our data (observations). Given our observations, we can organize each piece into matrices such as $[Y]$, $[X\beta]$, $[\varepsilon]$ etc. using initial estimates for the coefficients. We then algebraically organize SSE as a function of β (with the matrices) and minimize in order to estimate the coefficients with the least squared error.

observations → $\min[SSE(\beta)]$ → $[\hat{\beta}]$

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

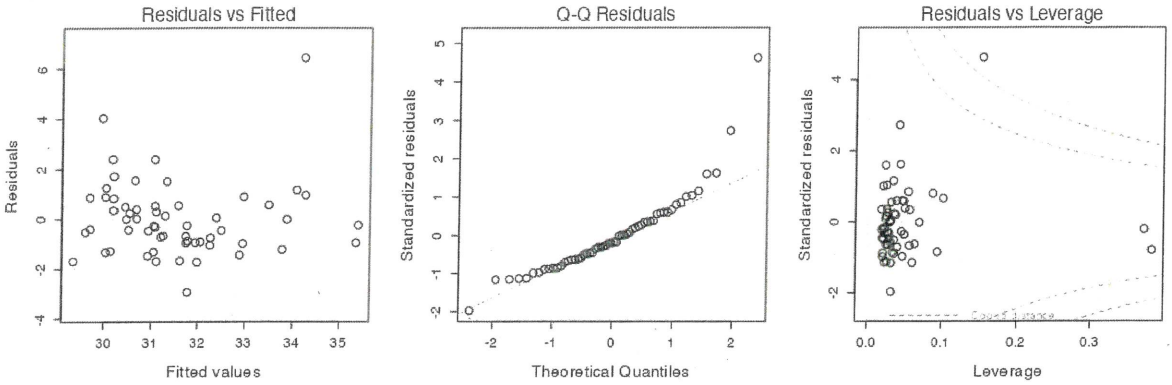

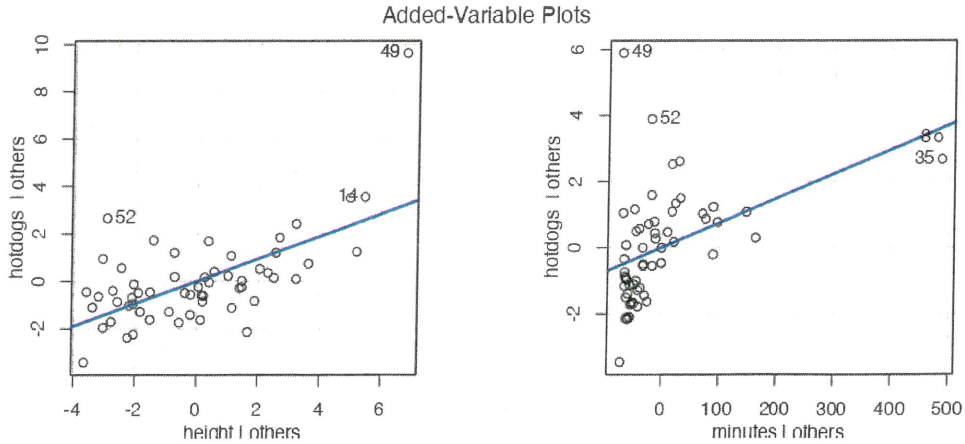

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

The estimated value of β_1 is 0.472892. This means that if all else is held constant, a one-unit increase in height will correspond to an approximately 0.47 increase hot dogs consumed in 10 minutes on average.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

$0.007417 \pm 3.963\,(0.001871)$

$0.00741477\overline{3}$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

The residuals vs. fitted plot shows no strong pattern, which implies linearity. However, the added-variable plot for height vs. the response shows a considerable positive correlation, so that's worth investigating.

(B) There is independence between each observation

We can confidently assume independence because the competitors competed individually, therefore one person's performance shouldn't have directly influenced another person's performance.

(C) The errors are normally distributed

The quantiles in the qq-plot follow in fitted line closely, which implies normality. However, the fact that both tails curve upward at the ends is worth investigating. The low Shapiro-wilk p-value also supports evidence that residuals are normally distributed.

(D) The errors have constant variance (homoscedasticity)

The residuals vs. fitted plot doesn't hold to any strong pattern, which means variance is fairly consistent throughout the range of observations. However, it may be worth investigating further because the Breusch-Pagan p-value isn't significantly low.

(E) There are no unduly influential points

There appear to be at least two points that are worth investigating for unusually high leverage.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

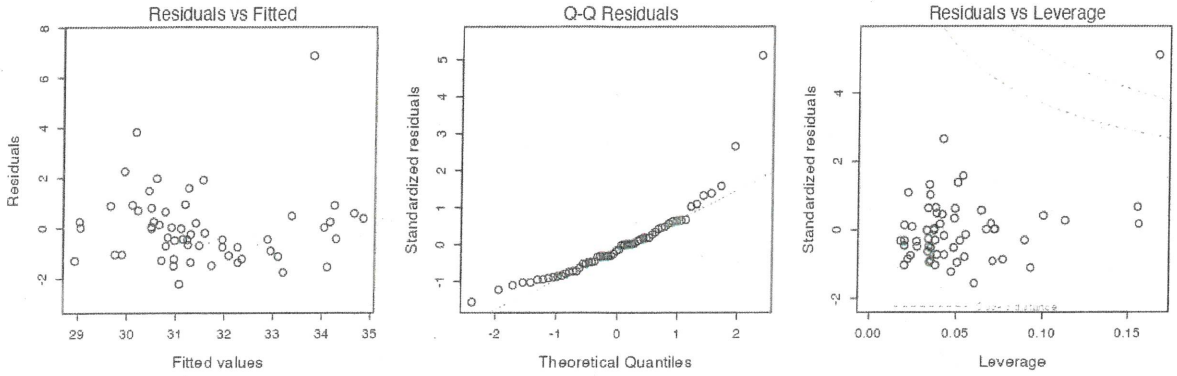

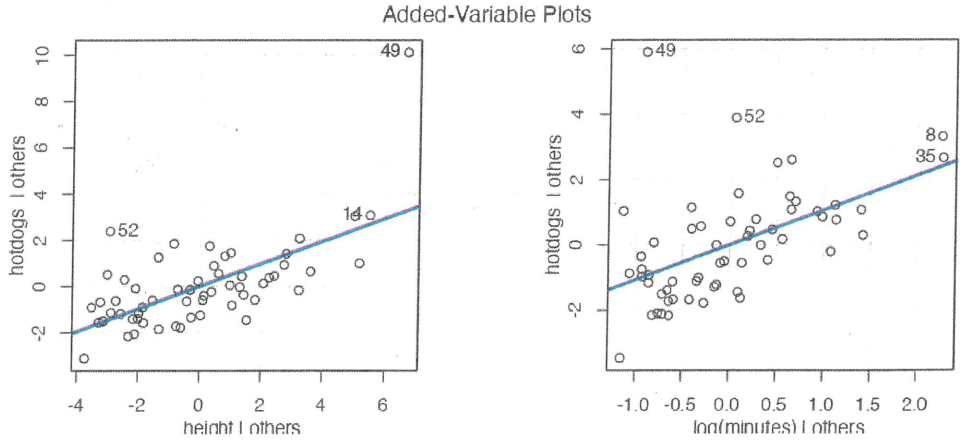

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would say we can't base our entire conclusion off of one p-value. In fact the smaller p-value for model (2) is closer to what we want anyway. We can't rely on this one test alone. We ought to compare how the models fit each of the assumptions, taking into account the use of each evaluation tool, in order to make the most informed conclusion.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

In terms of advantages, I like that model (2) has lower CV RMSE, and lower Breusch-pagan and shapiro-wilk p-values. This could indicate that it's better at predicting out-of-sample, and it better fits normality and equal variance assumptions. It also has eliminated pre-existing high-leverage points. Disadvantages, however, include larger coefficients, linearity in added-variable plots, and slightly less linear qq-plot observations. These could imply more variance, more multicollinearity, and residuals that are less normally distributed.

Question 16: State the estimated value of β_2 and provide an interpretation in context

The estimate of β_2 is 1.06702. This means that, holding all else constant, a 1% increase in minutes of training corresponds to an approximately 1.067 increase in hot dogs consumed in 10 minutes on average.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that $\log(79) = 4.369$. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$\text{hotdogs}_{\text{friend}} = -6.05589 + 0.48834(68) + 1.06702(4.369) + \epsilon$

$= 31.82 \approx 32$ hot dogs

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

$n - p - 1$

$58 - 2 - 1 = 55$

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31.82 \pm$

$2.004(0.204) = 0.4088$

$2.004(1.485) = 2.9759$

$(31.4112, 34.7959)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

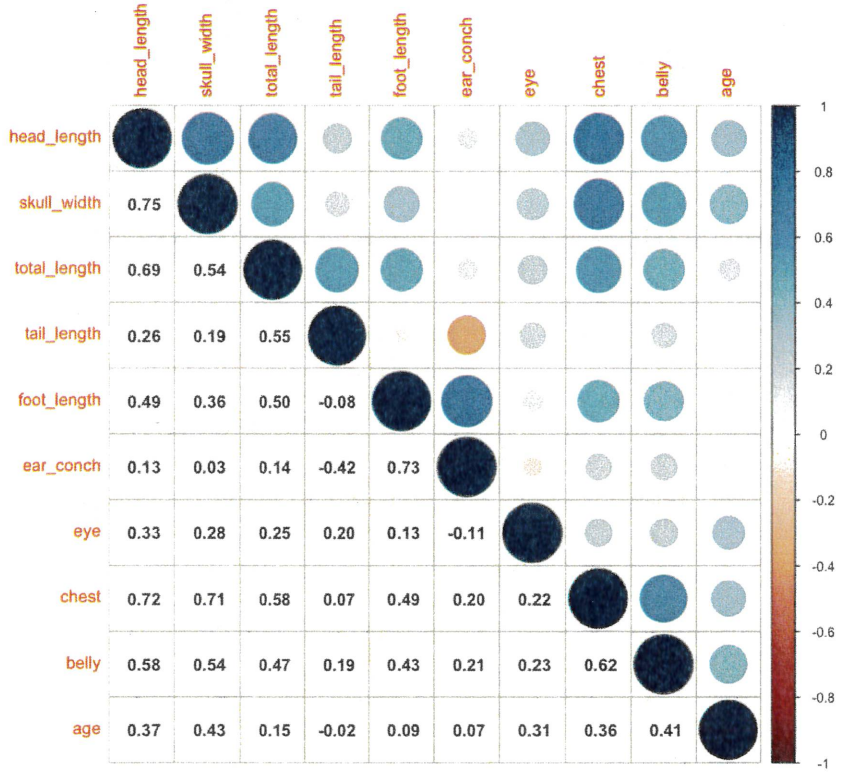

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$0.16559 \pm 2.01\,(0.10599)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that the population coefficient for belly width is as a predictor of age is between (-0.023, 0.443).

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

The F-test is telling us how well the covariates work together to predict age. While the individuals coefficients may not strongly predict age, as a conglomeration they are more effective at estimating the response.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

The 'Best Subsets' algorithm essentially takes every possible model combination of covariates and evaluates them based on their AIC metric. We select a model that has the lowest AIC.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

A Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 13

Bird, Ryan

211001693 rybird33 801-455-0697

Test: 637632 Form: A Wadsworth, Spencer Gordon

Entry: 01:53pm 10/20/2026 Checked out:

Employee: Ebony L.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

013

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are-allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

In Ridge and LASSO, We introduce Bias into our model to minimize the variance. This Bias - variance trade-off allows our model predictions to be more accurate

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$1.44\left(\frac{1}{3}\right) = .48$ $\quad 4/n = .04$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\hat{\beta} = (X^T X)^{-1} X^T y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\beta \rightarrow$ population parameter, an unknown that we try to estimate with $\hat{\beta}$.

$\hat{\beta} \rightarrow$ Estimate parameter using our sample data which we use to help fit our model

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

Least Squares takes the sum of our residuals by taking it's gradient and setting it equal to zero. Then we find our normal equation which gives us our $\hat{\beta}$ estimate. This estimate would be the minimum $\hat{\beta}$ that we then use for our model

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

n=58

Hot dogs → 10 min consumption
height → contenstants height
min → time training

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

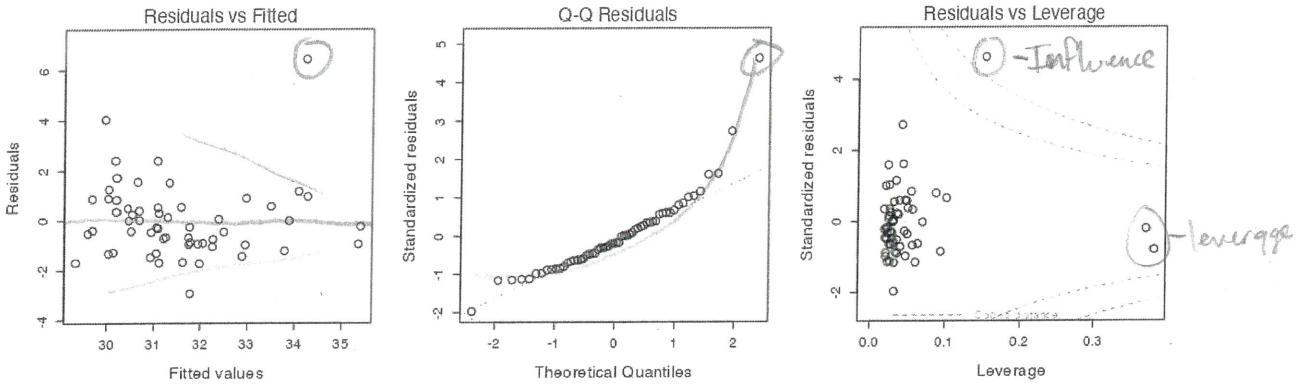

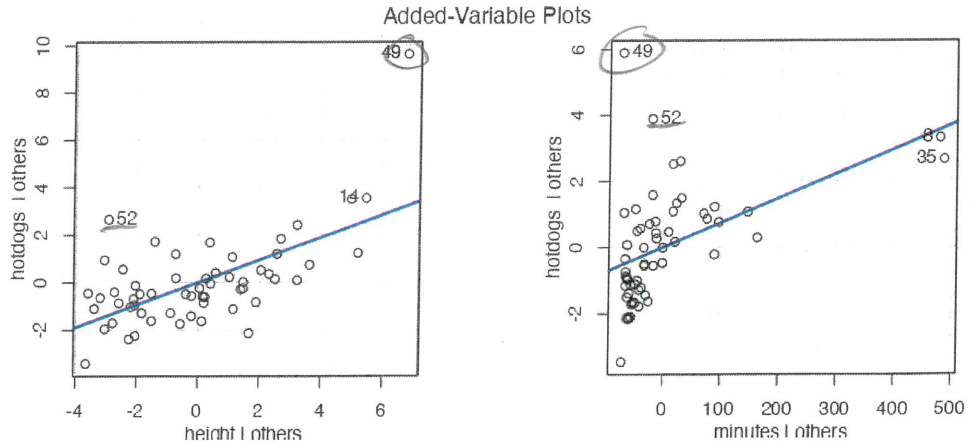

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = .472892$

Holding minutes constant, an increase in an inch of height is associated with an increase of .4728 Hot dogs in 10 minute window on average.

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

For the most part, there is a positive linear relationship in the Added-variable plots. However, points 49 & 52 would need investigation before I would assume linearity

(B) There is independence between each observation

From what I understand, there are no inherit flags with how the data was collected. I would like to see more randomness noted. However, from what we know we can move forward but would like to know more

(C) The errors are normally distributed

There is a clear violation found in our Q-Q Plot. We would need to investegate a few points, That being said, we also know that our n is larger. Additionally our Breush-pagan test says that they are approximately normal. However, wouldnt hurt to look

(D) The errors have constant variance (homoscedasticity)

This is the clearest violation to me. The residual vs. fitted plots are funnel shaped additionally some outliers. Would definetly need to investegate and under stand more. Plus our Shapiro-Wilk test is also telling us that our variance is not equal

(E) There are no unduly influential points

I would say that there is influential points. in our residual vs leverage, we see one point well beyond the cooks distance line. Would need to investegate that point more

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

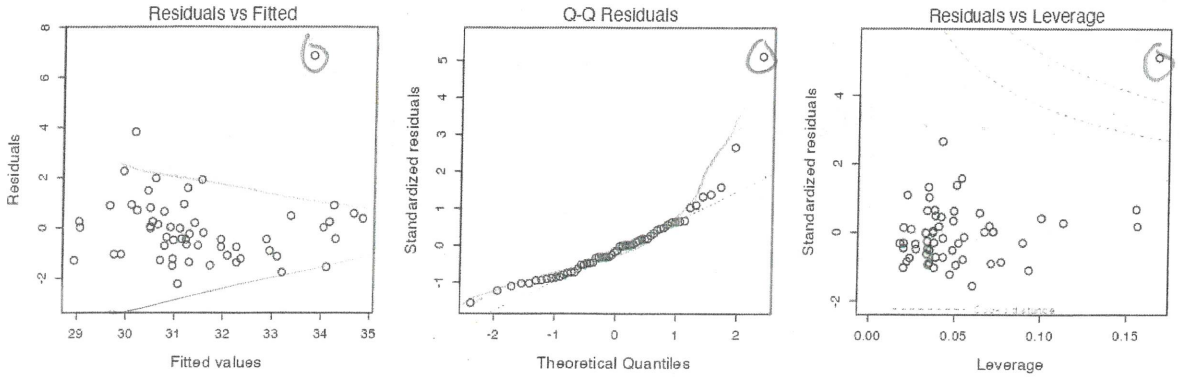

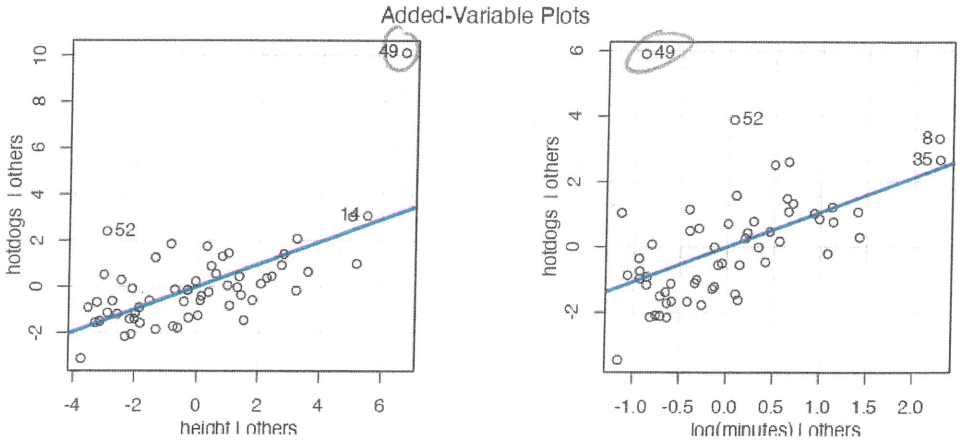

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

better

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would not rely on p-values to make assumptions. When we look at the QQplot for each, they are almost identical. p-values don't give you the full picture and shouldn't make harsh assumptions from them. I don't know if there is a superior model as LINE assumptions still seem violated.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advantages: Our leverage has gone down with some points, Our model fits linearity better, lower CV RMSE

Disadvantages: LINE assumptions still violated, Breush-Pagan p-value got worse, Intercept estimate is much lower, didn't help the model to perform better

Question 16: State the estimated value of β_2 and provide an interpretation in context

$\hat{\beta}_2 = 1.06702 \approx \frac{1.067}{100} = .01$

Holding height constant, a 1% increase in the number of minutes spent training, is associated with an .0106 increase in the number of Hotdogs consumed in 10 minutes on average

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 + .48834(68) + 1.06702 \log(79)$

$-6.05589 + 33.20712 + 4.66181 = 31.813$ Hot Dogs in 10 min

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

SE conf SE pred

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31.813 \pm 2.004(1.485)$

$31.813 + 2.97594 = 34.78894$

$31.813 - 2.97594 = 28.83706$

95% pred $= (28.84, 34.79)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

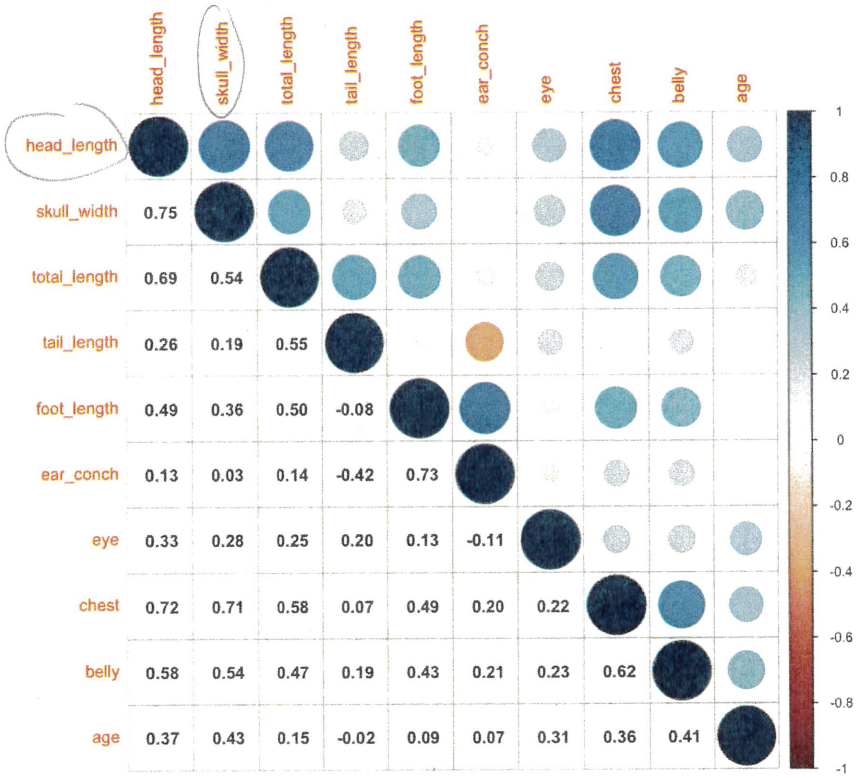

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

$\beta_{jp} \pm t^* SE_{\beta_j}$

$.16559 \pm 2.01(.10549)$

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident, that the average coefficient for belly will be between -0.023 & 0.443

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

the model as a whole is causing our test to find significance. If none of the coefficients are found to produce significant results. I would check for overfitting as well as multicollinearity. Both would help in this situation

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

Best subsets will make every model possible to see what model would be best. We would use AIC as the metric to determine the best algorith. AIC specifically differs from other metrics because of the penalty that it has for factors.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

Fit 1 Model
(A) prior to the for loop X loop → split, fit, predict, compute RMSE

(B) within the for loop

(C) after the for loop – Average & Compare

(D) never in the CV process

C Average RMSE across all test-training splits and report

C Average the predictions and record their value

D Choose the model with the largest RMSE

A Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

D Compute each model's AIC and BIC

D Fit each model using the testing data

B Fit each model using the training data

I Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 70

Ortega Avila, Luca

551253443 lucajose 801-648-1253

Test: 641689 Form: A
Entry: 08:25pm 10/21/2026
Employee: Julia L.

Wadsworth, Spencer Gordon
Checked out:

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

NS

40142034

070

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets.** Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

Because they approach the coefficients of the model to zero (ridge) when in lasso case it can actually take the coefficient to zero, eliminating them from the model.

So in datasets with redundant covariantes "shrunk" is how the influence of them

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\frac{(2.4)^2}{3+1}\left(\frac{0.25}{1-0.25}\right)$ 1.44 0.04 0.48

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

estimator $\hat{\beta} = (X^T X \hat{\beta}) = X^T Y$.

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

$\hat{\beta}$ is the estimation we try to make in the model while β is the actual fixed value of the parameter.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

$n = 58$

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

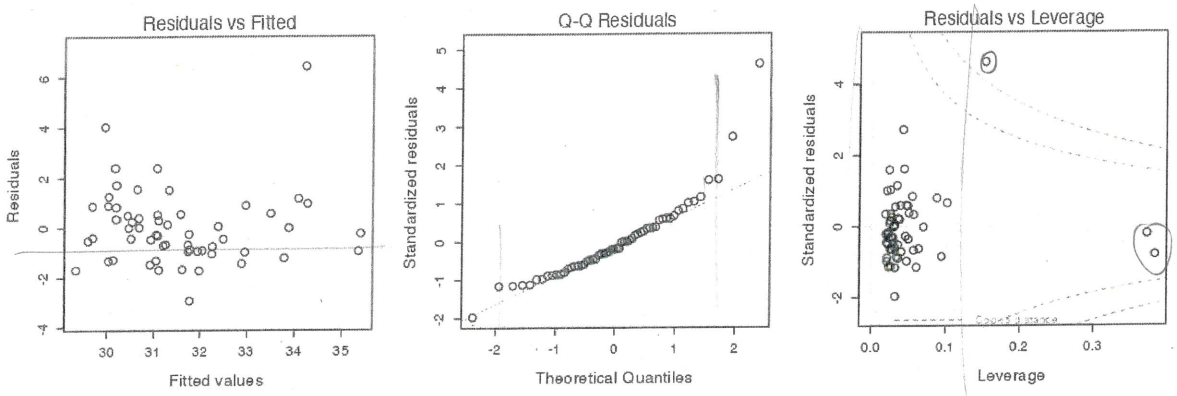

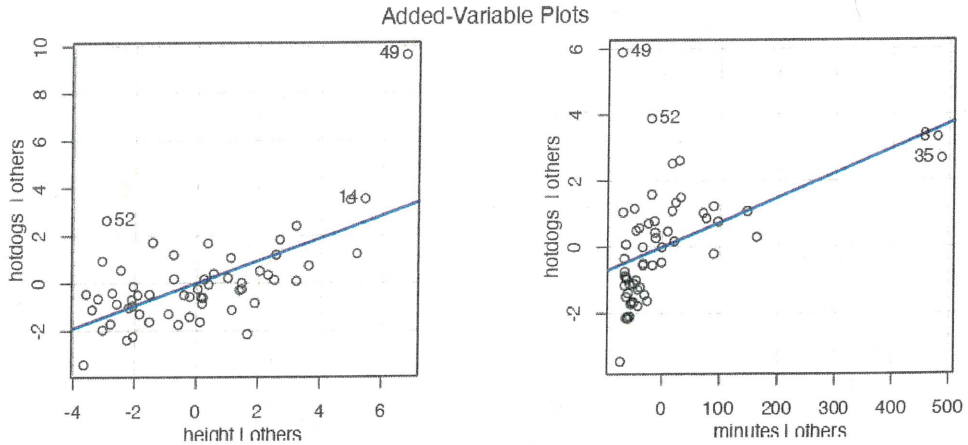

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

the estimated value of β_1 is 0.472892
since hush and Hotdog have a pretty positive relationship.
thats the influence or coff of β_1 in the whole model

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

the residual vs Fitted plot shaw a pretty horizontal
Tendency at the 0 value

(B) There is independence between each observation

Yes! Based on the design of the study I would state that
the heigh and time: training per week are pretty independent!

(C) The errors are normally distributed

No really, the Q-Q plot stow an asending line but
around smaller than -2 and bigger than ~1.5 it lost the linear shape

(D) The errors have constant variance (homoscedasticity)

Yes with 3 minor data dat. out of the umbral.

(E) There are no unduly influential points

Yes there are! I make them in the
Residual vs leverage graph

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

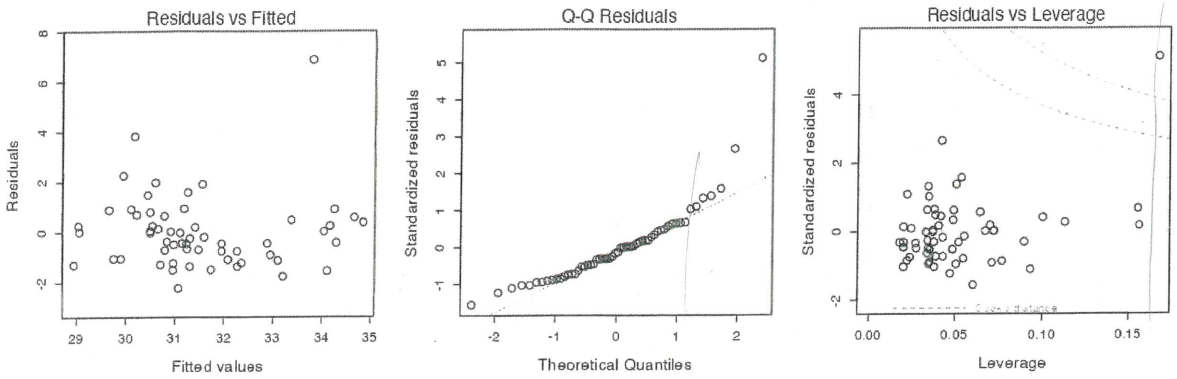

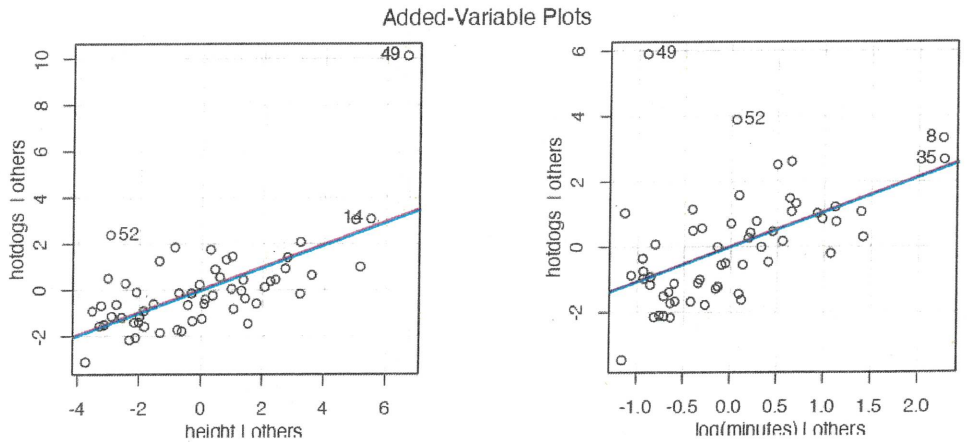

Coefficient	Estimate	Std. Error	t value	Pr($>\|t\|$)
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

No! model 2 simplify the study and observation of the model. Since the p-value of model 2 0.028 / we could past some normality but the influence of minutes b_2 could be Recognize better!

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Advanteges are the study of b_2 log(minutes) is clearly more simple and easy to work with the Data. and Recognize patters

Disadvantage the transformation didn't help with the overall Normality model how you can see in the new Q-Q Residual.

Question 16: State the estimated value of β_2 and provide an interpretation in context

the estimated value of B_2 is 1.06702 in the context that how much influence in the whole model, how many minutes they throw on how many Hotdog would eat in 10 min.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$\hat{Hotdogs} = -6.05589 + 0.48834(68) + 1.06702(4.369)$

model 2 ≈ 31.8

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55 $\quad n - p - 1 = 58 - 1 - 1 = 56$

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$\hat{\beta}_1 \pm t \times SE(\hat{\beta}_1)$

$31.8 \pm (2.004 \times 0.30294)$

$, 32.408)$

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

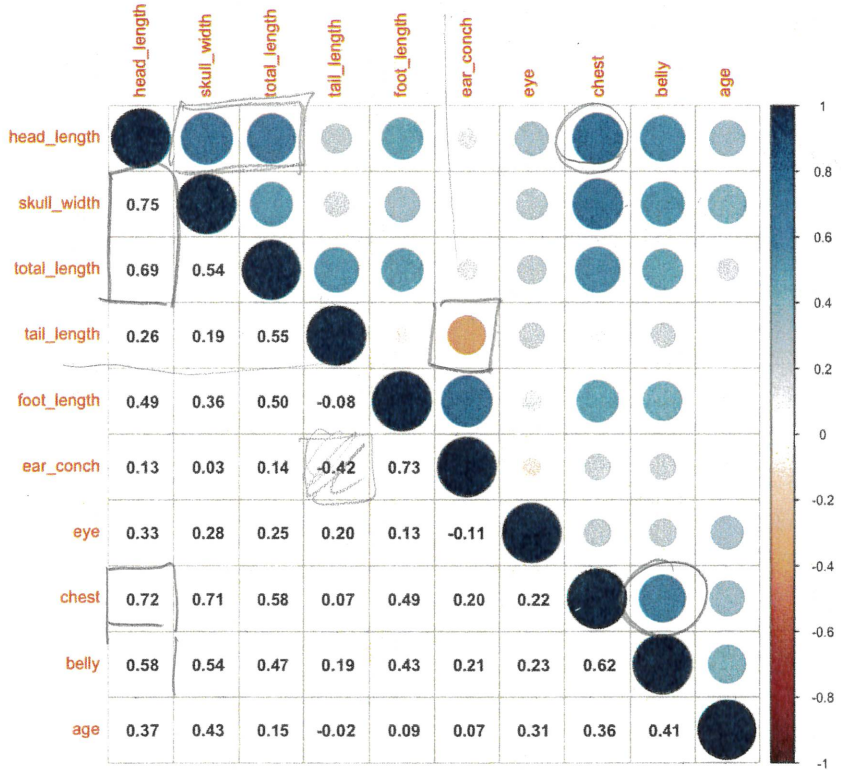

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

0.16559 ± t·

0.16559 ± 2.01 ×

0.16559 ± 2.01 × 0.10599

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that if we repeat this same research the sample will fall around the interval (−0.023, 0.443)

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

I think all coefficient measure similar values around "individual size" any of them will be significant in the study. Since some of them overlap or could have a high redundance.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

$$AIC = n \ln\left(\frac{SSE}{n}\right) + 2p^*$$

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

C Average the predictions and record their value

D Choose the model with the largest RMSE

B Choose the model with the smallest RMSE

A Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

B Compute each model's AIC and BIC

C Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

C Refit each model using the full data set and compare the coefficient estimates to the true values

A Split the data into test and training sets

A Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 68

Canning, Brayden

728498995 bcannin1 385-630-7497

Test: 641631 Form: A Wadsworth, Spencer Gordon

Entry: 08:01pm 10/21/2026 Checked out:

Employee: Addie T.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034
068

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

They are shrinkage methods because they shrink some predictors down to zero (Lasso) and or all of them close to zero (Ridge). This is accomplished by multiplying λL_R or λL_{1SE}

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

5.76/4 = 1.44 .25/.75 = .33 = .4752 > .04

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

Avg < 4/n

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\vec{\beta} = (X^T X)^{-1} X^T Y$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

β comes from a population

$\vec{\beta}$ comes from samples of a population

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

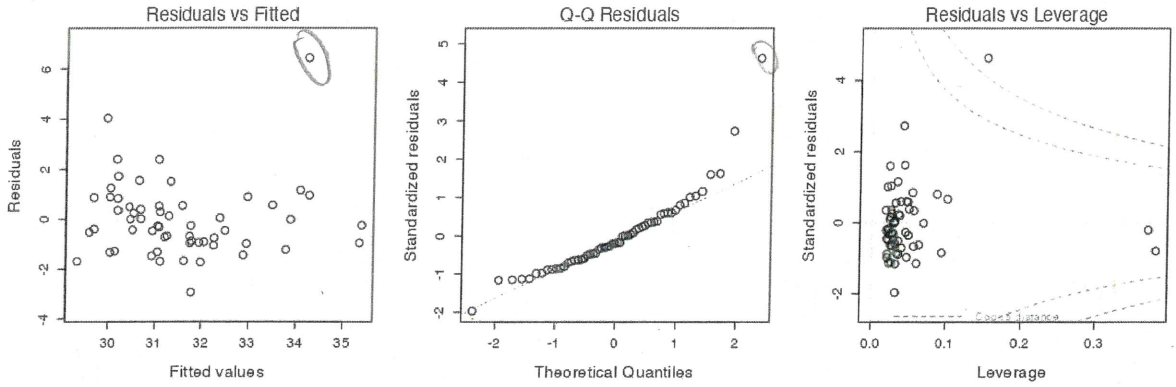

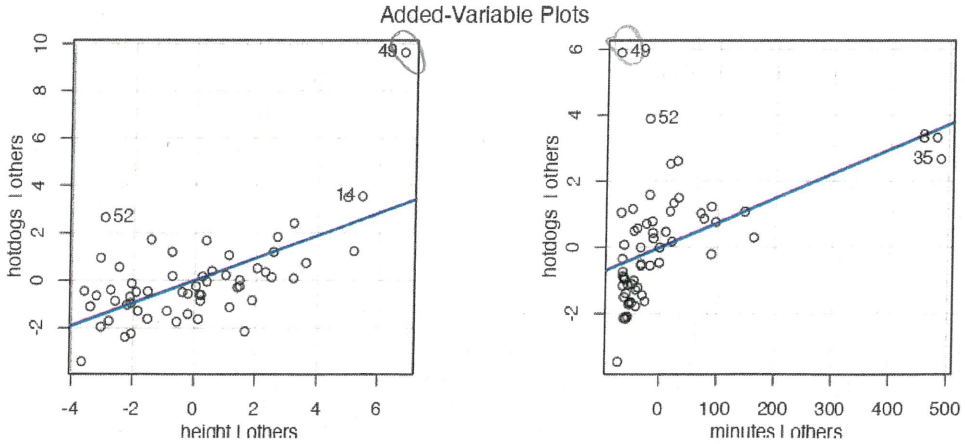

Coefficient	Estimate	Std. Error	t value	Pr($>\|t\|$)
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\hat{\beta}_1 = 0.4729$

holding all else constant, every inch increased in a competitor is associated with a 0.4729 increase in dogs consumed

Question 12: **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

linear, a line at zero would represent the residuals without any bending or curves

(B) There is independence between each observation

does it talk about random sampling just overall was collective - skeptical

(C) The errors are normally distributed

~~yes, line would be flat on the residual vs fitted graph~~

no, big curve going left on the Q-Q

(D) The errors have constant variance (homoscedasticity)

yes, line would be flat on the residual vs fitted graph,

(E) There are no unduly influential points

There is one point really out there on all of the graphs - could pull data

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

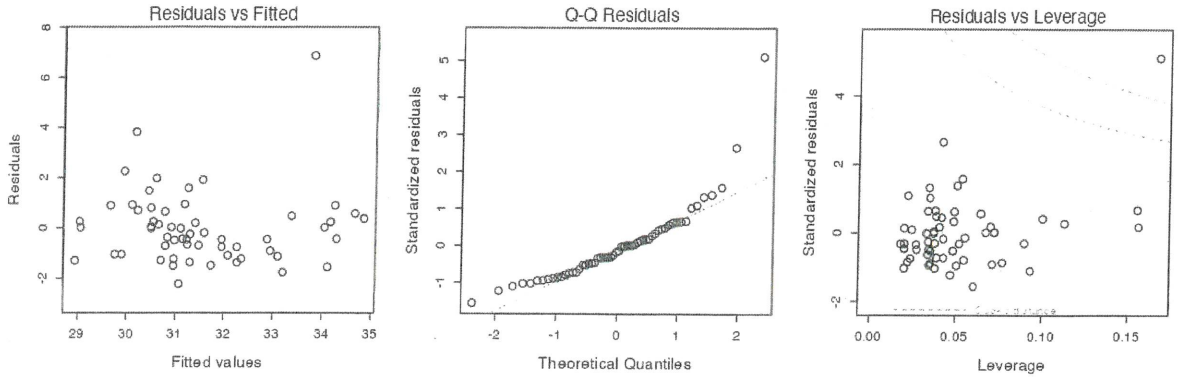

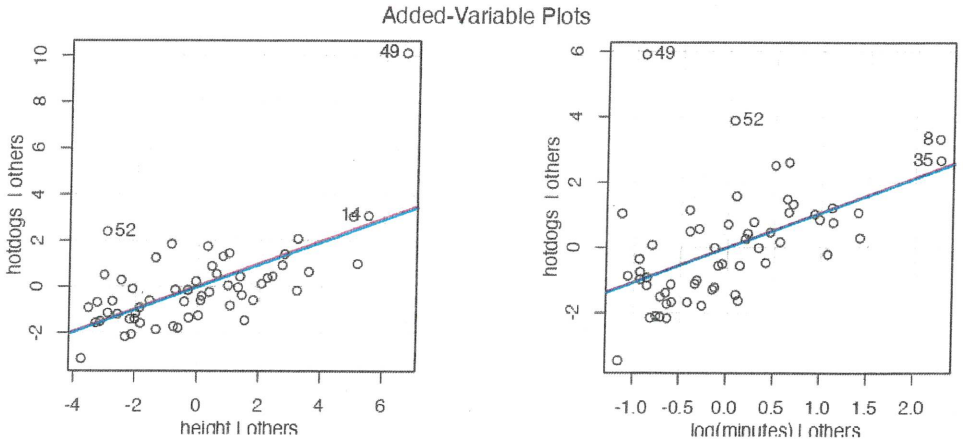

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

The smaller $p(0.028) < .5$ → violates Breusch-Pagan

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

Question 16: State the estimated value of β_2 and provide an interpretation in context

$B_2 = 1.06702$

Holding all else constant, a ____ percent increase in log(minutes) would be associated with ____ increase (amount of) in days consumed

2

MEAN

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$-6.05589 + 0.43834(68) + 1.06702(4.369) + 5.43620$

37.24 dogs consumed

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55

(B) 57

(C) 56

(D) 48

Question 19: You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

We are 95% confident that the amount of dogs Dr Wadsworth can consume is in the interval (34.234, 40.246)

$37.24 \pm 2.004\sqrt{0.204^2 + 1.485^2}$

$0.0416 \cdot 2.205$

$\sqrt{2.2466}$

$2.004\ (1.5)$

3.006

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.22377    7.34105  -1.801   0.0778 .
head_length   0.07077    0.10120   0.699   0.4877
skull_width   0.16559    0.10599   1.562   0.1247
total_length -0.04752    0.10167  -0.467   0.6423
tail_length  -0.05049    0.17030  -0.296   0.7681
foot_length  -0.21499    0.10733  -2.003   0.0507 .
ear_conch     0.15730    0.10231   1.538   0.1306
eye           0.43736    0.22549   1.940   0.0582 .
chest         0.01883    0.18154   0.104   0.9178
belly         0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

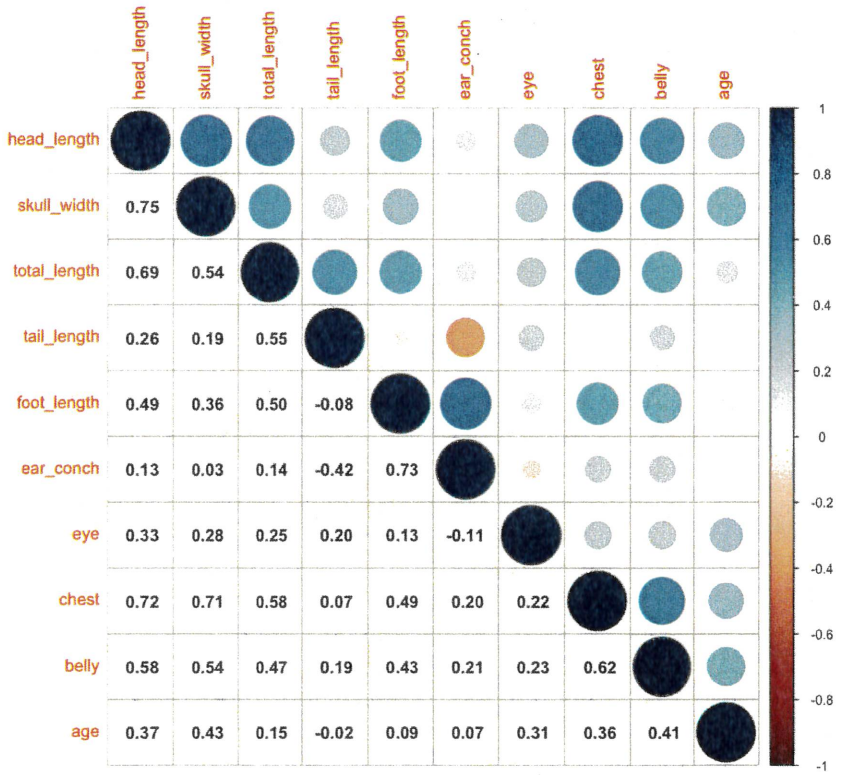

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

CI: $\hat{\beta} \pm t^* \cdot SE$ 0.16659 ± 2.01

0.16659 ± 2.01

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

we are 95% confident that the true mean of belly width is in the interval (-0.023, 0.443)

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

Multicollinearity meaning that though we can see the model is effective from the F-test, we aren't sure which of the coefficients is actually driving it

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

great question... umm... usually... algorithm means trying to find the best model and they use metrics to score those algorithms to use

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

C Average the predictions and record their value

A Choose the model with the largest RMSE

A Choose the model with the smallest RMSE

B Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

A Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

D Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

D Use the models fitted to the training data to get predictions for the test data

END OF EXAM

STAT 330-001

40142034 66

Prusse, Anna Elizabeth

341818956 aeprusse 303-717-1533

Test: 641576 Form: A Wadsworth, Spencer Gordon

Entry: 07:46pm 10/21/2026 Checked out:

Employee: Ila B.

SPECIAL INSTRUCTIONS

TC Calc
Notes: 3x5
Dictionary
Take a deep breath. Be careful to read instructions thoroughly. Answer questions completely.

Please keep this sheet **VISIBLE** on your desk at all times. When you are finished, *staple this sheet* to the front of your test booklet and return all materials to the Grading Table. You may use the space below for notes or scratch paper.

PLEASE TURN OFF YOUR CELL PHONE

40142034
D66

STAT 330 - Exam 1

Section 001 (Wadsworth), Section 002 (White)
October 16th - 21st, 2026

Instructions

- You may not refer to your notes, other people, or other resources. You can, however, use one 3 inch by 5 inch note card that you create (writing on both sides is allowed).
- Do not discuss any aspect of this exam with anyone except a Stat 330 TA or instructor until scores are posted.
- You are-allowed to use a calculator for this exam.
- In Parts II and III, **make sure to clearly show your work for any calculations that you make and to give your answers in context of the data sets**. Place your final answer inside the box, if one is provided.
- There are 24 questions. Make sure you answer all of them.
- There is extra scratch paper at the end of the exam.

PART I

Question 1: (Circle one answer) Partial regression (or added-variable) plots illustrate the effect of adding an additional variable to the model after accounting for the variables that are already included in the model.

(A) TRUE

(B) FALSE

Question 2: (TRUE/FALSE) If the relationship between x and y is nonlinear, a simple linear regression model may still be feasible after transforming the data.

- TRUE
- FALSE

Question 3: (Circle one answer) When adding an additional variable to a linear regression model, the model's R^2 ...

(A) always increases or remains unchanged.

(B) always decreases or remains unchanged.

(C) may increase or decrease depending of the additional variable.

Question 4: (Circle all that apply) Which of the following statements describe possible effects of multicollinearity?

(A) It can cause the linearity assumption to be violated.

(B) It can cause estimated coefficients to have signs opposite those of their true population coefficients.

(C) It can result in larger standard errors, and thus greater uncertainty in estimates.

(D) It can cause the coefficient of determination (R^2) to become negative.

Question 5: (Circle all that apply) Which of the following are true statements about multiple linear regression (MLR)?

(A) BIC will usually favor a model with more predictors than the AIC criterion when performing model selection.

(B) Cook's distance can be used to identify influential observations.

(C) Transformations of the x's or y may help alleviate problems with linearity, equal variance, or normality.

(D) The least-squares estimate of β can be obtained by fitting a simple linear regression for each x_k only.

(E) The assumptions on the error variables (i.e., the ϵ_i) in a multiple linear regression model are exactly the same as they are for simple linear regression.

Question 6: (Circle one answer) What is the primary benefit of using cross-validation (CV)?

(A) CV reduces variance in the sampling distribution of regression coefficients by splitting the training and test data.

(B) CV gives a more realistic representation of model performance by predicting observations that were not used to fit the model.

(C) CV ensures model assumptions are met which eliminates a time-consuming step of the modeling process.

(D) CV is the most efficient way to calculate the shrunken coefficient estimates in LASSO and ridge regression models.

Question 7: (Circle all that apply) In datasets with high multicollinearity and/or redundant covariates, what advantages do LASSO and/or ridge regression have in comparison to the traditional ordinary least squares (OLS) approach to model fitting?

(A) LASSO can set coefficients to be exactly equal to zero, and thus serves as an automated approach to variable selection

(B) Their coefficient estimates are less biased than those of OLS

(C) They produce more accurate confidence intervals and hypothesis tests

(D) Ridge regression often handles multicollinearity well

(E) They can handle datasets where the number of predictors exceeds the number of observations

(F) They can take advantage of the bias-variance tradeoff to produce more stable and accurate predictions

Question 8: Why are LASSO and ridge regression sometimes referred to as "shrinkage" methods? Specifically, what is being shrunk and how is this accomplished?

lasso and ridge regression are sometimes referred to as "shrinkage" methods because that's exactly what they do! they shrink the coefficients by either dropping some of the predictors or setting the coefficients to be close or equal to zero.

Question 9: A multiple linear regression model with $p = 3$ predictors is fit using $n = 100$ observations. Observation i has standardized residual $r_i = 2.4$ and leverage $h_{ii} = 0.25$.

Use

$$D_i = \frac{r_i^2}{p+1}\left(\frac{h_{ii}}{1-h_{ii}}\right).$$

$\left(\frac{(2.4)^2}{4}\right)\left(\frac{0.25}{1-0.25}\right) = (1.44)(0.333) = 0.48$

(Circle one answer) Which conclusion is most appropriate?

(A) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(B) $D_i = 0.48$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(C) $D_i = 0.48$. The observation is not influential because its leverage is less than 0.50.

(D) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation is potentially influential and should be investigated.

(E) $D_i = 0.2$. Because this exceeds $4/n = 0.04$, the observation must be deleted before reporting the model.

(F) $D_i = 0.2$. The observation is not influential because its leverage is less than 0.50.

(G) Cook's distance cannot be calculated without knowing the observation's original residual.

Question 10:

(A) What is the formula for $\widehat{\beta}$?

$\widehat{\beta} = (X^TX)^{-1}X^TY$

(B) What is the conceptual difference between $\widehat{\beta}$ and β?

the difference between $\widehat{\beta}$ and β is that β is the true coefficent of the population versus $\widehat{\beta}$ which is the estimated coefficent from the data in our sample. we use $\widehat{\beta}$ to find out about β.

(C) Explain how least squares can be used to estimate regression coefficients. You do not need to derive anything, but you should be clear about how this is done conceptually.

least squares can be used to estimate regression coefficients through matrix multiplication. the inverse of X^TX equals the sums of the products of the predictors. X^TY is the sums of products of each predictor with y. when multiplied together, we get our estimated regression coefficients.

PART II

Some consider competitive eaters to be among the most inspiring individuals in modern times. We are interested in knowing what factors relate to a competitive eater's ability to consume. In a recent state fair, a contest was held where contestants competed to see who could eat the most hotdogs in 10 minutes. Data were collected on 58 amateur competitive eaters, most of whom were from the 4 largest counties in the state. For each eater, the total number of hotdogs consumed in a single 10 minute sitting was recorded. Each contestant's height in inches was also recorded, and each contestant was asked "How many minutes per week do you spend training for a contest?" The variables are `hotdogs`, `height`, and `minutes`. Model (1) is used for analysis. Use the plots and model fit output below to answer questions 11-13.

Model 1

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \texttt{minutes}_i + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{1}$$

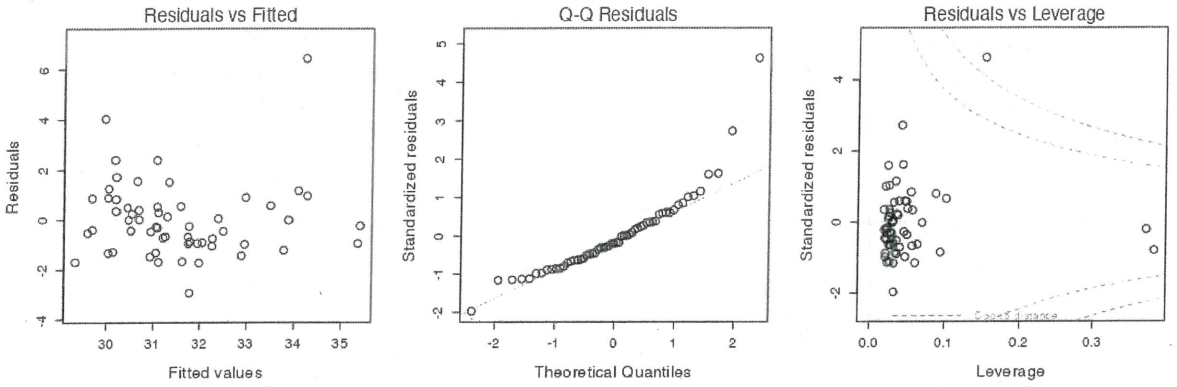

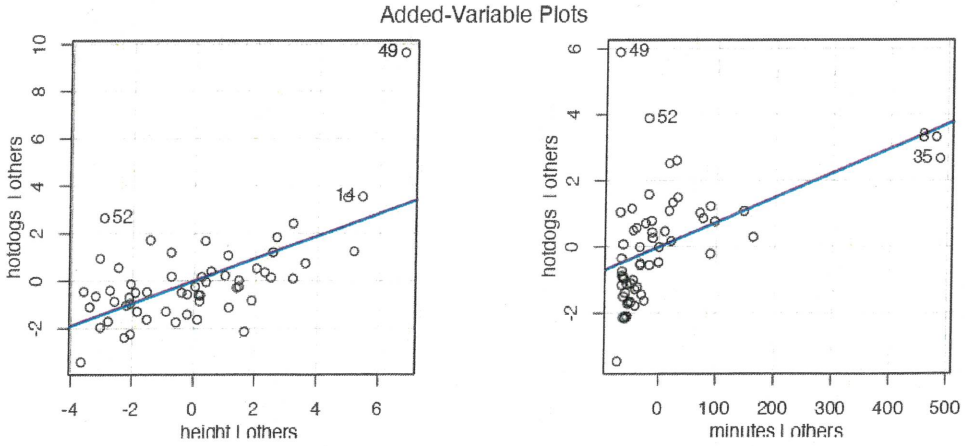

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-1.295392	5.467711	-0.237	0.813602
`height`	0.472892	0.080352	5.885	2.46e-07
`minutes`	0.007417	0.001871	3.963	0.000215

Other Summaries	Value
Breusch-Pagan p-value	0.071
Shapiro-Wilk p-value	2.43e-8
CV error	1.84623

Question 11: State the estimated value of β_1, and interpret its value in the context of the problem

$\beta_{1,height} = 0.472892$

among competitive eaters with the same height, each additional one inch in height is associated with 0.4729 higher average total number of hotdogs consumed.

* **Question 12:** **(Circle all that apply)** Which of the following are true?

(A) The 95% confidence interval of β_2 does **not** contain 0 because p is low we reject that $\beta_2 = 0$

(B) The 95% confidence interval of β_2 does contain 0

(C) There is an 81.4% probability that $\beta_0 = 0$

(D) $\beta_1 = 0.4729$

Question 13: Given what you know, provide a brief assessment of how well the data satisfy each of the following multiple linear regression assumptions:

(A) The relationships between the explanatory and response variables are linear

the data meets the linearity assumption because the added variable plots are somewhat linear.

(B) There is independence between each observation

the data meets the independence assumption because each competitive eater does not affect another.

(C) The errors are normally distributed

the data meets the normality assumption because the shapiro-wilk test gives a small p-value.

(D) The errors have constant variance (homoscedasticity)

the data does not meet the equal variance assumption because of the residuals vs fitted plots and the high p-value from the breusch-pagan test.

(E) There are no unduly influential points

the data meets this assumption because although there are observations with large residuals and low leverage and observations with high leverage and low residual, there are no observations with high influence because there are none with both high leverage and large residual.

Model 2

Model (2) below represents another possible multiple linear regression model for the same data. Here the variable `minutes` was log-transformed prior to fitting the model. Use model (2) and the output below to answer questions 14-16

$$\texttt{hotdogs}_i = \beta_0 + \beta_1 \texttt{height}_i + \beta_2 \log(\texttt{minutes}_i) + \epsilon_i; \quad \epsilon_i \overset{iid}{\sim} N(0, \sigma^2) \tag{2}$$

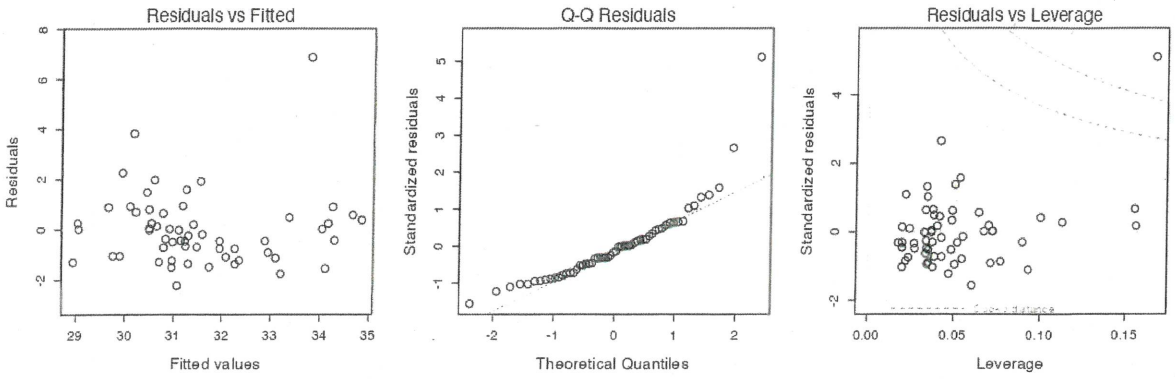

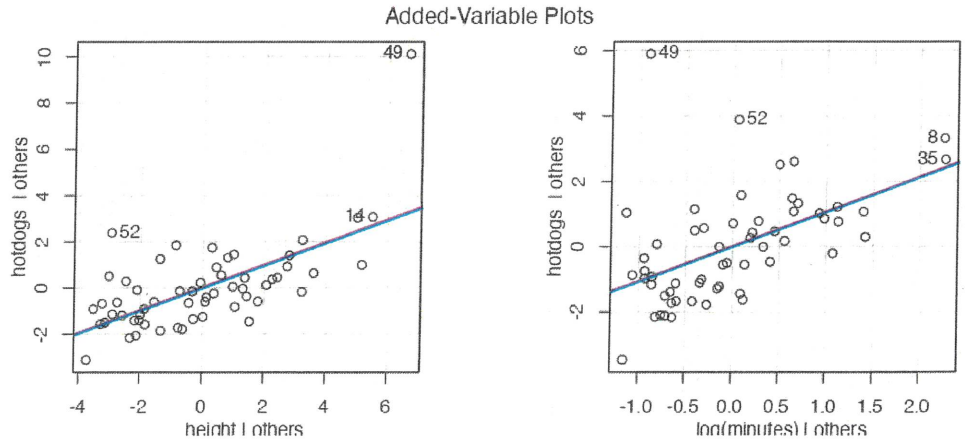

Coefficient	Estimate	Std. Error	t value	$\Pr(>\|t\|)$
Intercept	-6.05589	5.43620	-1.114	0.27
`height`	0.48834	0.07821	6.244	6.48e-08
log(`minutes`)	1.06702	0.24079	4.431	4.51e-05

Other Summaries	Value
Breusch-Pagan p-value	0.028
Shapiro-Wilk p-value	3.43e-5
CV RMSE	1.66702

Question 14: A Breusch-Pagan test for model (1) results in a p-value of 0.071 while a Breusch-Pagan test for model (2) results in a p-value of 0.028. A researcher concludes that model (1) is clearly superior to model (2). How would you respond?

I would disagree with the researcher because when using the breusch-pagan test, you want the p-value to be low. model two has a smaller p-value so we can be confident the data has equal variance.

Question 15: Explain what advantage(s) and disadvantage(s) model (2) provides over model (1)

an advantage model two provides is that when they used a level log transformation, it made it so the equal variance assumption was met. a disadvantage model two provides is that because they used a transformation, we can't directly compare the two models without back transforming.

Question 16: State the estimated value of β_2 and provide an interpretation in context

β_2 minutes = 1.06702

among competitive eaters, each additional minute per week spent practicing is associated with 1.06702 higher average total number of hotdogs consumed.

Use the following for questions 17-19

Your best friend reveals that he will be competing in the hotdog eating contest at the next state fair. You know that he is 68 inches tall and that he trains 79 minutes per week. Note that log(79) = 4.369. Use this information and the fit for model (2) to answer the following questions.

Question 17: Use model (2) to calculate the predicted number of hotdogs your friend can consume in 10 minutes

$\text{hotdogs} = -6.06589 + (68)(0.48834) + (1.06702)(4.369)$

$\text{hotdogs} = -6.06589 + 33.20712 + 4.666181 =$

$31.81 \rightarrow$ **31 hotdogs**

Question 18: To create a confidence or prediction interval for the expected number of hotdogs your friend can consume in 10 minutes, what is the correct number of degrees of freedom for the critical t^* used to calculate the margin of error?

(A) 55 $\quad 68 - 2 - 1 = 65$

(B) 57

(C) 56

(D) 48

* **Question 19:** You're given that $s\sqrt{x^*(X^\top X)^{-1}x^*} = 0.204$, $s\sqrt{1 + x^*(X^\top X)^{-1}x^*} = 1.485$, and $t^*_{df,0.975} = 2.004$. Give a 95% *prediction* interval for the number of hotdogs your friend can consume in 10 minutes.

$31 + (1.485)(2.004) = 31.81 + 2.97594 = 34.78594$

$31 - (1.489)(2.004) = 31.81 - 2.97094 = 28.83406$

(28.834, 34.7859)

Part III - Application

Questions 20-24 concern the following dataset:

In 1995, Australian zoologists conducted a study of local possum populations. They trapped dozens of specimens, determined their ages, and took measurements of several physical attributes from each possum. The process of determining a wild animal's age can be challenging and imprecise, and usually requires a detailed analysis of the animal's teeth, skin, skeletal structure, or DNA. Among other questions, researchers were interested in whether a possum's age could be predicted using only easy to obtain physical measurements. For this analysis we will focus on the 59 male possums from the sample. The variables in the dataset are as follows (see below). A multiple linear regression model to predict age using all other variables was fit to these data, resulting in the following model summary and coefficient estimates (see below)

Variable	Description
head_length	in mm
skull_width	in mm
total_length	in cm
tail_length	in cm
foot_length	in mm
ear_conch	ear conch circumference (in mm)
eye	width of right eye (in mm)
chest	chest girth (in cm)
belly	belly girth (in cm)
age	in years

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -13.22377    7.34105  -1.801   0.0778 .
head_length    0.07077    0.10120   0.699   0.4877
skull_width    0.16559    0.10599   1.562   0.1247
total_length  -0.04752    0.10167  -0.467   0.6423
tail_length   -0.05049    0.17030  -0.296   0.7681
foot_length   -0.21499    0.10733  -2.003   0.0507 .
ear_conch      0.15730    0.10231   1.538   0.1306
eye            0.43736    0.22549   1.940   0.0582 .
chest          0.01883    0.18154   0.104   0.9178
belly          0.20983    0.11580   1.812   0.0761 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.655 on 49 degrees of freedom
Multiple R-squared:  0.3538,    Adjusted R-squared:  0.2352
F-statistic: 2.982 on 9 and 49 DF,  p-value: 0.006552
```

A correlation plot (shown below) of these variables points towards some level of multicollinearity in the data. The maximum VIF among these variables is approximately 4 with an average VIF of approximately 3.

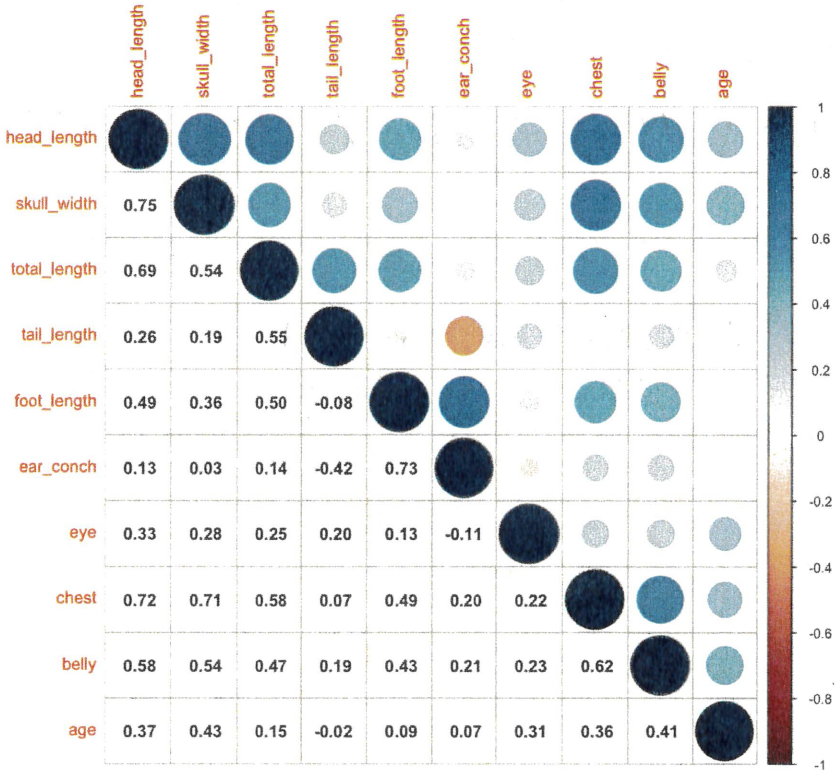

Question 20:

(a) Using the information from the summary table above, produce a 95% confidence interval for the coefficient on `skull_width`. The 0.975 critical value of a t-distribution with 49 degrees of freedom is 2.01. You may report your answer in the form $a \pm b * c$.

0.10599 ± 1.562 × 0.10599

0.10599 ± 1.562 × 0.10599

(b) A 95% confidence interval for $\beta_{\texttt{belly}}$, the coefficient on belly width, is (-0.023, 0.443). Interpret this interval within the context of the data.

We are 95% confident that the average coefficent for belly width for all male possums is between -0.023 and 0.443 cm.

Question 21: Note from the model summary that while the overall F-test of model significance produces a p-value of 0.0065, none of the individual model coefficients are found to have a significant effect at the $\alpha = 0.05$ level. What do you think is driving this "contradiction"? (It's not actually a contradiction.)

this "contradiction" of the overall F-test being significant, while none of the individual model coefficients are significant is because when all combined, you can't tell which predictors have more or less of an impact.

Question 22: Briefly describe how the "best subsets" algorithm using AIC as its criterion performs variable selection.

the best subsets algorithm using AIC as its criterion performs variable selection by rewarding fit and penalizing the number of predictors. the lower the number for AIC, the better.

Question 23: Best subsets using AIC as a criterion identified the following model with five variables:
`lm(age ~ skull_width + foot_length + ear_conch + eye + belly, data = possum)`
This model produces the following summary and coefficient estimates:

```
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.76661    4.97828  -2.966  0.00451 **
skull_width   0.20221    0.07990   2.531  0.01439 *
foot_length  -0.22936    0.09385  -2.444  0.01789 *
ear_conch     0.18239    0.08653   2.108  0.03979 *
eye           0.45001    0.21687   2.075  0.04285 *
belly         0.21405    0.10316   2.075  0.04287 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.609 on 53 degrees of freedom
Multiple R-squared:  0.3396,    Adjusted R-squared:  0.2773
F-statistic:  5.45 on 5 and 53 DF,  p-value: 0.0004056
```

vs 0.2362

Is this model guaranteed to have the highest adjusted-R^2 among all possible models.

(A) TRUE

(B) FALSE

Question 24: Suppose we want to compare the out-of-sample predictive performance of the full model and the model chosen by best subsets, and ultimately select the model with the smaller cross-validated RMSE.

Cross-validation using repeated random training-test splits can be structured around a for loop over some number of iterations (let's say 100), with some setup happening prior to the start of the loop, the bulk of the code within the loop, and some code after the loop has ended to process the results. To be clear, we **ARE NOT** using k-fold cross validation.

For each of the following steps in the cross-validation process, identify (by writing in the blank space) whether it occurs:

(A) prior to the for loop

(B) within the for loop

(C) after the for loop

(D) never in the CV process

C Average RMSE across all test-training splits and report

D Average the predictions and record their value

D Choose the model with the largest RMSE

C Choose the model with the smallest RMSE

C Compare the predictions from each model (that was fitted to training data) to the observed test data and compute RMSE

C Compute each model's AIC and BIC

B Fit each model using the testing data

B Fit each model using the training data

A Initialize vectors for each model that you will use to store their RMSEs

A Refit each model using the full data set and compare the coefficient estimates to the true values

B Split the data into test and training sets

B Use the models fitted to the training data to get predictions for the test data

END OF EXAM